

(单选题) Q1.A solutions architect is optimizing a website for an upcoming musical event. Videos of the performances will be streamed in real time and then will be available on demand. The event is expected to attract a global online audience. Which service will improve the performance of both the real-time and on-demand streaming ?

- A.Amazon CloudFront
- B.AWS Global Accelerator
- C.Amazon Route 53
- D.Amazon S3 Transfer Acceleration

(单选题) Q2.Organizers for a global event want to put daily reports online as static HTML pages. The pages are expected to generate millions of views from users around the world. The files are stored in an Amazon S3 bucket. A solutions architect has been asked to design an efficient and effective solution. Which action should the solutions architect take to accomplish this ?

- A.Generate presigned URLs for the files
- B.Use cross-Region replication to all Regions
- C.Use the geoproximity feature of Amazon Route 53
- D.Use Amazon CloudFront with the S3 bucket as its origin

(单选题) Q3.A company has a multi-tier application that runs six front-end web servers in an Amazon EC2 Auto Scaling group in a single Availability Zone behind an Application Load Balancer (ALB). A solutions architect needs to modify the infrastructure to be highly available without modifying the application. Which architecture should the solutions architect choose that provides high availability ?

- A.Create an Auto Scaling group that uses three instances across each of two Regions
- B.Modify the Auto Scaling group to use three instances across each of two Availability Zones
- C.Create an Auto Scaling template that can be used to quickly create more instances in another Region
- D.Change the ALB in front of the Amazon EC2 instances in a round-robin configuration to balance traffic to the web tier

(单选题) Q4.A bicycle sharing company is developing a multi-tier architecture to track the location of its bicycles during peak operating hours.The company wants to use these data points in its existing analytics platform A solutions architect must determine the most viable multi-tier option to support this architecture.The data points must be accessible from the REST API.Which action meets these requirements for storing and retrieving location data ?

- A.Use Amazon Athena with Amazon S3
- B.Use Amazon API Gateway with AWS Lambda
- C.Use Amazon Quick Sight with Amazon Redshift
- D.Use Amazon API Gateway with Amazon Kinesis Data Analytics

(单选题) Q5.A solutions architect is designing the cloud architecture for a new application being deployed on AWS.The process should run in parallel while adding and removing application nodes as needed based on the number of jobs to be processed.The processor application is stateless.The solutions architect must ensure that the application is loosely coupled and the job items are durably stored.Which design should the solutions architect use ?

- A.Create an Amazon SNS topic to send the jobs that need to be processed.Create an Amazon Machine Image (AMI) that consists of the processor application.Create a launch configuration that uses the AMI.Create an Auto Scaling group using the launch configuration.Set the scaling policy for the Auto Scaling group to add and remove nodes based on CPU usage
- B.Create an Amazon SQS queue to hold the jobs that need to be processed.Create an Amazon Machine Image (AMI) that consists of the processor application.Create a launch configuration that uses the AMI.Create an Auto Scaling group using the launch configuration.Set the scaling policy for the Auto Scaling group to add and remove nodes based on network usage
- C.Create an Amazon SQS queue to hold the jobs that need to be processed.Create an Amazon Machine Image (AMI) that consists of the processor application.Create a launch template that uses the AMI.Create an Auto Scaling group using the launch template.Set the scaling policy for the Auto Scaling group to add and remove nodes based on the number of items in the SQS queue
- D.Create an Amazon SNS topic to send the jobs that need to be processed.Create an Amazon Machine Image (AMI) that consists of the processor application.Create a

launch template that uses the AMI.Create an Auto Scaling group using the launch template.Set the scaling policy for the Auto Scaling group to add and remove nodes based on the number of messages published to the SNS topic.

(单选题) Q6.A security team wants to limit access to specific services or actions in all of the team's AWS accounts.All accounts belong to a large organization in AWS Organizations.The solution must be scalable and there must be a single point where permissions can be maintained.What should a solutions architect do to accomplish this ?

- A.Create an ACL to provide access to the services or actions.
- B.Create a security group to allow accounts and attach it to user groups
- C.Create cross-account roles in each account to deny access to the services or actions.
- D.Create a service control policy in the root organizational unit to deny access to the

(单选题) Q7.A company hosts an application on multiple Amazon EC2 instances.The application processes messages from an Amazon SQS queue writes to an Amazon RDS table and deletes the message from the queue Occasional duplicate records are found in the RDS table.The SQS queue does not contain any duplicate messages.What should a solutions architect do to ensure messages are being processed once only ?

- A.Use the CreateQueue API call to create a new queue
- B.Use the AddPermission API call to add appropriate permissions
- C.Use the ReceiveMessage API call to set an appropriate wait time.
- D.Use the ChangeMessageVisibility API call to increase the visibility timeout

(单选题) Q8.An Amazon EC2 administrator created the following policy associated with an IAM group containing several users.

```
{  
  "Version": "2020-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Action": "ec2:TerminateInstance",
```

```
"Resource": "",
"Condition": {
  "IpAddress": {
    "aws:SourceIP": "10.100.100.0/24"
  }
}

{
  "aws:SourceIP": "10.100.100.0/24",
  "Effect": "Deny", "Action": "ec2:",
  "resource": "*",
  "condition": {
    "StringNotEquals": {
      "ec2:Region": "us-east-1"
    }
  }
}
```

What is the effect of this policy ?

- A. Users can terminate an EC2 instance in any AWS Region except us-east-1.
- B. Users can terminate an EC2 instance with the IP address 10.100.100.1 in the us-east-1 Region.
- C. Users can terminate an EC2 instance in the us-east-1 Region when the user's source IP is 10.100.100.254.
- D. Users cannot terminate an EC2 instance in the us-east-1 Region when the user's source IP is 10.100.100.254.

(单选题) Q9. An application runs on Amazon EC2 instances across multiple Availability Zones. The instances run in an Amazon EC2 Auto Scaling group behind an Application Load Balancer.

The application performs best when the CPU utilization of the EC2 instances is at or near 40%. What should a solutions architect do to maintain the desired performance across all instances in the group ?

- A. Use a simple scaling policy to dynamically scale the Auto Scaling group
- B. Use a target tracking policy to dynamically scale the Auto Scaling group
- C. Use an AWS Lambda function to update the desired Auto Scaling group capacity

D.Use scheduled scaling actions to scale up and scale down the Auto Scaling group

(多选题) Q10.A solutions architect is designing a two-tier web application.The application consists of a public-facing web tier hosted on Amazon EC2 in public subnets.The database tier consists of Microsoft SQL Server running on Amazon EC2 in a private subnet Security is a high priority for the company.How should security groups be configured in this situation ? (Select TWO)

A.Configure the security group for the web tier to allow inbound traffic on port 443 from 0.0.0.0/0

B.Configure the security group for the web tier to allow outbound traffic on port 443 from 0.0.0.0/0

C.Configure the security group for the database tier to allow inbound traffic on port 1433 from the security group for the web tier

D.Configure the security group for the database tier to allow outbound traffic on ports 443 and 1433 to the security group for the web tier

E.Configure the security group for the database tier to allow inbound traffic on ports 443 and 1433 from the security group for the web tier

(单选题) Q11.A company is hosting a web application on AWS using a single Amazon EC2 instance that stores user-uploaded documents in an Amazon EBS volume.For better scalability and availability the company duplicated the architecture and created a second EC2 instance and EBS volume in another Availability Zone : placing both behind an Application Load Balancer.After completing this change users reported that each time they refreshed the website they could see one subset of their documents or the other but never all of the documents at the same time.

What should a solutions architect propose to ensure users see all of their documents at once"

A.Copy the data so both EBS volumes contain all the documents.

B.Configure the Application Load Balancer to direct a user to the server with the documents.

C.Copy the data from both EBS volumes to Amazon EFS.Modify the application to

save new documents to Amazon EFS.

D.Configure the Application Load Balancer to send the request to both servers.Return each document from the correct server.

(单选题) Q12.An application runs on Amazon EC2 instances in private subnets.The application needs to access an Amazon DynamoDB table.What is the MOST secure way to access the table while ensuring that the traffic does not leave the AWS network ?

- A.Use a VPC endpoint for DynamoDB.
- B.Use a NAT gateway in a public subnet.
- C.Use a NAT instance in a private subnet.
- D.Use the internet gateway attached to the VPC.

(单选题) Q13.A media company is evaluating the possibility of moving its systems to the AWS Cloud.The company needs at least 10 TB of storage with the maximum possible I/O performance for video processing.300 TB of very durable storage for storing media content, and 900 TB of storage to meet requirements for archival media that is not in use anymore.Which set of services should a solutions architect recommend to meet these requirements ?

- A.Amazon EBS for maximum performance, Amazon S3 for durable data storage, and Amazon S3 Glacier for archival storage
- B.Amazon EBS for maximum performance.Amazon EFS for durable data storage, and Amazon S3 Glacier for archival storage
- C.Amazon EC2 instance store for maximum performance, Amazon EFS for durable data storage, and Amazon S3 for archival storage
- D.Amazon EC2 instance store for maximum performance, Amazon S3 for durable data storage, and Amazon S3 Glacier for archival storage

(单选题) Q14.A company has a large Microsoft SharePoint deployment running on-premises that requires Microsoft Windows shared file storage.The company wants to migrate this workload to the AWS Cloud and is considering various storage options.The storage solution must be highly available and integrated with Active

Directory for access control.Which solution will satisfy these requirements ?

- A.Configure Amazon EFS storage and set the Active Directory domain for authentication.
- B.Create an SMB file share on an AWS Storage Gateway file gateway in two Availability Zones.
- C.Create an Amazon S3 bucket and configure Microsoft Windows Server to mount it as a volume.
- D.Create an Amazon FSx for Windows File Server file system on AWS and set the Active Directory domain for authentication.

(单选题) Q15.A company runs an application using Amazon ECS.The application creates resized versions of an original image and then makes Amazon S3 API calls to store the resized images in Amazon S3.How can a solutions architect ensure that the application has permission to access Amazon S3 ?

- A.Update the S3 role in AWS IAM to allow read/write access from Amazon ECS, and then relaunch the container.
- B.Create an IAM role with S3 permissions, and then specify that role as the taskRoleArn in the task definition.
- C.Create a security group that allows access from Amazon ECS to Amazon S3, and update the launch configuration used by the ECS cluster.
- D.Create an IAM user with S3 permissions, and then relaunch the Amazon EC2 instances for the ECS cluster while logged in as this account.

(单选题) Q16.A company's web application is running on Amazon EC2 instances behind an Application Load Balancer.The company recently changed its policy, which now requires the application to be accessed from one specific country only.Which configuration will meet this requirement ?

- A.Configure the security group for the EC2 instances.
- B.Configure the security group on the Application Load Balancer.
- C.Configure AWS WAF on the Application Load Balancer in a VPC.
- D.Configure the network ACL for the subnet that contains the EC2 instances.

(单选题) Q17.A company's website provides users with downloadable historical performance reports.The website needs a solution that will scale to meet the company's website demands globally.The solution should be cost effective, limit the ? provisioning of Into and provide the fastest possible response time.Which combination should a solutions architect recommend to meet these requirements ?

- A.Amazon CloudFront and Amazon S3
- B.AWS Lambda and Amazon Dynamo
- C.Application Load Balancer with Amazon EC2 Auto Scaling
- D.Amazon Route 53 with internal Application Load Balances

(单选题) Q18.A company has an Amazon EC2 instance running on a private subnet that needs to access a public websites to download patches and updates.The company does not want external websites to see the EC2 instance IP address or initiate connection to it.How can a solution architect achieve this objective ?

- A.Create a site-to-site VPN connection between the private subnet and the network in which the public site is deployed
- B.Create a NAT gateway in a public subnet Route outbound traffic from the private subnet through the NAT gateway
- C.Create a network ACL for the private subnet where the EC2 instance deployed only allows access from the IP address range of the public website
- D.Create a security group that only allows connections from the IP address range of the public website.Attach the security group to the EC2 instance.

(单选题) Q19.A development team needs to host a website that will be accessed by other teams.The website contents.consist of HTML.CSS, client side JavaScript, and images.Which method is the MOST cost-effective for hosting the website ?

- A.Containerize the website and host it in AWS Fargate
- B.Create an Amazon S3 bucket and host the website there.
- C.Deploy a web server on an Amazon EC2 instance to host the website.
- D.Configure an Application Load Balancer with an AWS Lambda target that uses the Express is framework

(单选题) Q20. Company is designing a website that uses an Amazon S3 bucket to store static images. The company wants all future requests have faster response times while reducing both latency and cost. Which service configuration should a solutions architect recommend ?

- A. Deploy a NAT server in front of Amazon S3.
- B. Deploy Amazon CloudFront in front of Amazon S3.
- C. Deploy a Network Load Balancer in front of Amazon S3.
- D. Configure Auto Scaling to automatically adjust the capacity of the website.

(单选题) Q21. A company's dynamic website is hosted using on-premises servers in the United States. The company is launching its product in Europe and it wants to optimize site loading times for new European users. The site's backend must remain in the United States. The product is being launched in a few days, and an immediate solution is needed. What should the solutions architect recommend ?

- A. Launch an Amazon EC2 instance in us-east-1 and migrate the site to it
- B. Move the website to Amazon S3. Use cross-Region replication between Regions.
- C. Use Amazon CloudFront with a custom origin pointing to the on-premises servers
- D. Use an Amazon Route 53 geoproximity routing policy pointing to on-premises servers

(单选题) Q22. A company has recently updated its internal security standards. The company must now ensure all Amazon S3 buckets and Amazon Elastic Block Store (Amazon EBS) volumes are encrypted with keys created and periodically rotated by internal security specialists. The company is looking for a native, software-based AWS service to accomplish this goal. What should a solutions architect recommend as a solution ?

- A. Use AWS Secrets Manager with customer master keys (CMKs) to store master key material and apply a routine to create a new CMK periodically and replace it in AWS Secrets Manager.
- B. Use AWS Key Management Service (AWS KMS) with customer master keys (CMKs) to store master key material and apply a routine to re-create a new key periodically and replace it in AWS KMS.

C.Use an AWS CloudHSM cluster with customer master keys (CMKs) to store master key material and apply a routine a re-create a new key periodically and replace it in the CloudHSM cluster nodes.

D.Use AWS Systems Manager Parameter Store with customer master keys (CMKs) keys to store master key material and apply a routine to re-create a new periodically and replace it in the Parameter Store.

(单选题) Q23.A solution architect must design a solution that uses Amazon CloudFront with an Amazon S3 to store a static website.The company security policy requires that all websites traffic be inspected by AWS WAF.How should the solution architect company with these requirements ?

A.Configure an S3 bucket policy to accept requests coming from the AWS WAF Amazon Resource Name (ARN) only

B.Configure Amazon CloudFront to forward all incoming requests to AWS WAF before requesting content from the S3 origin,

C.Configure a security group that allows Amazon CloudFront IP addresses to access Amazon S3 only Associate Aws WAF to CloudFront.

D.Configure Amazon CloudFront and Amazon S3 to use an origin access identity (OAI) to restrict access to the S3 bucket.Enable Aws WAF on the distribution.

(单选题) Q24.A solutions architect is using Amazon S3 to design the storage architecture of a new digital media application.The media files must be resilient to the loss of an Availability Zone Some files are accessed frequently while other files are rarely accessed in an unpredictable pattern.The solutions architect must minimize the costs of storing and retrieving the media files.Which storage option meets these requirements ?

A.S3 Standard

B.S3 Intelligent-Tiering

C.S3 Standard-Infrequent Access (S3 Standard-IA)

D.S3 One Zone-Infrequent Access (S3 One Zone-IA)

Answer:B

(单选题) Q25.An application allows users at a company's headquarters to access product data.The product data is stored in an Amazon RDS MySQL DB instance.The operations team has isolated an application performance slowdown and wants to separate read traffic from write traffic.A solutions architect needs to optimize the application's performance quickly.What should the solutions architect recommend ?

A.Change the existing database to a Multi-AZ deployment.Serve the read requests from the primary Availability Zone.

B.Change the existing database to a Multi-AZ deployment.Serve the read requests from the secondary Availability Zone.

C.Create read replicas for the database.Configure the read replicas with half of the compute and storage resources as the source database.

D.Create read replicas for the database.Configure the read replicas with the same compute and storage resources as the source database.

(单选题) Q26.A company that hosts its web application on AWS wants to ensure all Amazon EC2 instances.Amazon RDS DB instances and Amazon Redshift clusters are configured with tags.

The company wants to minimize the effort of configuring and operating this check.What should a solutions architect do to accomplish this"

A.Use AWS Config rules to define and detect resources that are not property tagged

B.Use Cost Explorer to display resources that are not properly tagged Tag those resources manually.

C.Write API calls to check all resources for proper tag allocation.Periodically run the code on an EC2 instance.

D.Write API calls to check all resources for proper tag allocation.Schedule an AWS Lambda function through Amazon CloudWatch to periodically run the code

(单选题) Q27.A company stores call recordings on a monthly basis Statistically, the recorded data may be referenced randomly within a year but accessed rarely after 1 year.Files that are newer than 1 year old must be queried and retrieved as quickly as possible.A delay in retrieving older files is acceptable A

solutions architect needs to store the recorded data at a minimal cost.Which solution is MOST cost-effective ?

- A.Store individual files in Amazon S3 Glacier and store search metadata in object tags created in S3 Glacier.Query S3 Glacier tags and retrieve the files from S3 Glacier
- B.Store individual files in Amazon S3 Use lifecycle policies to move the files to Amazon S3 Glacier after 1 year.Query and retrieve the files from Amazon S3 or S3 Glacier.
- C.Archive individual files and store search metadata for each archive in Amazon S3.Use lifecycle policies to move the files to Amazon S3 Glacier after 1 year.Query and retrieve the files by searching for metadata from Amazon S3
- D.Archive individual files in Amazon S3.Use lifecycle policies to move the files to Amazon S3 Glacier after 1 year.Store search metadata in Amazon DynamoDB Query the files from DynamoDB and retrieve them from Amazon S3 or S3 Glacier

(单选题) Q28.A company has a highly dynamic batch processing job that uses many Amazon EC2 instances to complete it.The job is stateless in nature, can be started and stopped at any given time with no negative impact, and typically takes upwards of 60 minutes total to complete.The company has asked a solutions architect to design a scalable and cost-effective solution that meets the requirements of the job.What should the solutions architect recommend ?

- A.Implement EC2 Spot Instances
- B.Purchase EC2 Reserved Instances
- C.Implement EC2 On-Demand Instances
- D.Implement the processing on AWS Lambda

(单选题) Q29.A company runs a production application on a fleet of Amazon EC2 instances.The application reads the data from an Amazon SQS queue and processes the messages in parallel.The message volume is unpredictable and often has intermittent traffic.This application should continually process messages without any downtime Which solution meets these requirements MOST cost-effectively ?

- A.Use Spot Instances exclusively to handle the maximum capacity required
- B.Use Reserved Instances exclusively to handle the maximum capacity required
- C.Use Reserved Instances for the baseline capacity and use Spot Instances to handle

additional capacity

D.Use Reserved instances for the baseline capacity and use On-Demand Instances to handle additional capacity

(单选题) Q30.A company provides an API to its users that automates inquiries for tax computations based on item prices.The company experiences a larger number of inquiries during the holiday season only that cause slower response times.A solution architect needs to design a solution that is scalable and elastic.What should the solutions architect do to accomplish this ?

A.Provide an API hosted on an Amazon EC2 instance.The EC2 instance performs the required computations when the API request is made.

B.Design a REST API using Amazon API Gateway that accepts the item names, API Gateway passes item names to AWS Lambda for tax computations.

C.Create an Application Load Balancer that has two Amazon EC2 instances behind it.The EC2 instances will compute the tax on the received item names.

D.Design a REST API using Amazon API Gateway that connects with an API hosted on an Amazon EC2 instance, API Gateway accepts and passes the item names to the EC2 instance for tax computations.

(单选题) Q31.A solutions architect is designing a VPC with public and private subnets.The VPC and subnets use IP 4 CIDR blocks.There is one public subnet and one private subnet in each of three Availability Zone (AZs) for high availability.An internet gateway is used to provide internet access for the public subnets.The private subnets require access to the internet to allow Amazon EC2 instances to download software updates.What should the solutions architect do to enable Internet access for the private subnets ?

A.Create three NAT gateways, one for each public subnet in each AZ.Create a private route table for each AZ that forwards non-VPC traffic to the NAT gateway in its AZ. B.Create three NAT gateways, one for each private subnet in each AZ.Create a private route table for each AZ that forwards non-VPC traffic to the NAT gateway in its AZ. C.Create second internet gateway on one of the private subnets.Update the route table for the private subnets that forward non-VPC traffic to the private internet

gateway.

D.Create an egress-only internet gateway on one of the public subnets.Update the route table for the private subnets that forward non-VPC traffic to the egress-only internet gateway.

(单选题) Q32.A company is preparing to launch a public-facing web application in the AWS Cloud.The architecture consists of Amazon EC2 instances within a VPC behind an Elastic Load Balancer (ELB) .A third-party service is used for the DNS.The company's solutions architect must recommend a solution to detect and protect against large-scale DDoS attacks.Which solution meets these requirements ?

- A.Enable Amazon Guard Duty on th account
- B.Enable Amazon Inspector on the EC2 instances
- C.Enable AWS Shield and assign Amazon Route 53 to it.
- D.Enable AWS Shield Advancd and assign the ELB to it.

(单选题) Q33.A company has a dynamic web application hostes on two Amazon EC2 instances.The company has its own SSL certificate, which is on each instance to perform SSL termination.

There has been an increase in traffic recently, and the operations team determined that SSL encryption and decryption is causing the compute capacity of the web servers to reach their maximum limit.What should a solutions architect do to increase the application's performance ?

- A.Create a new SSL certificate using AWS Certificate Manager (ACM) .Install the ACM certificate on each instance.
- B.Create an Amazon S3 bucket Migrate the SSL certificate to the S3 bucket.Configure the EC2 instances to reference the bucket for SSL termination.
- C.Create another EC2 instance as a proxy server.Migrate the SSL certificate to the new instance and configure it to direct connctions to the existing EC2 instances.
- D.Import the SSL certificate into AWS Crtificate Manager (ACM) .Create an Application Load Balancer with an HTTPS listener that uses the SSL certificate from ACM.

(单选题) Q34.A company maintains a searchable repository of items on its website.The data is stored in an Amazon RDS for MySQL database table that contains over 10 million rows.The database has 2 TB of General Purpose SSD (gp2) storage.There are millions of updates against this data every day through the company's website.The company has noticed some operations are taking 10 seconds or longer, and has determined that the database storage performance is bottleneck.Which solution addresses the performance issues ?

- A.Change the storage type to Provisioned IOPS SSD (io1) .
- B.Change the instance to a memory-optimized instance class.
- C.Change the instance to a burstable performance DB instance class.
- D.Enable Multi-AZ RDS read replicas with MySQL native asynchronous replication.

(单选题) Q35.A company has an application that generates a large number of files, each approximately 5 MB in size.The files are stored in Amazon S3.Company policy requires the files to be stored for 4 years before they can be deleted.Immediate accessibility is always required as the files contain critical business data that is not easy to reproduce.The files are frequently accessed in the first 30 days of the object creation but are rarely accessed after the first 30 days.Which storage solution is MOST cost effective ?

- A.Create an S3 bucket lifecycle policy to move files from S3 Standard to S3 Glacier 30 days from object creation.Delete the files 4 years after the object creation.
- B.Create an S3 bucket lifecycle policy to move files from S3 Standard to S3 One Zone-Infrequent Access (S3 One Zone-IA) 30 days from object creation.Delete the files 4 years after the object creation.
- C.Create an S3 bucket lifecycle policy to move files from S3 Standard to S3 Standard-Infrequent Access (S3 Standard-IA) 30 days from object creation.Delete the files 4 years after the object creation.
- D.Create an S3 bucket lifecycle policy to move files from S3 Standard to S3 Standard-Infrequent Access (S3 Standard-IA) 30 days from object creation.Move the file to S3 Glacier 4 years after object creation.

(单选题) Q36.A company is preparing to store confidential data in Amazon S3.For

compliance reasons, the data must be encrypted at rest. Encryption key usage must be logged for auditing purposes. Key must be rotated every year. Which solution meets these requirements and is the MOST operationally efficient?

- A. Server-side encryption with customer-provided keys (SSE-C)
- B. Server-side encryption with Amazon S3 managed keys (SSE-S3)
- C. Server-side encryption with AWS KMS (SSE-KMS) customer master keys (CMKs) with manual rotation.
- D. Server-side encryption with AWS KMS (SSE-KMS) customer master keys (CMKs) with automatic rotation.

(单选题) Q37. A company wants to migrate its MySQL database from on-premises to AWS. The company recently experienced a database outage that significantly impacted the business. To ensure this does not happen again, the company wants a reliable database solution on AWS that minimizes data loss and stores every transaction on at least two nodes. Which solution meets these requirements?

- A. Create an Amazon RDS DB instance with synchronous replication to three nodes in three Availability Zones.
- B. Create an Amazon RDS MySQL DB instance with Multi-AZ functionality enabled to synchronously replicate the data.
- C. Create an Amazon RDS MySQL DB instance with Multi-AZ and create a read replica in a separate AWS Region that synchronously replicates the data.
- D. Create an Amazon EC2 instance with a MySQL engine installed that triggers an AWS Lambda function to synchronously replicate the data to an Amazon RDS MySQL DB instance.

(单选题) Q38. A company has an automobile sales website that stores its listings in a database on Amazon RDS. When an automobile is sold, the listing needs to be removed from the website and the data must be sent to multiple target systems. Which design should a solutions architect recommend?

- A. Create an AWS Lambda function triggered when the database on Amazon RDS is updated to send the information to an Amazon Simple Queue Service (Amazon SQS) queue for the targets to consume.

B.Create an AWS Lambda function triggered when the database on Amazon RDS is updated to send the information to an Amazon Simple Queue Service (Amazon SQS) FIFO queue for the targets to consume.

C.Subscribe to an RDS event notification and send an Amazon Simple Queue Service (Amazon SQS) queue fanned out to multiple Amazon Simple Notification Service (Amazon SNS) topics.Use AWS Lambda functions to update the targets.

D.Subscribe to an RDS event notification and send an Amazon Simple Notification Service (Amazon SNS) topic fanned out to multiple Amazon Simple Queue Service (Amazon SQS) queues Use AWS Lambda functions to update the targets.

(单选题) Q39.A company wants to build a scalable key management infrastructure to support developers who need to encrypt data in their applications.What should a solutions architect do to reduce the operational burden ?

A.Use multi-factor authentication (MFA) to protect the encryption keys

B.Use AWS Key Management Service (AWS KMS) to protect the encryption keys

C.Use AWS Certificate Manager (ACM) to create, store and assign the encryption keys

D.Use an IAM policy to limit the scope of users who have access permissions to protect the encryption keys

(多选题) Q40.A company is running a multi-tier web application on premises.The web application is containerized and runs on a number of Linux hosts connected to a PostgreSQL database that contains user records.The operational overhead of maintaining the infrastructure and capacity planning is limiting the company's growth.A solutions architect must improve the application's infrastructure.Which combination of actions should the solutions architect take to accomplish this ? (Select TWO.)

A.Migrate the PostgreSQL database to Amazon Aurora

B.Migrate the web application to be hosted on Amazon EC2 instances.

C.Set up an Amazon CloudFront distribution for the web application content.

D.Set up Amazon ElastiCache between the web application and the PostgreSQL database

E.Migrate the web application to be hosted on AWS Fargate with Amazon Elastic Container Service (Amazon ECS)

(单选题) Q41.A company runs an online marketplace web application on AWS.The application serves hundreds of thousands of users during peak hours.The company needs a scalable, near-real-time solution to share the details of millions of financial transactions with several other internal applications.Transactions also need to be processed to remove sensitive data before being stored in a document database for low-latency retrieval.What should a solutions architect recommend to meet these requirements ?

A.Store the transactions data into Amazon DynamoDB.Set up a rule in DynamoDB to remove sensitive data from every transaction upon write.Use DynamoDB Streams to share the transactions data with other applications.

B.Stream the transactions data into Amazon Kinesis Data Firehose to store data in Amazon DynamoDB and Amazon S3.Use AWS Lambda integration with Kinesis Data Firehose to remove sensitive data.Other applications can consume the data stored in Amazon S3.

C.Stream the transactions data into Amazon Kinesis Data Streams.Use AWS Lambda integration to remove sensitive data from every transaction and then store the transactions data in Amazon DynamoDB.Other applications can consume the transactions data off the Kinesis data stream.

D.Store the batched transactions data in Amazon S3 as files.Use AWS Lambda to process every file and remove sensitive data before updating the files in Amazon S3.The Lambda function then stores the data in Amazon DynamoDB.Other applications can consume transaction files stored in Amazon S3.

(多选题) Q42.A company uses an Amazon S3 bucket to store static images for its website.The company configured permissions to allow access to Amazon S3 objects by privileged users only.

What should a solutions architect do to protect against data loss ? (Select TWO.)

A.Enable versioning on the S3 bucket.

B.Enable access logging on the S3 bucket.

- C.Enable server-side encryption on the S3 bucket.
- D.Configure an S3 lifecycle rule to transition objects to Amazon S3 Glacier
- E.Use MFA Delete to require multi-factor authentication to delete an object.

(单选题) Q43.A solutions architect is developing a multiple-subnet VPC architecture.The solution will consist of six subnets in two Availability Zones.The subnets are defined as public, private, and dedicated for databases.Only the Amazon EC2 instances running in the private subnets should be able to access a database.Which solution meets these requirements ?

- A.Create a new route table that excludes the route to the public subnets'CIDR blocks.Associate the route table to the database subnets.
- B.Create a security group that denies ingress from the security group used by instances in the public subnets.Attach the security group to an Amazon RDS DB instance.
- C.Create a security group that allows ingress from the security group used by instances in the private subnets.Attach the security group to an Amazon RDS DB instance.
- D.Create a new peering connection between the public subnets and the private subnets.Create a different peering connection between the private subnets and the database subnets.

(单选题) Q44.A solutions architect is designing the architecture of a new application being deployed to the AWS Cloud.The application will run on Amazon EC2 On-Demand Instances and will automatically scale across multiple Availability Zones.The EC2 instances will scale up and down frequently throughout the day.An Application Load Balancer (ALB) will handle the load distribution.The architecture needs to support distributed session data management.The company is willing to make changes to code if needed.What should the solutions architect do to ensure that the architecture supports distributed session data management ?

- A.Use Amazon ElastiCache to manage and store session data.
- B.Use session affinity (sticky sessions) of the ALB to manage session data.
- C.Use Session Manager from AWS Systems Manager to manage the session.

D.Use the GetSessionToken API operation in AWS Security Token Service (AWS STS) to manage the session.

(单选题) Q45.A solutions architect is designing a solution that requires frequent updates to a website that is hosted on Amazon S3 with versioning enabled For compliance reasons, the older versions of the objects will not be accessed frequently and will need to be deleted after 2 years.What should the solutions architect recommend to meet these requirements at the LoWEsT cost ?

A.Use S3 batch operations to replace object tags Expire the objects based on the modified tags.

B.Configure an S3 Lifecycle policy to transition older versions of objects to S3 Glacier.Expire the objects after 2 years.

C.Enable S3 Event Notifications on the bucket that sends older objects to the Amazon Simple Queue Service (Amazon SQS) queue for further processing

D.Replicate older object versions to a new bucket.Use an S3 Lifecycle policy to expire the objects in the new bucket after 2 years.

(多选题) Q46.A company is storing sensitive user information in an Amazon S3 bucket.The company wants to provide secure access to this bucketfrom the application tier running on Amazon EC2 instances inside a VPC.Which combination of steps should a solutions architect take to accomplish this ? (Select TWO.)

A.Configure a VPC gateway endpoint for Amazon S3 within the VPC.

B.Create a bucket policy to make the objects in the S3 bucket public.

C.Create a bucket policy that limits access to only the application tier running in the VPC.

D.Create an IAM user with an S3 access policy and copy the IAM credentials to the EC2 instance.

D.Create a NAT instance and have the EC2 instances use the NAT instance to access the S3 bucket.

(单选题) Q47.A company is preparing to deploy a new serverless workload.A solutions architect needs to configure permissions for invoking an AWS Lambda

function. The function will be triggered by an Amazon EventBridge (Amazon CloudWatch Events) rule. Permissions should be configured using the principle of least privilege. Which solution will meet these requirements?

A. Add an execution role to the function with `lambda:InvokeFunction` as the action and `*` as the principal.

B. Add an execution role to the function with `lambda:InvokeFunction` as the action and `Service:events.amazonaws.com` as the principal.

C. Add a resource-based policy to the function with `lambda:*` as the action and `Service:events.amazonaws.com` as the principal.

D. Add a resource-based policy to the function with `lambda:invokeFunction` as the action and `Service:events.amazonaws.com` as the principal.

(单选题) Q48. A company has a service that produces event data. The company wants to use AWS to process the event data as it is received. The data is written in a specific order that must be maintained throughout processing. The company wants to implement a solution that minimizes operational overhead. How should a solutions architect accomplish this?

A. Create an Amazon Simple Queue Service (Amazon SQS) FIFO queue to hold messages. Set up an AWS Lambda function to process messages from the queue.

B. Create an Amazon Simple Notification Service (Amazon SNS) topic to deliver notifications containing payloads to process. Configure an AWS Lambda function as a subscriber.

C. Create an Amazon Simple Queue Service (Amazon SQS) standard queue to hold messages. Set up an AWS Lambda function to process messages from the queue independently.

D. Create an Amazon Simple Notification Service (Amazon SNS) topic to deliver notifications containing payloads to process. Configure an Amazon Simple Queue Service (Amazon SQS) queue as a subscriber.

(单选题) Q49. A company uses Amazon S3 to store its confidential audit documents. The S3 bucket uses bucket policies to restrict access to audit team IAM user credentials according to the principle of least privilege. Company managers are

worried about accidental deletion of documents in the S3 bucket and want a more secure solution.What should a solutions architect do to secure the audit documents ?

- A.Enable the versioning and MFA Delete features on the S3 bucket
- B.Enable multi-factor authentication (MFA) on the IAM user credentials for each audit team IAM user account.
- C.Add an S3 Lifecycle policy to the audit team's IAM user accounts to deny the s3 : DeleteObject action during audit dates.
- D.Use AWS Key Management Service (AWS KMS) to encrypt the S3 bucket and restrict audit team IAM user accounts from accessing the KMS key.

(单选题) Q50.A company receives 10 TB of instrumentation data each day from several machines located at a single factory.The data consists of JSON files stored on a storage area network (SAN) in an on-premises data center located within the factory.The company wants to send this data to Amazon S3 where it can be accessed by several additional systems that provide critical near-real-time analytics.A secure transfer is important because the data is considered sensitive.Which solution offers the MOST reliable data transfer ?

- A.AWS DataSync over public internet
- B.AWS DataSync over AWS Direct Connect
- C.AWS Database Migration Service (AWS DMS) over public internet
- D.AWS Database Migration Service (AWS DMS) over AWS Direct Connect

(单选题) Q51.A company owns an asynchronous API that is used to ingest use requests and, based on the request type, dispatch requests to the appropriate microservice for processing.The company is using Amazon API Gateway to deploy the API front end, and an AWS Lambda function that invokes Amazon DynamoDB to store user requests before dispatching them to the processing microservices.The company provisioned as much DynamoDB throughput as its budget allows, but the company is still experiencing availability issues and is losing user requests. What should a solutions architect do to address this issue without impacting existing users ?

- A.Add throttling on the API Gateway with server-side throttling limits
- B.Use DynamoDB Accelerator (DAX) and Lambda to buffer writes to DynamoDB
- C.Create a secondary index in DynamoDB for the label with the user requests.
- D.Use the Amazon Simple Queue Service (Amazon SQS) queue and Lambda to buffer writes to DynamoDB.

(单选题) Q52.A company has an application that ingests incoming messages.These messages are then quickly consumed by dozens of other applications and microservices.The number of messages varies drastically and sometimes spikes as high as 100 000 each second The company wants to decouple the solution and increase scalability.Which solution meets these requirements ?

- A.Persist the messages to Amazon Kinesis Data Analytics.All the applications will read and process the messages.
- B.Deploy the application on Amazon EC2 instances in an Auto Scaling group, which scales the number of EC2 instances based on CPU metrics.
- C.Write the messages to Amazon Kinesis Data Streams with a single shard.All applications will read from the stream and process the messages.
- D.Publish the messages to an Amazon Simple Notification Service (Amazon SNS) topic with one or more Amazon Simple Queue Service (Amazon SQS) subscriptions All applications then process the messages from the queues.

(单选题) Q53.A company wants to move a multi-tiered application from on premises to the AWS Cloud to improve the application's performance.The application consists of application tiers that communicate with each other by way of RESTful services.Transactions are dropped when one tier becomes overloaded.A solutions architect must design a solution that resolves these issues and modernizes the application.Which solution meets these requirements and is the MOST operationally efficient ?

- A.Use Amazon API Gateway and direct transactions to the AWS Lambda functions as the application layer.Use Amazon Simple Queue Service (Amazon SQS) as the communication layer between application services.
- B.Use Amazon CloudWatch metrics to analyze the application performance history

to determine the servers' peak utilization during the performance failures. Increase the size of the application server's Amazon EC2 instances to meet the peak requirements.

C. Use Amazon Simple Notification Service (Amazon SNS) to handle the messaging between application servers running on Amazon EC2 in an Auto Scaling group. Use Amazon CloudWatch to monitor the SNS queue length and scale up and down as required.

D. Use Amazon Simple Queue Service (Amazon SQS) to handle the messaging between application servers running on Amazon EC2 in an Auto Scaling group. Use Amazon CloudWatch to monitor the SQS queue length and scale up when communication failures are detected.

(单选题) Q54. A company recently started using Amazon Aurora as the data store for its global ecommerce application. When large reports are run, developers report that the ecommerce application is performing poorly. After reviewing metrics in Amazon CloudWatch, a solutions architect finds that the ReadIOPS and CPU Utilization metrics are spiking when monthly reports run. What is the MOST cost-effective solution?

- A. Migrate the monthly reporting to Amazon Redshift
- B. Migrate the monthly reporting to an Aurora Replica.
- C. Migrate the Aurora database to a larger instance class.
- D. Increase the Provisioned IOPS on the Aurora instance.

(单选题) Q55. A company has two applications: a sender application that sends messages with payloads to be processed and a processing application intended to receive the messages with payloads. The company wants to implement an AWS service to handle messages between the two applications. The sender application can send about 1,000 messages each hour. The messages may take up to 2 days to be processed. If the messages fail to process, they must be retained so that they do not impact the processing of any remaining messages. Which solution meets these requirements and is the MOST operationally efficient?

- A. Set up an Amazon EC2 instance running a Redis database. Configure both applications to use the instance. Store, process, and delete the

messages, respectively. B. Use an Amazon Kinesis data stream to receive the messages from the sender application, integrate the processing application with the Kinesis Client Library (KCL). C. Integrate the sender and processor applications with an Amazon Simple Queue Service (Amazon SQS) ; queue. Configure a dead-letter queue to collect the messages that failed to process.

D. Subscribe the processing application to an Amazon Simple Notification Service (Amazon SNS) topic to receive notifications to process. Integrate the sender application to write to the SNS topic.

(单选题) Q56. A company wants to use high performance computing (HPC) infrastructure on AWS for financial risk modeling. The company's HPC workloads run on Linux. Each HPC workflow runs on hundreds of Amazon EC2 Spot Instances, is short-lived and generates thousands of output files that are ultimately stored in persistent storage for analytics and long term future use. The company seeks a cloud storage solution that permits the copying of on-premises data to long-term persistent storage to make data available for processing by all EC2 instances. The solution should also be a high performance file system that is integrated with persistent storage to read and write datasets and output files. Which combination of AWS services meets these requirements ?

A. Amazon FSx for Lustre integrated with Amazon S3

B. Amazon FSx for Windows File Server integrated with Amazon S3

C. Amazon S3 Glacier integrated with Amazon Elastic Block Store (Amazon EBS)

D. Amazon S3 bucket with a VPC endpoint integrated with an Amazon Elastic Block Store (Amazon EBS) General Purpose SSD (gp2) volume.

(单选题) Q57. A company recently migrated a message processing system to AWS. The system receives messages into an ActiveMQ queue running on an Amazon EC2 instance. Messages are processed by a consumer application running on Amazon EC2. The consumer application processes the messages and writes results to a MySQL database running on Amazon EC2. The company wants this application to be highly available with low operational complexity. Which architecture offers the HIGHEST availability ?

A.Add a second ActiveMQ server to another Availability Zone.Add an additional consumer EC2 instance in another Availability Zone.Replicate the MySQL database to another Availability Zone.

B.Use Amazon MQ with active/standby brokers configured across two Availability Zones.Add an additional consumer EC2 instance in another Availability Zone.Replicate the MySQL database to another Availability Zone.

C.Use Amazon MQ with active/standby brokers configured across two Availability Zones.Add an additional consumer EC2 instance in another Availability Zone.Use Amazon RDS for MySQL with Multi-AZ enabled.

D.Use Amazon MQ with active/standby brokers configured across two Availability Zones.Add an Auto Scaling group for the consumer EC2 instances across two Availability Zones Use Amazon RDS for MySQL with Multi-AZ enabled.

(单选题) Q58.A company runs an application in the AWS Cloud and uses Amazon DynamoDB as the database.The company deploys Amazon EC2 instances to a private network to process data from the database.The company uses two NAT instances to provide connectivity to DynamoDB.The company wants to retire the NAT instances.A solutions architect must implement a solution that provides connectivity to DynamoDB and that does not require ongoing management.What is the Most cost-effective solution that meets these requirements ?

A.Create a gateway VPC endpoint to provide connectivity to DynamoDB

B.Configure a managed NAT gateway to provide connectivity to DynamoDB

C.Establish an AWS Direct Connect connection between the private network and DynamoDB

D.Deploy an AWS PrivateLink endpoint service between the private network and DynamoDB

(单选题) Q59.A company runs a web-based portal that provides users with global breaking news, local alerts, and weather updates.The portal delivers each user a personalized view by using a mixture of static and dynamic content.Content is served over HTTPS through an API server running on an Amazon EC2 instance behind an Application Load Balancer (ALB) .The company wants the portal to provide this

content to its users across the world as quickly as possible. How should a solutions architect design the application to ensure the LEAST amount of latency for all users?

- A. Deploy the application stack in a single AWS Region. Use Amazon CloudFront to serve all static and dynamic content by specifying the ALB as an origin.
- B. Deploy the application stack in two AWS Regions. Use an Amazon Route 53 latency routing policy to serve all content from the ALB in the closest Region.
- C. Deploy the application stack in a single AWS Region. Use Amazon CloudFront to serve the static content. Serve the dynamic content directly from the ALB.
- D. Deploy the application stack in two AWS Regions. Use an Amazon Route 53 geolocation routing policy to serve all content from the ALB in the closest Region.

(单选题) Q60. A company is running a batch application on Amazon EC2 instances. The application consists of a backend with multiple Amazon RDS databases. The application is causing a high number of reads on the databases. A solutions architect must reduce the number of database reads while ensuring high availability. What should the solutions architect do to meet this requirement?

- A. Add Amazon RDS read replicas.
- B. Use Amazon ElastiCache for Redis.
- C. Use Amazon Route 53 DNS caching.
- D. Use Amazon ElastiCache for Memcached.

(多选题) Q61. A company has a Microsoft .NET application that runs on an on-premises Windows Server. The application stores data by using an Oracle Database Standard Edition server. The company is planning a migration to AWS and wants to minimize development changes while moving the application. The AWS application environment should be highly available. Which combination of actions should the company take to meet these requirements? (Select TWO)

- A. Refactor the application as serverless with AWS Lambda functions running .NET Core.
- B. Rehost the application in AWS Elastic Beanstalk with the .NET platform in a Multi-AZ deployment.

C.Replatform the application to run on Amazon EC2 with the Amazon Linux Amazon Machine Image (AMI) .

D.Use AWS Database Migration Service (AWS DMS) to migrate from the Oracle database to Amazon DynamoDB in a Multi-AZ deployment

E.Use AWS Database Migration Service (AWS DMS) to migrate from the Oracle database to Oracle on Amazon RDS in a Multi-AZ deployment

(单选题) Q62.A company hosts its multi-tier applications on AWS.For compliance, governance, auditing, and security, the company must track configuration changes on its AWS resources and record a history of API calls made o these resources.What should a solutions architect do to meet these requirements ?

A.Use AWS CloudTrail to track configuration changes and AWS Config to record API calls

B.Use AWS Config to track configuration changes and AWS CloudTrail to record API calls

C.Use AWS Config to track configuration changes and Amazon CloudWatch to record API calls

D.Use AWS CloudTrail to track configuration changes and Amazon CloudWatch to record API calls

(多选题) Q63.A company recently launched Linux-based application instances on Amazon EC2 in a private subnet and launched a Linux-based bastion host on an Amazon EC2 instance in a public subnet of an VPC.A solution architect needs to connect from the on-premises network, through the company's internet connection, to the bastion host, and to the application servers.The solution architect must make sure that the security groups of all the EC2 instances will allow that access.Which combination of steps should the solutions architect take to meet these requirements ? (select TWO)

A.Replace the current security group of the bastion host with one that only allows inbound access from the application instances.

B.Replace the current security group of the bastion host with one that only allows inbound access from the internal IP range for the company.

- C.Replace the current security group of the bastion host with one that only allows inbound access from the external IP range for the company
- D.Replace the current security group of the application instances with one that allows inbound SSH access from only the private IP address of the bastion host.
- E.Replace the current security group of the application instances with one that allows inbound SSH access from only the public IP address of the bastion host.

(单选题) Q64.A company observes an increase in Amazon EC2 costs in its most recent bill.The billing team notices unwanted vertical scaling of instance types for a couple of EC2 instances.A solutions architect needs to create a graph comparing the last 2 months of EC2 costs and perform an in-depth analysis to identify the root cause of the vertical scaling.How should the solutions architect generate the information with the LEAST operational overhead ?

- A.Use AWS Budgets to create a budget report and compare costs based on instance types.
- B.Use Cost Explorer's granular filtering feature to perform an in-depth analysis of EC2 costs based on instance types.
- C.Use graphs from the AWS Billing and Cost Management dashboard to compare EC2 costs based on instance types for the last 2 months.
- D.Use AWS Cost and Usage Report to create a report and send it to an Amazon S3 bucket.Use Amazon QuickSight Amazon S3 as a source to generate an interactive graph based on instance types.

(多选题) Q65.A company is running a publicly accessible serverless application that uses Amazon API Gateway and AWS Lambda.The application's traffic recently spiked due to fraudulent requests from botnets.Which steps should a solutions architect take to block requests from unauthorized users ? (Select Two.)

- A.Create a usage plan with an API key that is shared with genuine users only.
- B.Integrate logic within the Lambda function to ignore the requests from fraudulent addresses.
- C.Implement an AWS WAF rule to target malicious requests and trigger actions to filter them out.

D.Convert the existing public API to a private API.Update the DNS records to redirect users to the new API endpoint.

E.Create an IAM role for each user attempting to access the APLA user will assume the role when making the API call.

(单选题) Q66.A gaming company is designing a highly available architecture.The application runs on a modified Linux kernel and support only UDP-based traffic.The company needs the front-end tier to provide the best possible user experience.The tier must have low latency, route traffic to the nearest edge location, and possible static IP addresses for entry into the application endpoints.What should a solution architect do to meet these requirements ?

A.Configure Amazon Route 53 to forward requests to an Application Load Balancer.Use AWS Lambda for the application in AWS Application Auto Scaling.

B.Configure Amazon CloudFront to forward requests to a network Load Balancer.Use AWS Lambda for the application in a Aws Application Auto Scaling group

C.Configure AWS Global Accelerator to forward requests to a Network Load Balancer.Use Amazon EC2 instances for the application in an EC2 Auto Scaling group.

D.Configure Amazon API Gateway to forward requests to an Application Load Balancer.Use Amazon EC2 instances for the application in an EC2 Auto Scaling group.

(单选题) Q67.A company has an application that provides marketing services to stores.The services are based on previous purchases by store customers.The stores upload transaction data to the company through SFTP, and the data is processed and analyzed to generate new marketing offers.Some of the files can exceed 200 GB in size.Recently, the company discovered that some of the stores have uploaded files that contain personally identifiable information (PII) that should not have been included.The company wants administrators to be alerted if PII is shared again.

The company also wants to automate remediation.What should a solutions architect do to meet these requirements with the LEAST development effort ?

A.Use an Amazon S3 bucket as a secure transfer point.Use Amazon Inspector to scan the objects in the bucket If objects contain PII, trigger an S3 Lifecycle policy to remove the objects that contain PII.

B. Use an Amazon S3 bucket as a secure transfer point. Use Amazon Macie to scan the objects in the bucket. If objects contain PII, use Amazon Simple Notification Service (Amazon SNS) to trigger a notification to the administrators to remove the objects that contain PII.

C. Implement custom scanning algorithms in an AWS Lambda function. Trigger the function when objects are loaded into the bucket. If objects contain PII, use Amazon Simple Notification Service (Amazon SNS) to trigger a notification to the administrators to remove the objects that contain PII.

D. Implement custom scanning algorithms in an AWS Lambda function. Trigger the function when objects are loaded into the bucket. If objects contain PII, use Amazon Simple Email Service (Amazon SES) to Trigger a notification to the administrators and trigger an S3 Lifecycle policy to remove the objects that contain PII.

(单选题) Q68. An entertainment company is using Amazon DynamoDB to store media metadata. The application is read intensive and experiencing delays. The company does not have staff to handle additional operational overhead and needs to improve the performance efficiency of DynamoDB without reconfiguring the application. What should a solutions architect recommend to meet this requirement?

A. Use Amazon ElastiCache for Redis

B. Use Amazon DynamoDB Accelerate (DAX)

C. Replicate data by using DynamoDB global tables

D. Use Amazon ElastiCache for Memcached with Auto Discovery enabled

(单选题) Q69. A company runs a photo processing application that needs to frequently upload and download pictures from Amazon S3 buckets that are located in the same AWS Region. A solutions architect has noticed an increased cost in data transfer fees and needs to implement a solution to reduce these costs. How can the solutions architect meet this requirement?

A. Deploy Amazon API Gateway into a public subnet and adjust the route table to route S3 calls through it.

B. Deploy a NAT gateway into a public subnet and attach an endpoint policy that allows access to the S3 buckets.

- C. Deploy the application into a public subnet and allow it to route through an internet gateway to access the S3 Buckets
- D. Deploy an S3 VPC gateway endpoint into the VPC and attach an endpoint policy that allows access to the S3 buckets.

(单选题) Q70. A company needs guaranteed Amazon EC2 capacity in three specific Availability Zones in a specific AWS Region for an upcoming event that will last 1 week. What should the company do to guarantee the EC2 capacity?

- A. Purchase Reserved Instances that specify the Region needed.
- B. Create an On-Demand Capacity Reservation that specifies the Region needed.
- C. Purchase Reserved Instances that specify the Region and three Availability Zones needed.
- D. Create an On-Demand Capacity Reservation that specifies the Region and three Availability Zones needed.

(单选题) Q71. A company sells ringtones created from clips of popular songs. The files containing the ringtones are stored in Amazon S3 Standard and are at least 123 KB in size. The company has millions of files but downloads are infrequent for ringtones older than 90 days. The company needs to save money on storage while keeping the most accessed files readily available for its users. Which action should the company take to meet these requirements MOST cost-effectively?

- A. Configure S3 Standard-Infrequent Access (S3 Standard-IA) storage for the initial storage tier of the objects
- B. Move the files to S3 Intelligent-Tiering and configure it to move objects to a less expensive storage tier after 90 days
- C. Configure S3 inventory to manage objects and move them to S3 Standard-Infrequent Access (S3 Standard-IA) after 90 days
- D. Implement an S3 Lifecycle policy that moves the objects from S3 Standard to S3 Standard-Infrequent Access (S3 Standard-IA) after 90 days

(多选题) Q72. A company is designing a cloud communications platform that is driven by APIs. The application is hosted on Amazon EC2 instances behind a Network

Load Balancer (NLB) .The company uses Amazon API Gateway to provide external users with access to the application through APIs.The company wants to protect the platform against web exploits like SQL Injection and also wants to detect and mitigate large, sophisticated DDoS attacks.Which combination of solutions provides the MOST protection ? (Select TWO.)

- A.Use AWS WAF to protect the NLB
- B.Use AWS Shield Advanced with the NLB
- C.Use AWS WAF to protect Amazon API Gateway
- D.Use Amazon GuardDuty with AWS Shield Standard
- E.Use AWS Shield Standard with Amazon API Gateway

(单选题) Q73.A company has an AWS account used for software engineering.The AWS account has access to the company's on-premises data center through a pair of AWS Direct Connect connections All non-VPC traffic routes to the virtual private gateway.A development team recently created an AWS Lambda function through the console.The development team needs to allow the function to access a database that runs in a private subnet in the company's data center.Which solution will meet these requirements ?

- A.Configure the Lambda function to run in the VPC with the appropriate security group.
- B.Set up a VPN connection from AWS to the data center.Route the traffic from the Lambda function through the VPN
- C.Update the route tables in the VPC to allow the Lambda function to access the on-premises data center through direct connect.
- D.Create an Elastic IP address.Configure the Lambda function to send traffic through the Elastic IP address without an elastic network interface.

(单选题) Q74.A development team runs monthly resource-intensive tests on its general purpose Amazon RDS (or MySQL DB instance with Performance insights enabled.The testing lasts for 48 hours once a month and is the only process that uses the database.The team wants to reduce the cost of running the tests without reducing the compute and memory attributes of the DB instance.Which solution meets these

requirements MOST cost-effectively ?

A.Stop the DB instance when tests are completed Restart the DB instance when required

B.Use an Auto Scaling policy with me DB instance to automatically scale when tests are completed

C.Create a snapshot when tests are completed Terminate the DB instance and restore the snapshot when required

D.Modify the DB instance to a low-capacity instance when tests are completed Modify the DB instance again when required

(单选题) Q75.A company hosts its web applications in the AWS Cloud.The company configures Elastic Load Balancers to use certificates that are imported into AWS Certificate Manager (ACM) .The company's security team must be notified 30 days before the expiration of each certificate.What should a solutions architect recommend to meet this requirement ?

A.Add a rule in ACM to publish a custom message to an Amazon Simple Notification Service (Amazon SNS) topic every day beginning 30 days before any certificate will expire

B.Create an AWS Config rule that checks for certificates that will expire within 30 days.Configure Amazon EventBridge (Amazon CloudWatch Events) to invoke a custom alert by way of Amazon Simple Notification Service (Amazon SNS) when AWS Config reports a noncompliant resource.

C.Use AWS Trusted Advisor to check for certificates that will expire within 30 days.Create an Amazon CloudWatch alarm that is based on Trusted Advisor metrics for check status changes.Configure the alarm to send a custom alert by way of Amazon Simple Notification Service (Amazon SNS) .

D.Create an Amazon EventBridge (Amazon CloudWatch Events) rule to detect any certificates that will expire within 30 days Configure the rule to invoke an AWS Lambda function.Configure the Lambda function to send a custom alert by way of Amazon Simple Notification Service (Amazon SNS) .

(单选题) Q76.A medical records company is hosting an application on Amazon

EC2 instances. The application processes customer data files that are stored on Amazon S3. The EC2 instances are hosted in public subnets. The EC2 instances access Amazon S3 over the internet, but they do not require any other network access. A new requirement mandates that the network traffic for file transfers take a private route and not be sent over the internet. Which change to the network architecture should a solutions architect recommend to meet this requirement?

- A. Create a NAT gateway. Configure the route table for the public subnets to send traffic to Amazon S3 through the NAT gateway.
- B. Configure the security group for the EC2 instances to restrict outbound traffic so that only traffic to the S3 prefix list is permitted.
- C. Move the EC2 instances to private subnets. Create a VPC endpoint for Amazon S3, and link the endpoint to the route table for the private subnets.
- D. Remove the internet gateway from the VPC. Set up an AWS Direct Connect connection, and route traffic to Amazon S3 over the Direct Connect connection.

(多选题) Q77. A social media company allows users to upload images to its website. The website runs on Amazon EC2 instances. During upload requests, the website resizes the images to a standard size and stores the resized images in Amazon S3. Users are experiencing slow upload requests to the website. The company needs to reduce coupling within the application and improve website performance. A solutions architect must design the most operationally efficient process for image uploads. Which combination of actions should the solutions architect take to meet these requirements? (Select TWO.)

- A. Configure the application to upload images to S3 Glacier.
- B. Configure the web server to upload the original images to Amazon S3.
- C. Configure the application to upload images directly from each user's browser to Amazon S3 through the use of a presigned URL.
- D. Configure S3 Event Notifications to invoke an AWS Lambda function when an image is uploaded. Use the function to resize the image.
- E. Create an Amazon EventBridge (Amazon CloudWatch Events) rule that invokes an AWS Lambda function on a schedule to resize uploaded images.

(单选题) Q78.A company wants to run its critical applications in containers to meet requirements for scalability and availability.The company prefers to focus on maintenance of the critical applications.The company does not want to be responsible for provisioning and managing the underlying infrastructure that runs the containerized workload.What should a solutions architect do to meet these requirements ?

- A.Use Amazon EC2instances and install Docker on the instances.
- B.Use Amazon Elastic Container Service (Amazon ECS) on AmazonEc2 worker nodes.
- C.Use Amazon Elastic Container Service (Amazon ECS) on AWS Fargate.
- D.Use Amazon EC2 instances from an Amazon Elastic Container Service (Amazon ECS) -optimized Amazon Machine Image (AMI) .

(多选题) Q79.A solutions architect must design a highly available infrastructure for a website.The website is powered by Windows web servers that run on Amazon EC2 instances.The solutions architect must implement a solution that can mitigate a large-scale DDoS attack that originates from thousands of IP addresses.Downtime is not acceptable for the website.Which actions should the solutions architect take to protect the website from such an attack ? (Select Two)

- A.Use AWS Shield Advanced to stop the DDoS attack.
- B.Configure Amazon GuardDuty to automatically block the attackers.
- C.Configure the website to use Amazon CloudFront for both static and dynamic content.
- D.Use an AWS Lambda function to automatically add attacker IP addresses to VPC network ACLs.
- E.Use EC2Spot Instances in an Auto Scaling group with a target tracking scaling policy that is set to 80%CPU utilization.

(单选题) Q80.A company recently signed a contract with an AWS Managed Service Provider (MSP) Partner for help with an application migration initiative.A solutions architect needs to share an Amazon Machine Image (AMI) from an existing AWS account with the MSP Partner's AWS account The AMI is backed by Amazon Elastic Block Store (Amazon EBS) and uses a customer managed customer master

key (CMK) to encrypt EBS volume snapshots.What is the MOST secure way for the solutions architect to share the AMI with the MSP Partner's AWS account ?

A.Make the encrypted AMI and snapshots publicly available Modify the CMKs key policy to allow the MSP Partner's AWS account to use the key

B.Modify the launch Permission property of the AMI Share the AMI with the MSP Partner's AWS account only.Modify the CMKs key policy to allow the MSP Partner's AWS account to use the key

C.Modify the launch Permission property of the AMI.Share the AMI with the MSP Partner's AWS account only.Modify the CMKs key policy to trust a new CMK that is owned by the MSP Partner for encryption.

D.Export the AMI from the source account to an Amazon S3 bucket in the MSP Partner's AWS account.Encrypt the S3 bucket with a CMK that is owned by the MSP Partner.Copy and launch the AMI in the MSP Partner's AWS account.

(单选题) Q81.A company runs an on-premises application that is powered by a MySQL database.The company is migrating the application to AWS to increase the application's elasticity and availability The current architecture shows heavy read activity on the database during times of normal operation.Every 4 hours, the company's development team pulls a full export of the production database to populate a database in the staging environment.During this period, users experience unacceptable application latency.The development team is unable to use the staging environment until the procedure completes.A solutions architect must recommend replacement architecture that alleviates the application latency issue.The replacement architecture also must give the development team the ability to continue using the staging environment without delay.Which solution meets these requirements ?

A.Use Amazon Aurora MySQL with Multi-AZ Aurora Replicas for production Populate the staging database by implementing a backup and restore process that uses the mysqldump utility.

B.Use Amazon Aurora MySQL with Multi-AZ Aurora Replicas for production Use database cloning to create the staging database on-demand.

C.Use Amazon RDS for MySQL with a Multi-AZ deployment and read replicas for

production. Use the standby instance for the staging database.

D. Use Amazon RDS for MySQL with a Multi-AZ deployment and read replicas for production. Populate the staging database by implementing a backup and restore process that uses the mysqldump utility.

(单选题) Q82. A company is implementing a shared storage solution for a media application that is hosted in the AWS Cloud. The company needs the ability to use SMB clients to access data.

The solution must be fully managed. Which AWS solution meets these requirements? A. Create an AWS Storage Gateway volume gateway. Create a file share that uses the required client protocol. Connect the application server to the file share.

B. Create an AWS Storage Gateway tape gateway. Configure tapes to use Amazon S3. Connect the application server to the tape gateway.

C. Create an Amazon EC2 Windows instance. Install and configure a Windows file share role on the instance. Connect the application server to the file share.

D. Create an Amazon FSx for Windows File Server file system. Attach the file system to the origin server. Connect the application server to the file system.

(单选题) Q83. A company needs the ability to analyze the log files of its proprietary application. The logs are stored in JSON format in an Amazon S3 bucket. Queries will be simple and will run on-demand. A solutions architect needs to perform the analysis with minimal changes to the existing architecture. What should the solutions architect do to meet these requirements with the LEAST amount of operational overhead?

A. Use Amazon Redshift to load all the content into one place and run the SQL queries as needed.

B. Use Amazon CloudWatch Logs to store the logs. Run SQL queries as needed from the Amazon CloudWatch console.

C. Use Amazon Athena directly with Amazon S3 to run the queries as needed.

D. Use AWS Glue to catalog the logs. Use a transient Apache Spark cluster on Amazon EMR to run the SQL queries as needed.

(单选题) Q84.A company is concerned about the security of its public web application due to recent web attacks.The application uses an Application Load Balancer (ALB) .A solutions architect must reduce the risk of DDoS attacks against the application.What should the solutions architect do to meet this requirement ?

- A.Add an Amazon Inspector agent to the ALB.
- B.Configure Amazon Macie to prevent attacks.
- C.Enable AWS Shield Advanced to prevent attacks.
- D.Configure Amazon GuardDuty to monitor the ALB.

(单选题) Q85.A company hosts a static website on-premises and wants to migrate the website to AWS.The website should load as quickly as possible for users around the world.The company also wants the most cost-effective solution.What should a solutions architect do to accomplish this ?

- A.Copy the website content to an Amazon S3 bucket.Configure the bucket to serve static webpage content.Replicate the S3 bucket to multiple AWS Regions.
- B.Copy the website content to an Amazon S3 bucket.Configure the bucket to serve static webpage content.Configure Amazon CloudFront with the S3 bucket as the origin.
- C.Copy the website content to an Amazon EBS-backed.Amazon EC2 instance running Apache HTTP Server.Configure Amazon Route 53 geolocation routing policies to select the closest origin.
- D.Copy the website content to multiple Amazon EBS-backed.Amazon EC2 instances running Apache HTTP Server in multiple AWS Regions.Configure Amazon CloudFront geolocation routing policies to select the closest origin.

(单选题) Q86. A company has a customer relationship management (CRM) application that stores data in an Amazon RDS DB instance that runs Microsoft SQL Server. The company's IT staff has administrative access to the database. The database contains sensitive data. The company wants to ensure that the data is not accessible to the IT staff and that only authorized personnel can view the data.

What should a solutions architect do to secure the data?

- A. Use client-side encryption with an Amazon RDS managed key

- B. Use client-side encryption with an AWS Key Management Service (AWS KMS) customer managed key
- C. Use Amazon RDS encryption with an AWS Key Management Service (AWS KMS) default encryption key
- D. Use Amazon RDS encryption with an AWS Key Management Service (AWS KMS) customer managed key

(单选题) Q87.A solutions architect is designing an infrastructure to run a web application on Amazon EC2 instances.The company wants to ensure high availability and adjust the number of running instances based on demand to optimize costs.What should the solutions architect do to meet these requirements ?

- A.Configure an Application Load Balancer in front of an Auto Scaling group to deploy instances to multiple AWS Regions.
- B.Configure an Amazon CloudFront distribution in front of an Auto Scaling group to deploy instances to multiple AWS Regions.
- C.Configure an Application Load Balancer in front of an Auto Scaling group to deploy instances to multiple Availability Zones.
- D.Configure an Amazon API Gateway API in front of an Auto Scaling group to deploy instances to multiple Availability Zones.

(单选题) Q88.A global company currently gathers 100 GB of information daily from deployed sensors on smart devices with an average record size of 10 KB.A solutions architect must design a solution that allows real-time access to the data.What should the solutions architect do to meet these requirements ?

- A.Configure each sensor to upload data directly to an Amazon EC2 instance in each AWS Region.
- B.Configure the sensors to use Amazon Kinesis Data Streams for batching and data collection.
- C.Provision an Amazon Snowball appliance and install the Snowball client to optimize the transfer speed.
- D.Deploy the website on Amazon CloudFront and configure each device to upload data to the closest edge location.

(单选题) Q89.A company needs to meet a strict service level agreement (SLA) with its users for application availability.The primary application stack is hosted in the us-east-1 Region and a replica of the stack is in us-west-2 for disaster recovery.The company requires that no user traffic be served by the application stack in us-west-2 unless there is an issue with the application stack in us-east-1.What should a solutions architect do to meet the company's needs without manual intervention ?

- A.Configure a weighted routing policy in Amazon Route 53.Set the weight of us-east-1 to 100 and us-west-2 to 0.
- B.Configure a weighted routing policy in Amazon Route 53.Set the weight of us-west-2 to 50 and us-east-1to 50.
- C.Configure a failover routing policy in Amazon Route 53.Set us-east-1as the primary and us-west-2 as the secondary.
- D.Configure a failover routing policy in Amazon Route 53.Set us-west-2 as the primary and us-east-1 as the secondary.

(单选题) Q90.A company has an application running on a single server in one Availability Zone.A solutions architect needs to improve the availability of this application as it grows in importance.Which set of actions would accomplish this ?

- A.Deploy a second application instance in the same Availability Zone.Add the instances to Amazon Route 53in a round-robin configuration.
- B.Deploy a second application instance in a public subnet in the same Availability Zone.Attach an Elastic IP address to both instances for high availability.
- C.Redeploy the application instance in an Amazon EC2 Auto Scaling group with a minimum of 1 and a maximum of 1.Use AWS CloudFormation to automate the deployment.
- D.Deploy an Elastic Load Balancer and a second application instance in a second Availability Zone.Attach both instances to the Elastic Load Balancer and route all traffic through the load balancer.

(单选题) Q91.A company's HTTP application is behind a Network Load Balancer (NLB) .The NLB's target group is configured to use an Amazon EC2 Auto

Scaling group with multiple EC2 instances that run the web service. The company notices that the NLB is not detecting HTTP errors for the application. These errors require a manual restart of the EC2 instances that run the web service. The company needs to improve the application's availability without writing custom scripts or code. What should a solutions architect do to meet these requirements ?

A. Enable HTTP health checks on the NLB supplying the URL of the company's application

B. Add a cron job to the EC2 instances to check the local application's logs once each minute. If HTTP errors are detected, the application will restart

C. Replace the NLB with an Application Load Balancer. Enable HTTP health checks by supplying the URL of the company's application. Configure an Auto Scaling action to replace unhealthy instances

D. Create an Amazon CloudWatch alarm that monitors the UnhealthyHostCount metric for the NLB. Configure an Auto Scaling action to replace unhealthy instances when the alarm is in the ALARM state

(单选题) Q92. A company is implementing a shared storage solution for a gaming application that is hosted in an on-premises data center. The company needs the ability to use Lustre clients to access data. The solution must be fully managed. Which solution meets these requirements ?

A. Create an AWS Storage Gateway file gateway. Create a file share that uses the required client protocol. Connect the application server to the file share.

B. Create an Amazon EC2 Windows instance. Install and configure a Windows file share role on the instance. Connect the application server to the file share.

C. Create an Amazon Elastic File System (Amazon EFS) file system, and configure it to support Lustre. Attach the file system to the origin server. Connect the application server to the file system

D. Create an Amazon FSx for Lustre file system. Attach the file system to the origin server. Connect the application server to the file system.

(单选题) Q93. A company hosts more than 300 global websites and applications. The company requires a platform to analyze more than 30 TB of

clickstream data each day.What should a solutions architect do to transmit and process the clickstream data ? A.Design an AWS Data Pipeline to archive the data to an Amazon S3 bucket and run an Amazon EMR cluster with the data to generate analytics.

B.Create an Auto Scaling group of Amazon EC2instances to process the data and send it to an Amazon s3 data lake forAmazon Redshift to use for analysis.

C.Cache the data to Amazon CloudFront.Store the data in an Amazon S3 bucket.When an object is added to the s3 bucket run an AWS Lambda function to process the data for analysis

D.Collect the data from Amazon Kinesis Data Streams Use Amazon Kinesis Data Firehose to transmit the data to an Amazon S3 data lake.Load the data in Amazon Redshift for analysis

(单选题) Q94.A company has a production web application in which users upload documents through a web interface or a mobile app According to a new regulatory requirement, new documents cannot be modified or deleted after they are stored.What should a solutions architect do to meet this requirement ?

A.Store the uploaded documents in an Amazon S3 bucket with S3 Versioning and S3 Object Lock enabled

B.Store the uploaded documents in an Amazon S3 bucket Configure an s3 Lifecycle policy to archive the documentsperiodically.

C.Store the uploaded documents in an Amazon S3 bucket with S3 Versioning enabled.Configure an ACL to restrict all access to read-only.

D.Store the uploaded documents on an Amazon Elastic File System Amazon EFS volume Access the data by mounting the volume in read-only mode.

(多选题) Q95.An application development team is designing a microservice that will convert large images to smaller, compressed images.When a user uploads an image through the web interface, the microservice should store the image in an Amazon S3 bucket, process and compress the image with an AWS Lambda function, and store the image in its compressed form in a different S3 bucket.A solutions architect needs to design a solution that uses durable, stateless

components to process the images automatically. Which combination of actions will meet these requirements? (Select Two)

A. Create an Amazon Simple Queue Service (Amazon SQS) queue. Configure the S3 bucket to send a notification to the SQS queue when an image is uploaded to the S3 bucket.

B. Configure the Lambda function to use the Amazon Simple Queue Service (Amazon SQS) queue as the invocation source. When the SQS message is successfully processed, delete the message in the queue.

C. Configure the Lambda function to monitor the S3 bucket for new uploads. When an uploaded image is detected, write the file name to a text file in memory and use the text file to keep track of the images that were processed.

D. Launch an Amazon EC2 instance to monitor an Amazon Simple Queue Service (Amazon SQS) queue. When items are added to the queue, log the file name in a text file on the EC2 instance and invoke the Lambda function.

D. Configure an Amazon Event Bridge (Amazon CloudWatch Events) event to monitor the S3 bucket. When an image is uploaded, send an alert to an Amazon Simple Notification Service (Amazon SNS) topic with the application owner's email address for further processing.

(单选题) Q96. A company uses AWS Organizations to manage multiple AWS accounts for different departments. The management account has an Amazon S3 bucket that contains project reports. The company wants to limit access to this S3 bucket to only users of accounts within the organization in AWS Organizations. Which solution meets these requirements with the LEAST amount of operational overhead?

A. Add the `aws:PrincipalOrgID` global condition key with a reference to the organization ID to the S3 bucket policy.

B. Create an organizational unit (OU) for each department. Add the `aws:PrincipalOrgPaths` global condition key to the S3 bucket policy.

C. Use AWS CloudTrail to monitor the `CreateAccount`, `InviteAccountToOrganization`, `LeaveOrganization`, and `RemoveAccountFromOrganization` events. Update the S3 bucket policy accordingly.

D.Tag each user that needs access to the S3 bucket.Add the awsPrincipalTag global condition key to the S3 bucket policy.

(单选题) Q97.A company hosts an application on AWS Lambda functions that are invoked by an Amazon API Gateway API.The Lambda functions save customer data to an Amazon Aurora MySQL database Whenever the company upgrades the database, the Lambda functions fail to establish database connections until the upgrade is complete.The result is that customer data is not recorded for some of the event.A solutions architect needs to design a solution that stores customer data that is created during database upgrades Which solution will meet these requirements ?

A.Provision an Amazon RDS proxy to sit between the Lambda functions and the database.Configure the Lambda functions to connect to the RDS proxy.

B.Increase the run time of the Lambda functions to the maximum.Create a retry mechanism in the code that stores the customer data in the database.

C.Persist the customer data to Lambda local storage.Configure new Lambda functions to scan the local storage to save the customer data to the database.

D.Store the customer data in an Amazon Simple Queue Service (Amazon SQS) FIFO queue Create a new Lambda function that polls the queue and stores the customer data in the database.

(单选题) Q98.A company collects data for temperature, humidity, and atmospheric pressure in cities across multiple continents.The average volume of data that the company collects from each site daily is 500 GB Each site has a high-speed internet connection.The company wants to aggregate the data from all these global sites as quickly as possible in a single Amazon S3 bucket.The solution must minimize operational complexity.Which solution meets these requirements ?

A.Turn on S3 Transfer Acceleration on the destination S3 bucket Use multipart uploads to directly upload site data to the destination S3 bucket

B.Upload the data from each site to an S3 bucket in the closest Region.Use S3 Cross-Region Replication to copy objects to the destination S3 bucket.Then remove the data from the origin S3 bucket

C.Schedule AWS Snowball Edge Storage Optimized device jobs daily to transfer data

from each site to the closest Region. Use S3 Cross-Region Replication to copy objects to the destination S3 bucket.

D. Upload the data from each site to an Amazon EC2 instance in the closest Region. Store the data in an Amazon Elastic Block Store (Amazon EBS) volume. At regular intervals, take an EBS snapshot and copy it to the Region that contains the destination S3 bucket. Restore the EBS volume in that Region.

(单选题) Q99. A company has a Windows-based application that must be migrated to AWS. The application requires the use of a shared Windows file system attached to multiple Amazon EC2 Windows instances that are deployed across multiple Availability Zones. What should a solutions architect do to meet this requirement?

A. Configure AWS Storage Gateway in volume gateway mode. Mount the volume to each Windows instance

B. Configure Amazon FSx for Windows File Server. Mount the Amazon FSx file system to each Windows instance.

C. Configure a file system by using Amazon Elastic File System (Amazon EFS). Mount the EFS file system to each Windows instance.

D. Configure an Amazon Elastic Block Store (Amazon EBS) volume with the required size. Attach each EC2 instance to the volume. Mount the file system within the volume to each Windows instance

(单选题) Q100. A solutions architect is creating a new Amazon CloudFront distribution for an application. Some of the information submitted by users is sensitive. The application uses HTTPS but needs another layer of security. The sensitive information should be protected throughout the entire application stack, and access to the information should be restricted to certain applications. Which action should the solutions architect take?

A. Configure a CloudFront signed URL

B. Configure a CloudFront signed cookie.

C. Configure a CloudFront field-level encryption profile.

D. Configure CloudFront and set the Origin Protocol Policy setting to HTTPS Only for the viewer Protocol Policy

(单选题) Q101.A company has thousands of edge devices that collectively generate 1 TB of status alerts each day Each alert is approximately 2 KB in size.A solutions architect needs to implement a solution to ingest and store the alerts for future analysis.The company wants a highly available solution.However, the company needs to minimize costs and does not want to manage additional infrastructure.Additionally, the company wants to keep 14 days of data available for immediate analysis and archive any data older than 14 days.What is the MOST operationally efficient solution that meets these requirements ?

A.Create an Amazon Kinesis Data Firehose delivery stream to ingest the alerts.Configure the Kinesis Data Firehose stream to deliver the alerts to an Amazon S3 bucket Set up an s3 Lifecycle configuration to transition data to Amazon S3 Glacier after 14 days

B.Launch Amazon EC2 instances across two Availability Zones and place them behind an Elastic Load Balancer to ingest the alerts.Create a script on the EC2 instances that will store the alerts in an Amazon S3 bucket.Set up an S3 Lifecycle configuration to transition data to Amazon S3 Glacier after 14 days

C.Create an Amazon Kinesis Data Firehose delivery stream to ingest the alerts.Configure the Kinesis Data Firehose stream to deliver the alerts to an Amazon Elasticsearch Service (Amazon ES) cluster.Set up the Amazon ES cluster to take manual snapshots every day and delete data from the cluster that is older than 14 days.

D.Create an Amazon Simple Queue Service (Amazon SQS) standard queue to ingest the alerts, and set the message retention period to 14 days.Configure consumers to poll the SQS queue, check the age of the message, and analyze the message data as needed If the message is 14 days old, the consumer should copy the message to an Amazon S3 bucket and delete the message from the SQS queue

(单选题) Q102.A company is running an online transaction processing (OLTP) workload on AWS.This workload uses an unencrypted Amazon RDS DB instance in a Multi-AZ deployment.Daily database snapshots are taken from this instance.What should a solutions architect do to ensure the database and snapshots are always

encrypted moving forward ?

- A.Encrypt a copy of the latest DB snapshot.Replace existing DB instance by restoring the encrypted snapshot
- B.Create a new encrypted Amazon Elastic Block Store Amazon EBS) volume and copy the snapshots to it Enable encryption on the DB instance
- C.Copy the snapshots and enable encryption using AWS Key Management Service (AWS KMS) .Restore encrypted snapshot to an existing DB instance
- D.Copy the snapshots to an Amazon S3 bucket that is encrypted using server-side encryption with AwS Key Management Service (AwS KMS) managed keys (SSE-KMS)

(单选题) Q103.A company runs multiple applications in containers on Amazon Elastic Container Service (Amazon ECS instances.The applications run in an ECS cluster.Users access the applications through an Application Load Balancer (ALB) . The users report that the applications are slow every morning from 9-00 until 9 : 30 After that time period, no issues occur.What should a solutions architect recommend to resolve this issue ?

- A.Configure an Auto Scaling group with a step scaling policy.Attach the ECS instances to the Auto Scaling group
- B.Configure an Auto Scaling group with a scheduled scaling policy.Attach the ECS instances to the Auto Scaling group
- C.Create a second ECS cluster with two ECS instances Deploy some of the applications to the second ECS cluster
- D.Add a second ALB in front of ECS instances.Configure the second ALB to route traffic to some of the applications on the ECS instances.

(单选题) Q104.A company runs a containerized application on a Kubernetes cluster in an on-premises data center.The company is using a MongoDB database for data storage.The company wants to migrate some of these environments to AWS, but no code changes or deployment method changes are possible at this time.The company needs a solution that minimizes operational overhead.Which solution meets these requirements ?

- A.Use Amazon Elastic Container Service (Amazon ECS) with Amazon EC2 worker nodes for compute and MongoDB on EC2 for data storage
- B.Use Amazon Elastic Container Service (Amazon ECS) with AWS Fargate for compute and Amazon DynamoDB for data storage.
- C.Use Amazon Elastic Kubemnetes Service (Amazon EKS) with Amazon EC2 worker nodes for compute and Amazon DynamoDB for data storage
- D.Use Amazon Elastic Kubemnetes Service (Amazon EKS) with AWS Fargate for compute and Amazon DocumentDB (with MongoDB compatibility) for data storage

(单选题) Q105.A survey company has gathered data for several years from areas in the United States.The company hosts the data in an Amazon S3 bucket that is 3 TB in size and growing.The company has started to share the data with a European marketing firm that has S3 buckets.The company wants to ensure that its data transfer costs remain as lo was possible.Which solution will meet these requirements ?

- A.Configure the Requester Pays feature on the company's S3 bucket.
- B.Configure S3 Cross-Region Replication from the company's S3 bucket to one of the marketing firm's S3 buckets.
- C.Configure cross-account access for the marketing firm so that the marketing firm has access to the company's S3 bucket.
- D.Configure the company's S3 bucket to use S3intelligent-Tiering.Sync the S3 bucket to one of the marketing firm's S3 buckets.

(单选题) Q106.A company runs its infrastructure on AWS and has a registered base of 700000 users for its document management application.The company intends to create a product that converts large.pdf files to.jpg image files.The.pdf files average 5 MB in size.The company needs to store the original files and the converted files.A solutions architect must design a scalable solution to accommodate demand that will grow rapidly over time.Which solution meets these requirements MOST cost-effectively ?

- A.Save the.pdf files to Amazon S3.Configure an S3 PUT event to invoke an AWS Lambda function to convert the files to.jpg format and store them back in Amazon

S3.

B.Save the.pdf files to Amazon DynamoDB.Use the DynamoDB Streams feature to invoke an AWS Lambda function to convert the files to jpg format and store them back in DynamoDB.

C.Upload the pdf files to an AWS Elastic Beanstalk application that includes Amazon EC2 instances, Amazon Elastic Block Store (Amazon EBS) storage, and an Auto Scaling group.Use a program in the EC2 instances to convert the files to jpg format.Save the.pdf files and the.jpg files in the EBS store.

D.Upload the.pdf files to an AWS Elastic Beanstalk application that includes Amazon EC2 instances, Amazon Elastic File System (Amazon EFS) storage, and an Auto Scaling group.Use a program in the EC2 instances to convert the file to ipg format.Save the pdf files and the jpg files in the EBS store

(单选题) Q107.A company runs a stateless web application in production on a group of Amazon EC2 On-Demand Instances behind an Application Load Balancer.The application experiences heavy usage during an 8-hour period each business day.Application usage is moderate and steady overnight.Application usage is low during weekends.The company wants to minimize its EC2 costs without affecting the availability of the application.Which solution will meet these requirements ?

A.Use Spot Instances for the entire workload.

B.Use Reserved Instances for the baseline level of usage.Use Spot Instances for any additional capacity that the application needs.

C.Use On-Demand Instances for the baseline level of usage.Use Spot Instances for any additional capacity that the application needs.

D.Use Dedicated Instances for the baseline level of usage.Use On-Demand Instances for any additional capacity that the application needs.

(单选题) Q108.A company that has a global user base is launching a new service.The service is highly anticipated because of a long marketing campaign.The company has built the service on Amazon EC2 instances that run behind an Application Load Balancer (ALB) .The company uses an Amazon RDS for MySQL

DB instance and Amazon S3 to store data.Previously, the company experienced a large-scale cyberattack that originated from a specific country where the company has no user base.Before the launch of the service, the company wants to block all traffic that originates from that specific country:Which solution will meet this requirement ?

- A.Create an EC2 security group to deny any traffic that comes from the specific country's CIDR ranges.Attach the security group to the ALB
- B.Create an RDS security group to deny any traffic that comes from the specific country's CIDR ranges Attach the security group to the DB instance
- C.Create a Deny statement on the S3 bucket policy.Use a condition that specifies which country code to deny access to
- D.Create an AWS WAF web ACL that has a rule to block the traffic based on a geographical match condition.Associate the web ACL with the ALB

(单选题) Q109.A company runs an application on a large fleet of Amazon EC2 instances.The application reads and write entries into an Amazon DynamoDB table.The size of the DynamoDB table continuously grows, but the application needs only data from the last 30 days.The company needs a solution that minimizes cost and development effort.Which solution meets these requirements ?

- A.Use an AWS CloudFormation template to deploy the complete solution.Redeploy the CloudFormation stack every 30 days, and delete the original stack.
- B.Use an EC2 instance that runs a monitoring application from AWS Marketplace.Configure the monitoring application to use Amazon DynamoDB Streams to store the timestamp when a new item is created in the table.Use a script that runs on the EC2 instance to delete items that have a timestamp that is older than 30 days
- C.Configure Amazon DynamoDB Streams to invoke an AWS Lambda function when a new item is created in the table.Configure the Lambda function to delete items in the table that are older than 30 days
- D.Extend the application to add an attribute that has a value of the current timestamp plus 30 days to each new item that is created in the table.Configure DynamoDB to use the attribute as the TTL attribute.

(多选题) Q110.A solutions architect needs to help a company optimize the cost of running an application on AWS.The application will use Amazon EC2 instances.AWS Fargate and AWS Lambda for compute within the architecture.The EC2instances will run the data ingestion layer of the application.EC2 usage will be sporadic and unpredictable.Workloads that run on EC2 instances can be interrupted at any time.The application front end will run on Fargate, and Lambda will serve the API layer.The front-end utilization and API layer utilization will be predictable over the course of the next year.Which combination of purchasing options will provide the MOST cost-effective solution for hosting this application ? (Select TWO)

- A.Use Spot Instances for the data ingestion layer
- B.Use On-Demand Instances for the data ingestion layer
- C.Purchase a 1-year Compute Savings Plan for the front end and API layer
- D.Purchase 1-year All Upfront Reserved Instances for the data ingestion layer
- E.Purchase a 1-year EC2Instance Savings Plan for the front end and API layer

(单选题) Q111.A company has a legacy data processing application that runs on Amazon EC2instances Data is processed sequentially, but the order of results does not matter.The application uses a monolithic architecture.

The only way that the company can scale the application to meet increased demand is to increase the size of the instancesThe company's developers have decided to rewrite the application to use a microservices architecture on Amazon Elastic Container Service (Amazon ECS) .What should a solutions architect recommend for communication between the microservices ?

- A.Create an Amazon Simple Queue Service (Amazon SQS) queue.Add code to the data producers, and send data to the queue.Add code to the data consumers to process data from the queue
- B.Create an Amazon Simple Notification Service (Amazon SNS) topic.Add code to the data producers, and publish notifications to the topic.Add code to the data consumers to subscribe to the topic
- C.Create an AWS Lambda function to pass messages.Add code to the data producers to call the Lambda function with a data object.Add code to the data consumers to

receive a data object that is passed from the Lambda function

D.Create an Amazon DynamoDB table.Enable DynamoDB Streams.Add code to the data producers to insert data into the table.Add code to the data consumers to use the DynamoDB Streams API to detect new table entries and retrieve the data

(单选题) Q112.A company has applications that run on Amazon EC2 instances in a VPC.One of the applications needs to call the Amazon S3 API to store and read objects.According to the company's security regulations, no traffic from the applications is allowed to travel across the internet.Which solution will meet these requirements ?

- A.Configure an S3 gateway endpoint
- B.Create an S3 bucket in a private subnet
- C.Create an S3 bucket in the same AWS Region as the EC2 instances
- D.Configure a NAT gateway in the same subnet as the EC2 instances

(单选题) Q113.A company has a data ingestion workflow that includes the following components:--An Amazon Simple Notification Service (Amazon SNS) topic that receives notifications about new data deliveries.-An AWS Lambda function that processes and stores the data.The ingestion workflow occasionally fails because of network connectivity issues.When failure occurs, the corresponding data is not ingested unless the company manually reruns the job What should a solutions architect do to ensure that all notifications are eventually processed ?

- A.Configure the Lambda function for deployment across multiple Availability Zones
- B.Modify the Lambda function's configuration to increase the CPU and memory allocations for the function
- C.Configure the SNS topic's retry strategy to increase both the number of retries and the wait time between retries
- D.Configure an Amazon Simple Queue Service (Amazon SQS) queue as the on-failure destination.Modify the Lambda function to process messages in the queue

(多选题) Q114.A company wants to migrate an on-premises data center to AWS.The data center hosts an SFTP server that stores its data on an NFS-based file

system.The server holds 200 GB of data that needs to be transferred.The server must be hosted on an Amazon EC2 instance that uses an Amazon Elastic File System (Amazon EFS) file system.Which combination of steps should a solutions architect take to automate this task ? (Select TWO.)

- A.Launch the EC2 instance into the same Availability Zone as the EFS file system.
- B.Install an AWS DataSync agent in the on-premises data center
- C.Create a secondary Amazon Elastic Block Store (Amazon EBS) volume on the EC2 instance for the data.
- D.Manually use an operating system copy command to push the data to the EC2 instance
- E.Use AWS DataSync to create a suitable location configuration for the on-premises SFTP server.

(单选题) Q115.A company has several web servers that need to frequently access a common Amazon RDS MySQL Multi-AZ DB instance.The company wants a secure method for the web servers to connect to the database while meeting a security requirement to rotate user credentials frequently.Which solution meets these requirements ?

- A.Store the database user credentials in AWS Secrets Manager.Grant the necessary IAM permissions to allow the web servers to access AWS Secrets Manager
- B.Store the database user credentials in AWS Systems Manager OpsCenter.Grant the necessary IAM permissions to allow the web servers to access OpsCenter
- C.Store the database user credentials in a secure Amazon S3 bucket Grant the necessary IAM permissions to allow the web servers to retrieve credentials and access the database
- D.Store the database user credentials in files encrypted with AWS Key Management Service (AWS KMS) on the web server file system.The web server should be able to decrypt the files and access the database

(单选题) Q116.A company is hosting a static website on Amazon S3 and is using Amazon Route 53 for DNS.The website is experiencing increased demand from around the world.The company must decrease latency for users who access the

website.Which solution meets these requirements MOST cost-effectively ?

A.Replicate the S3 bucket that contains the website to all AWS Regions.Add Route 53 geolocation routing entries

B.Provision accelerators in AWS Global Accelerator.Associate the supplied IP addresses with the S3bucket.Edit the Route 53entries to point to the IP addresses of the accelerators

C.Add an Amazon CloudFront distribution in front of the S3 bucket.Edit the Route 53 entries to point to the CloudFront distribution

D.Enable S3 Transfer Acceleration on the bucket Edit the Route 53 entries to point to the new endpoint

(单选题) Q117.A company is developing a file-sharing application that will use an Amazon S3 bucket for storage.The company wants to serve all the files through an Amazon CloudFront distribution.The company does not want the files to be accessible through direct navigation to the S3 URL.What should a solutions architect do to meet these requirements ?

A.Write individual policies for each S3 bucket to grant read permission for only CloudFront access

B.Create an IAM user Grant the user read permission to objects in the S3 bucket Assign the user to CloudFront

C.Write an S3 bucket policy that assigns the CloudFront distribution ID as the Principal and assigns the target S3 bucket as the Amazon Resource Name (ARN)

D.Create an origin access identity (OAI) Assign the OAI to the CloudFront distribution.Configure the S3 bucket permissions so that only the OAI has read permission.

(单选题) Q118.A company wants to run applications in containers in the AWS Cloud.These applications are stateless and can tolerate disruptions within the underlying infrastructure.The company needs a solution that minimizes cost and operational overhead.What should a solutions architect do to meet these requirements ?

A.Use Spot Instances in an Amazon EC2 Auto Scaling group to run the application

containers

B.Use Spot Instances in an Amazon Elastic Kubernetes Service (Amazon EKS) managed node group

C.Use On-Demand Instances in an Amazon EC2Auto Scaling group to run the application containers

D.Use On-Demand Instances in an Amazon Elastic Kubernetes Service (Amazon EKS) managed node group

(单选题) Q119.A company wants to reduce the cost of its existing three-tier web architecture.The web, application, and database servers are running on Amazon EC2 instances for the development, test, and production environments.The EC2 instances average 30%CPU utilization during peak hours and 10%CPU utilization during non-peak hours.The production EC2instances run 24 hours a day.The development and test EC2 instances run for at least 8 hours each day.The company plans to implement automation to stop the development and test EC2 instances when they are not in use.Which EC2 instance purchasing solution will meet the companys requirements MOST cost-effectively ?

A.Use Spot Instances for the production EC2 instances.Use Reserved Instances for the development and test EC2 instances.

B.Use Reserved Instances for the production EC2 instances Use On-Demand Instances for the development and test EC2 instances.

C.Use Spot blocks for the production EC2 instances Use Reserved Instances for the development and testEc2instances

D.Use On-Demand Instances for the production EC2 instances.Use Spot blocks for the development and test EC2 instances.

(单选题) Q120.A company needs to retain application log fles for a critical application for 10years.The application team regularly accesses logs from the past month for troubleshooting.but logs older than 1 month are rarely accessed.The application generates more than 10 TB of logs per month.Which storage option meets these requirements MOST cost-effectively ?

A.Store the logs in Amazon S3.Use AWS Backup to move logs more than 1month old

to S3 Glacier Deep Archive

B.Store the logs in Amazon S3.Use S3Lifecycle policies to move logs more than 1month old to S3 Glacier Deep Archive

C.Store the logs in Amazon CloudWatch Logs.Use AWS Backup to move logs more than 1month old to S3 Glacier Deep Archive

D.Store the logs in Amazon CloudWatch Logs.Use Amazon S3 Lifecycle policies to move logs more than 1 month old to s3 Glacier Deep Archive

(单选题) Q121.A company is migrating a distributed application to AWS.The application serves variable workloads.The legacy platorm consists of a primary server that coordinates jobs across multiple compute nodes.The company wants to modernize the application with a solution that maximizes resiliency and scalability How should a solutions architect design the architecture to meet these requirements ?

A.Configure an Amazon Simple Queue Service (Amazon SQS) queue as a destination for the jobs.Implement the compute nodes with Amazon EC2 instances that are managed in an Auto Scaling group.Configure EC2 Auto Scaling to use scheduled scaling

B.Configure an Amazon Simple Queue Service (Amazon SQS) queue as a destination for the jobs.Implement the compute nodes with Amazon EC2 instances that are managed in an Auto Scaling group.Configure EC2 Auto Scaling based on the size of the queue

C.Implement the primary server and the compute nodes with Amazon EC2 instances that are managed in an Auto Scaling group.Configure AWS CloudTrail as a destination for the jobs.Configure EC2 Auto Scaling based on the load onthe primary server.

D.Implement the primary server and the compute nodes with Amazon EC2 instances that are managed in an Auto Scaling group.Configure Amazon EventBridge (Amazon CloudWatch Events) as a destination for the jobs.Configure EC2 Auto Scaling based on the load on the compute nodes.

(单选题) Q122.A company is storing backup files by using Amazon S3 Standard

storage. The files are accessed frequently for 1 month. However, the files are not accessed after 1 month. The company must keep the files indefinitely. Which storage solution will meet these requirements MOST cost-effectively ?

- A. Configure S3 Intelligent-Tiering to automatically migrate objects
- B. Create an S3 Lifecycle configuration to transition objects from S3 Standard to S3 Glacier Deep Archive after 1 month
- C. Create an S3 Lifecycle configuration to transition objects from S3 Standard to S3 Standard-Infrequent Access (S3 Standard-IA) after 1 month.
- D. Create an S3 Lifecycle configuration to transition objects from S3 Standard to S3 One Zone-Infrequent Access (S3 One Zone-IA) after 1 month

(单选题) Q123. A global company is using Amazon API Gateway to design REST APIs for its loyalty club users in the us-east-1 Region and the ap-southeast-2 Region. A solutions architect must design a solution to protect these API Gateway managed REST APIs across multiple accounts from SQL injection and cross-site scripting attacks. Which solution will meet these requirements with the LEAST amount of administrative effort ?

- A. Set up AWS WAF in both Regions. Associate Regional web ACLs with an API stage
- B. Set up AWS Firewall Manager in both Regions. Centrally configure AWS WAF rules.
- C. Set up AWS Shield in both Regions. Associate Regional web ACLs with an API stage.
- D. Set up AWS Shield in one of the Regions. Associate Regional web ACLs with an API stage

(单选题) Q124. A company delivers files in Amazon S3 to certain users who do not have AWS credentials. These users must be given access for a limited time. What should a solutions architect do to securely meet these requirements ?

- A. Configure public access on an Amazon S3 bucket.
- B. Generate a presigned URL to share with the users
- C. Encrypt files using AWS KMS and provide keys to the users
- D. Create and assign IAM roles that will grant GetObject permissions to the users

(单选题) Q125. A company is building an interactive knowledge base system to

help its call center staff work more efficiently.The system will be deployed on AWS and will be used 24 hours a day, 7 days a week.The company has more than 10 years of customer support transcripts stored in plaintext files.The company needs a database solution in which the company can bulk load existing transcripts and can individually load new transcripts as they are captured.The solution must provide search functionality to identify transcripts in which keywords and phrases occur.Which solution meets these requirements MOST cost-effectively ?

- A.Amazon Athena
- B.Amazon RDS
- C.Amazon OpenSearch Service (Amazon Elasticsearch Service)
- D.Amazon DynamoDB

(单选题) Q126.A company recently experienced a DDoS attack on its web application.The application is hosted on Amazon EC2 instances and is architected for high availability with Elastic Load Balancers and Auto Scaling groups.The DDoS attack lasted for an extended period of time and resulted in additional cost to the company.The company wants to acquire financial protection from future DDoS attacks Which solution will provide this protection MOST cost-effectively ?

- A.AWS WAF
- B.AWS Shield Standard
- C.AWS Shield Advanced
- D.Amazon Detective

(单选题) Q127.A gaming company hosts a browser-based application on AWS.The users of the application consume a large number of videos and images that are stored in Amazon S3.This content is the same for all users.The application has increased in popularity, and millions of users worldwide are accessing these media files.The company wants to provide the files to the users while reducing the load on the origin.Which solution meets these requirements MOST cost-effectively ?

- A.Deploy an AWS Global Accelerator accelerator in front of the web servers
- B.Deploy an Amazon CloudFront web distribution in front of the S3 bucket
- C.Deploy an Amazon ElastiCache for Redis instance in front of the web servers

D. Deploy an Amazon ElastiCache for Memcached instance in front of the web servers

(多选题) Q128. A company runs its two-tier ecommerce website on AWS. The web tier consists of a load balancer that sends traffic to Amazon EC2 instances. The database tier uses an Amazon RDS DB instance. The EC2 instances and the RDS DB instance should not be exposed to the public internet. The EC2 instances require internet access to complete payment processing of orders through a third-party web service. The application must be highly available. Which combination of configuration options will meet these requirements? (Select TWO.)

A. Use an Auto Scaling group to launch the EC2 instances in private subnets. Deploy an RDS Multi-AZ DB instance in private subnets

B. Configure a VPC with two private subnets and two NAT gateways across two Availability Zones. Deploy an Application Load Balancer in the private subnets

C. Use an Auto Scaling group to launch the EC2 instances in public subnets across two Availability Zones. Deploy an RDS Multi-AZ DB instance in private subnets

D. Configure a VPC with one public subnet, one private subnet, and two NAT gateways across two Availability Zones. Deploy an Application Load Balancer in the public subnet.

E. Configure a VPC with two public subnets, two private subnets, and two NAT gateways across two Availability Zones. Deploy an Application Load Balancer in the public subnets

(单选题) Q129. A company is serving an application through Amazon CloudFront. The company wants to protect the application from common SQL injection and cross-site scripting attacks. The company also wants the ability to block IP addresses and apply rate limiting. Which AWS service meets these requirements?

A. Amazon GuardDuty

B. AWS Shield

C. Amazon Inspector

D. AWS WAF

(单选题) Q130.A solutions architect is designing a new hybrid architecture to extend a company's on-premises infrastructure to AWS.The company requires a highly available connection with consistent low latency to an AWS Region.The company needs to minimize costs and is willing to accept slower traffic if the primary connection fails.What should the solutions architect do to meet these requirements ?

- A.Provision an AWS Direct Connect connection to a Region.Provision a VPN connection as a backup if the primary Direct Connect connection fails
- B.Provision a VPN tunnel connection to a Region for private connectivity.Provision a second VPN tunnel for private connectivity and as a backup if the primary VPN connection fails
- C.Provision an AWS Direct Connect connection to a Region.Provision a second Direct Connect connection to the same Region as a backup if the primary Direct Connect connection fails
- D.Provision an AWS Direct Connect connection to a Region.Use the Direct Connect failover attribute from the AWS CLI to automatically create a backup connection if the primary Direct Connect connection fails

(单选题) Q131.A company wants to store large amounts of data in Amazon S3 buckets.Numerous applications access the data, and the access pattern of the data is irregular.A solutions architect must recommend a cost-effective storage solution that does not affect performance or require operational overhead.Which solution meets these requirements ?

- A.AWS Trusted Advisor
- B.S3 Analytics
- C.S3 Intelligent-Tiering
- D.Cost Explorer

(单选题) Q132.A company needs to store its accounting records in Amazon S3.The records must be immediately accessible for 1 year and then must be archived for an additional 9 years.No one at the company, including administrative users and root users, can be able to delete the records during the entire 10-year period.The

records must be stored with maximum resiliency.Which solution will meet these requirements ?

- A.Store the records in S3 Glacier for the entire 10-year period.Use an access control policy to deny deletion of the records for a period of 10 years
- B.Store the records by using S3 Intelligent-Tiering.Use an IAM policy to deny deletion of the records.After 10 years, change the IAM policy to allow deletion
- C.Use an S3 Lifecycle policy to transition the records from S3 Standard to S3 Glacier Deep Archive after 1 year, Use S3 Object Lock in compliance mode for a period of 10 years
- D.Use an S3 Lifecycle policy to transition the records from S3 Standard to S3 One Zone-Infrequent Access (S3 One Zone-IA) after 1 year.Use S3 Object Lock in governance mode for a period of 10 years

(单选题) Q133.A company is designing an application where users upload small files into Amazon S3.After a user uploads a file, the file requires one-time simple processing to transform the data and save the data in JSON format for later analysis.Each file must be processed as quickly as possible after it is uploaded.Demand will vary.On some days, users will upload a high number of files.On other days, users will upload a few files or no files.Which solution meets these requirements with the LEAST operational overhead ?

- A.Configure Amazon EMR to read text files from Amazon S3.Run processing scripts to transform the data.Store the resulting JSON file in an Amazon Aurora DB cluster.
- B.Configure Amazon S3 to send an event notification to an Amazon Simple Queue Service (Amazon SQS) queue.Use Amazon EC2 instances to read from the queue and process the data Store the resulting JSON file in Amazon DynamoDB.
- C.Configure Amazon S3 to send an event notification to an Amazon Simple Queue Service (Amazon SQS) queue.Use an AWS Lambda function to read from the queue and process the data.Store the resulting JSON file in Amazon DynamoDB.
- D.Configure Amazon EventBridge (Amazon CloudWatch Events) to send an event to Amazon Kinesis Data Streams when a new file is uploaded.Use an AWS Lambda function to consume the event from the stream D) data Store the resulting JSON file in an Amazon Aurora DB cluster.

(单选题) Q134.A company hosts a marketing website in an on-premises data center.The website consists of static documents and runs on a single server.An administrator updates the website content infrequently and usesan SFTP client to upload new documents.The company decides to host its website on AWS and to use Amazon CloudFront.The company's solutions architect creates a CloudFront distribution.The solutions architect must design the most cost-effective and resilient architecture for website hosting to serve as the CloudFront origin.Which solution will meet these requirements ?

- A.Create a virtual server by using Amazon Lightsail Configure the web server in the Lightsail instance.Upload website content by using an SFTP client.
- B.Create an AWS Auto Scaling group for Amazon EC2 instances.Use an Application Load Balancer Upload website content by using an SFTP client.
- C.Create a private Amazon S3 bucket.Use an S3 bucket policy to allow access from a CloudFront origin access identity (OAI) .Upload website content by using the AWS CL.
- D.Create a public Amazon S3 bucket.Configure AWS Transfer for SFTP Configure the S3 bucket for website hosting.Upload website content by using the SFTP client.

(单选题) Q135.A company has a website hosted on AWS.The website is behind an Application Load Balancer (ALB) that is configured to handle HTTP and HTTPS separately.The company wants to forward all requests to the website so that the requests will use HTTPS.What should a solutions architect do to meet this requirement ?

- A.Update the ALB's network ACL to accept only HTTPS traffic
- B.Create a rule that replaces the HTTP in the URL with HTTPS
- C.Create a listener rule on the ALB to redirect HTTP traffic to HTTPS.
- D.Replace the ALB with a Network Load Balancer configured to use Server Name Indication (SNI)

(单选题) Q136.A startup company wants to decouple its three-tier application stack into microservices and use container technologies to build out the microservices.The company does not want to provision and manage the compute

infrastructure for its containers on AWS.What should a solutions architect recommend to meet these requirements ?

A.Provision an Amazon Elastic Container Service (Amazon ECS) cluster.Aach ECS nodes to an Amazon EC2 Auto Scaling group to host containers.

B.Provision an Amazon Elastic Container Service (Amazon ECS) cluster with AWS Fargate.Deploy containers to Fargate tasks.

C.Create AWS Lambda functions to host microservices.Integrate the functions with an Amazon API Gateway API to redirect application traffic.

D.Provision an Amazon Elastic Kubernetes Service (Amazon EKS) cluster.Attach EKS nodes to an Amazon EC2Auto Scaling group to host containers.

(单选题) Q137.A company needs to store data in Amazon S3 and must prevent the data from being changed.The company wants new objects that are uploaded to Amazon S3 to remain unchangeable for a nonspecific amount of time until the company decides to modify the objects.Only specific users in the companys AWS account can have the ability to delete the objects.What should a solutions architect do to meet these requirements ?

A.Create an S3 Glacier vault.Apply a write-once, read-many (WORM) vault lock policy to the objects.

B.Create an S3 bucket with S3 Object Lock enabled Enable versioning.Set a retention period of 100 years.Use governance mode as the S3 buckets default retention mode for new objects.

C.Create an s3 bucket.Use AWS CloudTrail to track any S3API events that modify the objects.Upon notification, restore the modied objects from any backup versions that the company has.

D.Create an S3 bucket with S3 Object Lock enabled.Enable versioning.Add a legal hold to the objects.Add the s3 : PutObjectlegalHold permission to the IAM policies of users who need to delete the objects.

(单选题) Q138.A global company hosts its web application on Amazon EC2 instances behind an Application Load Balancer (ALB) .The web application has static data and dynamic data.The company stores its static data in an Amazon S3

bucket The company wants to improve performance and reduce latency for the static data and dynamic data. The company is using its own domain name registered with Amazon Route 53. What should a solutions architect do to meet these requirements ?

A. Create an Amazon CloudFront distribution that has the S3 bucket and the ALB as origins. Configure Route 53 to route traffic to the CloudFront distribution.

B. Create an Amazon CloudFront distribution that has the ALB as an origin. Create an AWS Global Accelerator standard accelerator that has the S3 bucket as an endpoint. Configure Route 53 to route traffic to the CloudFront distribution.

C. Create an Amazon CloudFront distribution that has the S3 bucket as an origin. Create an AWS Global Accelerator standard accelerator that has the ALB and the CloudFront distribution as endpoints. Create a custom domain name that points to the accelerator DNS name. Use the custom domain name as an endpoint for the web application.

D. Create an Amazon CloudFront distribution that has the ALB as an origin. Create an AWS Global Accelerator standard accelerator that has the S3 bucket as an endpoint. Create two domain names. Point one domain name to the CloudFront DNS name for dynamic content. Point the other domain name to the accelerator DNS name for static content. Use the domain names as endpoints for the web application.

(单选题) Q139. A company recently migrated to AWS and wants to implement a solution to protect the traffic that flows in and out of the production VPC. The company had an inspection server in its on-premises data center.

The inspection server performed specific operations such as traffic flow inspection and traffic filtering. The company wants to have the same functionalities in the AWS Cloud. Which solution will meet these requirements ?

A. Use Amazon GuardDuty for traffic inspection and traffic filtering in the production VPC.

B. Use Traffic Mirroring to mirror traffic from the production VPC for traffic inspection and filtering.

C. Use AWS Network Firewall to create the required rules for traffic inspection and traffic filtering for the production VPC.

D. Use AWS Firewall Manager to create the required rules for traffic inspection and

traffic filtering for the production VPC.

(单选题) Q140.A company has an ordering application that stores customer information in Amazon RDS for MySQL.During regular business hours, employees run one-time queries for reporting purposes.Timeouts are occurring during order processing because the reporting queries are taking along time to run.The company needs to eliminate the timeouts without preventing employees from performing queries.What should a solutions architect do to meet these requirements ?

- A.Create a read replica.Move reporting queries to the read replica.
- B.Create a read replica.Distribute the ordering application to the primary DB instance and the read replica.
- C.Migrate the ordering application to Amazon DynamoDB with on-demand capacity.
- D.Schedule the reporting queries for non-peak hours.

(单选题) Q141.A company stores its application logs in an Amazon CloudWatch Logs log group.A new policy requires the company to store all application logs in Amazon OpenSearch Service (Amazon Elasticsearch Service) in near-real time.Which solution will meet this requirement with the LEAST operational overhead ? A.Configure a CloudWatch Logs subscription to stream the logs to Amazon OpenSearch Service (Amazon Elasticsearch Service) .

- B.Create an AWS Lambda ftunction.Use the log group to invoke the function to write the logs to Amazon OpenSearch Sevice (Amazon Elasticsearch Service)
- C.Create an Amazon Kinesis Data Firehose delivery stream.Configure the log group as the delivery stream's source.Configure Amazon OpenSearch Service (Amazon Elasticsearch Service) as the delivery stream's destination.
- D.Instail and configure Amazon Kinesis Agent on each application server to deliver the logs to Amazon Kinesis Data Streams.Configure Kinesis Data Streams to deliver the logs to Amazon OpenSearch Service (Amazon Elasticsearch Service)

(多选题) Q142.A company stores data in an Amazon Aurora PostgreSQL DB cluster.The company must store all the data for 5 years and must delete all the data after 5 years.The company also must indefinitely keep audit logs of actions that are

performed within the database. Currently, the company has automated backups configured for Aurora. Which combination of steps should a solutions architect take to meet these requirements? (Select TWO.)

- A. Take a manual snapshot of the DB cluster.
- B. Create a lifecycle policy for the automated backups.
- C. Configure automated backup retention for 5 years.
- D. Configure an Amazon CloudWatch Logs export for the DB cluster
- E. Use AWS Backup to take the backups and to keep the backups for 5 years.

(单选题) Q143. A rapidly growing ecommerce company is running its workloads in a single AWS Region. A solutions architect must create a disaster recovery (DR) strategy that includes a different AWS Region. The company wants its database to be up to date in the DR Region with the least possible latency. The remaining infrastructure in the DR Region needs to run at reduced capacity and must be able to scale up if necessary. Which solution will meet these requirements with the LOWEST recovery time objective (RTO)?

- A. Use an Amazon Aurora global database with a pilot light deployment.
- B. Use an Amazon Aurora global database with a warm standby deployment.
- C. Use an Amazon RDS Multi-AZ DB instance with a pilot light deployment.
- D. Use an Amazon RDS Multi-AZ DB instance with a warm standby deployment

(多选题) Q144. A company produces batch data that comes from different databases. The company also produces live stream data from network sensors and application APIs. The company needs to consolidate all the data into one place for business analytics. The company needs to process the incoming data and then stage the data in different Amazon S3 buckets. Teams will later run one-time queries and import the data into a business intelligence tool to show key performance indicators (KPIs). Which combination of steps will meet these requirements with the LEAST operational overhead? (Select TWO.)

- A. Use Amazon Athena for one-time queries. Use Amazon QuickSight to create dashboards for KPIs.
- B. Use Amazon Kinesis Data Analytics for one-time queries. Use Amazon QuickSight to

create dashboards for KPIs.

C.Create custom AWS Lambda functions to move the individual records from the databases to an Amazon Redshift cluster.

D.Use an AWS Glue extract, transform, and load (ETL) job to convert the data into JSON format.Load the data into multiple Amazon OpenSearch Service (Amazon Elasticsearch Service) clusters.

E.Use blueprints in AWS Lake Formation to identify the data that can be ingested into a data lake.Use AWS Glue to crawl the source, extract the data, and load the data into Amazon S3 in Apache Parquet format.

(单选题) Q145.A company runs its ecommerce application on AWS.Every new order is published as a message in a RabbitMQ queue that runs on an Amazon EC2 instance in a single Availability Zone.These messages are company needs to redesign its architecture to provide the highest availability with the least operational overheadWhat should a solutions architect do to meet these requirements ?

A.Migrate the queue to a redundant pair (active/standby) of RabbitMQ instances on Amazon MQ.Create a Multi-AZ Auto Scaling group for EC2 instances that host the application.Create another Multi-AZ Auto Scaling group for EC2 instances that host the PostgreSQL database.

B.Migrate the queue to a redundant pair (active/standby) of RabbitMQ instances on Amazon MQ.Create a Multi-AZ Auto Scaling group for EC2 instances that host the application.Migrate the database to run on a Multi-AZ deployment of Amazon RDS for PostgreSQL.

C.Create a Multi-AZ Auto Scaling group for EC2 instances that host the RabbitMQ queue.Create another Multi-AZ Auto Scaling group for EC2 instances that host the application.Migrate the database to run on a Multi-AZ deployment of Amazon RDS for PostgreSQL.

D.Create a Multi-AZ Auto Scaling group for EC2 instances that host the RabbitMQ queue.Create another Multi-AZ Auto Scaling group for EC2 instances that host the application.Create a third Muti-AZ Auto Scaling group for EC2 instances that host the PostgreSQL database,

(单选题) Q146.A company is deploying a new public web application to AWS.The application will run behind an Application Load Balancer (ALB) .The application needs to be encrypted at the edge with an SSL/TLS certificate that is issued by an external certificate authority (CA) .The certificate must be rotated each year before the certificate expires.What should a solutions architect do to meet these requirements? A.Use AWS Certificate Manager (ACM) to issue an SSL/TLS certificate.Apply the certificate to the ALB.Use the managed renewal feature to automatically rotate the certificate.

B.Use AWS Certificate Manager (ACM) to issue an SSL/TLS certificate.Import the key material from the certificate.Apply the certificate to the ALB.Use the managed renewal feature to automatically rotate the certificate.

C.Use AWS Certificate Manager (ACM) Private Certificate Authority to issue an SSL/TLS certificate from the root CA.Apply the certificate to the ALB.Use the managed renewal feature to automatically rotate the certificate.

D.Use AWS Certificate Manager (ACM) to import an SSL/TLS certificate.Apply the certificate to the ALB.Use Amazon EventBridge (Amazon CloudWatch Events) to send a notification when the certificate is nearing expiration.Rotate the certificate manually.

(单选题) Q147.A company uses NFS to store large video files in on-premises network attached storage.Each video file ranges in size from 1MB to 500 GB.The total storage is 70 TB and is no longer growing.The companydecides to migrate the video files to Amazon S3.The company must migrate the video files as soon as possible while using the least possible network bandwidth.Which solution will meet these requirements ?

A.Create an S3 bucket.Create an IAM role that has permissions to write to the S3 bucket.Use the AwS CLi to copy all files locally to the S3 bucket.

B.Create an AWS Snowball Edge job.Receive a Snowball Edge device on premise.Use the Snowball Edge client to transfer data to the device.Return the device so that AWS can import the data into Amazon S3.

C.Deploy an S3 File Gateway on premises.Create a public service endpoint to connect to the S3 File Gateway.Create an S3 bucket.Create a new NFS file share on the S3 File

Gateway.Point the new file share to the S3 bucket.Transfer the data from the existing NFS file share to the S3File Gateway.

D.Setup an AWS Direct Connect connection between the on-premises network and AWS.Deploy an S3 File Gateway on premises.Create a public virtual interface (ViF) to connect to the S3 File Gateway.Create an S3 bucket.Create a new NFS file share on the S3 File Gateway.Point the new file share to the S3 bucket.Transfer the data from the existing NFS file share to the S3 File Gateway.

(单选题) Q148.A company uses Amazon S3 as its data lake.The company has a new partner that must use SFTP to upload data files.A solutions architect needs to implement a highly available SFTP solution that minimizes operational overhead.Which solution will meet these requirements ?

A.Use AWS Transfer Family to configure an SFTP-enabled server with a publicly accessible endpoint.Choose the S3 data lake as the destination.

B.Use Amazon S3 File Gateway as an SFTP server.Expose the S3 File Gateway endpoint URL to the new partner.Share the S3 File Gateway endpoint with the new partner.

C.Launch an Amazon EC2 instance in a private subnet in a VPC.Instruct the new partner to upload files to the EC2 instance by using a VPN.Run a cron job script on the EC2 instance to upload files to the S3 data lake.

D.Launch Amazon EC2 instances in a private subnet in a VPC.Place a Network Load Balancer (NLB) in front of the EC2 instances.Create an SFTP listener port for the NLB.Share the NLB hostname with the new partner.Run a cron job script on the EC2 instances to upload files to the S3 data lake.

(单选题) Q149.A company has migrated a two-tier application from its on-premises data center to the AWS Cloud.The data tier is a Multi-AZ deployment of Amazon RDS for Oracle with 12TB of General Purpose SSD Amazon Elastic Block Store (Amazon EBS) storage.The application is designed to process and store documents in the database as binary large objects (blobs) with an average document size of 6 MB.The database size has grown over time, reducing the performance and increasing the cost of storage.The company must improve the database performance and needs a solution that is highly available and

resilient.Which solution will meet these requirements MOST cost-effectively ?

A.Reduce the RDS DB instance size.Increase the storage capacity to 24 TiB.Change the storage type to Magnetic.

B.Increase the RDS DB instance size.Increase the storage capacity to 24 TiB.Change the storage type to Provisioned IOPS.

C.Create an Amazon S3 bucket.Update the application to store documents in the S3 bucket.Store the object metadata in the existing database.

D.Create an Amazon DynamoDB table.Update the application to use DynamoDB.Use AWS Database Migration Service (AWS DMS) to migrate data from the Oracle database to DynamoDB.

(单选题) Q150.A company has a small Python application that processes JSON documents and outputs the results to an on-premises SQL database.The application runs thousands of times each day.The company wants to move the application to the AWS Cloud.The company needs a highly available solution that maximizes scalability and minimizes operational overhead.Which solution will meet these requirements ?

A.Place the JSON documents in an Amazon S3 bucket.Run the Python code on multiple Amazon EC2 instances to process the documents.Store the results in an Amazon Aurora DB cluster.

B.Place the JSON documents in an Amazon S3 bucket.Create an AWS Lambda function that runs the Python code to process the documents as they arrive in the S3 bucket.Store the results in an Amazon Aurora DB cluster.

C.Place the JSON documents in an Amazon Elastic Block Store (Amazon EBS) volume.Use the EBS Multi-Attach feature to attach the volume to multiple Amazon EC2 instances.Run the Python code on the EC2 instances to process the documents.Store the results on an Amazon RDS DB instance.

D.Place the JSON documents in an Amazon Simple Queue Service (Amazon SQS) queue as messages.Deploy the Python code as a container on an Amazon Elastic Container Service (Amazon ECS) cluster that is configured with the Amazon EC2 launch type.Use the container to process the SQS messages.Store the results on an Amazon RDS DB instance.

(单选题) Q151.A company is running several business applications in three separate VPCs within the us-east-1 Region.The applications must be able to communicate between VPCs.The applications also must be able to consistently send hundreds of gigabytes of data each day to a latency-sensitive application that runs in a single on-premises data center.A solutions architect needs to design a network connectivity solution that maximizes cost-effectiveness.Which solution meets these requirements ?

- A.Configure three AWS Site-to-Site VPN connections from the data center to AWS.Establish connectivity by configuring one VPN connection for each VPC.
- B.Launch a third-party virtual network appliance in each VPC.Establish an IPsec VPN tunnel between the data center and each virtual appliance.
- C.Set up three AWS Direct Connect connections from the data center to a Direct Connect gateway in us-east-1.Establish connectivity by configuring each VPC to use one of the Direct Connect connections.
- D.Set up one AWS Direct Connect connection from the data center to AWS.Create a transit gateway, and attach each VPC to the transit gateway.Establish connectivity between the Direct Connect connection and the transit gateway

(单选题) Q152.A company needs to store data from its healthcare application.The application's data frequently changes.A new regulation requires audit access a tall levels of the stored data.The company hosts the application on an on-premises infrastructure that is running out of storage capacity.A solutions architect must securely migrate the existing data to AWS while satisfying the new regulation.Which solution will meet these requirements ?

- A.Use AWS DataSync to move the existing data to Amazon S3.Use AWS CloudTrail to log data events.
- B.Use AWS Snowcone to move the existing data to Amazon S3.Use AWS CloudTrail to log management events.
- C.Use Amazon S3 Transfer Acceleration to move the existing data to Amazon S3.Use AWS CloudTrail to log data events.
- D.Use AWS Storage Gateway to move the existing data to Amazon S3.Use AWS CloudTrail to log management events.

(单选题) Q153.A company has an ecommerce checkout workflow that writes an order to a database and calls a service to process the payment.Users are experiencing timeouts during the checkout process.When users resubmit the checkout form, multiple unique orders are created for the same desired transaction.How should a solutions architect refactor this workflow to prevent the creation of multiple orders ?

A.Configure the web application to send an order message to Amazon Kinesis Data Firehose.Set the payment service to retrieve the message from Kinesis Data Firehose and process the order.

B.Create a rule in AWS CloudTrail to invoke an AWS Lambda function based on the logged application path request.Use Lambda to query the database, call the payment service, and pass in the order information.

C.Store the order in the database.Send a message that includes the order number to Amazon Simple Notification Service (Amazon SNS).Set the payment service to poll Amazon SNS, retrieve the message, and process the order.

D. Store the order in the database.Send a message that includes the order number to an Amazon Simple Queue Service (Amazon SQS) FIFO queue.Set the paymentservice to retrieve the message and process the order.Delete the message from the queue.

(多选题) Q154.A company used an Amazon RDS for MySQL DB instance during application testing.Before terminating the DB instance at the end of the test cycle, a solutions architect created two backups.The solutions architect created the first backup by using the mysqldump utility to create a database dump.The solutions architect created the second backup by enabling the final DB snapshot option on RDS termination.The company is now planning for a new test cycle and wants to create a new DB instance from the most recent backup.The company has chosen a MySQL-compatible edition of Amazon Aurora to host the DB instance.Which solutions will create the new DB instance ? (Select TWO.)

A.Import the RDS snapshot directly into Aurora.

B.Upload the RDS snapshot to Amazon S3.Then import the RDS snapshot into Aurora.

C.Upload the database dump to Amazon S3.Then import the database dump into Aurora.

D.Use AWS Database Migration Service (AWS DMS) to import the RDS snapshot into Aurora.

E.Upload the database dump to Amazon S3.Then use AWS Database Migration Service (AWS DMS) to import the database dump into Aurora

(单选题) Q155.A company is creating an application that runs on containers in a VPC.The application stores and accesses data in an Amazon S3 bucket.During the development phase, the application will store and access 1 TB of data in Amazon S3 each day.The company wants to minimize costs and wants to prevent traffic from traversing the internet whenever possible.Which solution will meet these requirements ?

A.Enable S3 Intelligent-Tiering for the S3 bucket.

B.Enable S3 Transfer Acceleration for the S3 bucket.

C.Create a gateway VPC endpoint for Amazon S3.Associate this endpoint with all route tables in the VPC.

D.Create an interface endpoint for Amazon S3in the VPC.Associate this endpoint with all route tables in the VPC.

(多选题) Q156.A rapidly growing global ecommerce company is hosting its web application on AWS.The web application includes static content and dynamic content.The website stores online transaction processing (OLTP) data in an Amazon RDS database.The website's users are experiencing slow page loads.Which combination of actions should a solutions architect take to resolve this issue ? (Select TWO.)

A.Configure an Amazon Redshift cluster.

B.Set up an Amazon CloudFront distribution.

C.Host the dynamic web content in Amazon S3.

D.Create a read replica for the RDS DB instance.

E.Configure a Multi-AZ deployment for the RDS DB instance.

(多选题) Q157.A company needs to store contract documents.A contract lasts for 5 years.During the 5-year period, the company must ensure that the documents cannot be overwritten or deleted.The company needs to encrypt the documents at rest and rotate the encryption keys automatically every year.Which combination of steps should a solutions architect take to meet these requirements with the LEAST operational overhead ? (Select TWO.)

- A.Store the documents in Amazon S3.Use S3 Object Lock in governance mode.
- B.Store the documents in Amazon S3.Use S3 Object Lock in compliance mode.
- C.Use server-side encryption with Amazon S3 managed encryption keys (SSE-S3) . Configure key rotation.
- D.Use server-side encryption with AWS Key Management Service (AWS KMS) customer managed keys.Configure key rotation.
- E.Use server-side encryption with AWS Key Management Service (AWS KMS) customer provided (imported) keys.Configure key rotation.

(单选题) Q158.A solutions architect is migrating a document management workload to AWS.The workload keeps 7 TiB of contract documents on a shared storage file system and tracks them on an external database.Most of the documents are stored and retrieved eventually for reference in the future.The application cannot be modified during the migration, and the storage solution must be highly available.Documents are retrieved and stored by web servers that run on Amazon EC2 instances in an Auto Scaling group.The Auto Scaling group can have up to 12 instances.Which solution meets these requirements MOST cost-effectively ?

- A.Provision an enhanced networking optimized EC2 instance to serve as a shared NFS storage system.
- B.Create an Amazon S3 bucket that uses the S3 Standard-Infrequent Access (S3Standard-IA) storage class.Mount the S3 bucket to the EC2 instances in the Auto Scaling group.
- C.Create an SFTP server endpoint by using AWS Transfer for SFTP and an Amazon S3 bucket.Configure the EC2 instances in the Auto Scaling group to connect to the SFTP server.
- D.Create an Amazon Elastic File System (Amazon EFS) file system that uses the EFS

Standard-Infrequent Access (EFS Standard-IA) storage class. Mount the file system to the EC2 instances in the Auto Scaling group.

(单选题) Q159. A company is using a SQL database to store movie data that is publicly accessible. The database runs on an Amazon RDS Single-AZ DB instance. A script runs queries at random intervals each day to record the number of new movies that have been added to the database. The script must report a final total during business hours. The company's development team notices that the database performance is inadequate for development tasks when the script is running. A solutions architect must recommend a solution to resolve this issue. Which solution will meet this requirement with the LEAST operational overhead?

- A. Modify the DB instance to be a Multi-AZ deployment.
- B. Create a read replica of the database. Configure the script to query only the read replica.
- C. Instruct the development team to manually export the entries in the database at the end of each day.
- D. Use Amazon ElastiCache to cache the common queries that the script runs against the database.

(多选题) Q160. A gaming company is moving its public scoreboard from a data center to the AWS Cloud. The company uses Amazon EC2 Windows Server instances behind an Application Load Balancer to host its dynamic application. The company needs a highly available storage solution for the application. The application consists of static files and dynamic server-side code. Which combination of steps should a solutions architect take to meet these requirements? (Select TWO.)

- A. Store the static files on Amazon S3. Use Amazon CloudFront to cache objects at the edge.
- B. Store the static files on Amazon S3. Use Amazon ElastiCache to cache objects at the edge.
- C. Store the server-side code on Amazon Elastic File System (Amazon EFS). Mount the EFS volume on each EC2 instance to share the files.
- D. Store the server-side code on Amazon FSx for Windows File Server. Mount the FSx

for Windows File Server volume on each EC2 instance to share the files.

E.Store the server-side code on a General Purpose SSD (gp2) Amazon Elastic Block Store (Amazon EBS) volume.Mount the EBS volume on each EC2 instance to share the files.

(单选题) Q161.A company has an application that places hundreds of.csv files into an Amazon S3 bucket every hour.The files are 1GB in size.Each time a file is uploaded, the company needs to convert the file to Apache Parquet format and place the outputfile into an S3 bucket.Which solution will meet these requirements with the LEAST operational overhead ?

A.Create an AWS Lambda function to download the.csv files, convert the files to Parquet format, and place the output files in an S3 bucket.Invoke the Lambda function for each S3 PUT event

B.Create an Apache Spark job to read the.csv files, convert the files to Parquet format, and place the output files in an S3 bucket.Create an AWS Lambda function for each S3PUTevent to invoke the Spark job.

C.Create an AWS Glue table and an AWS Glue crawler for the S3 bucket where the application places the.csv files.Schedule an AWS Lambda function to periodically use Amazon Athena to query the AWS Glue table, convert the query results into Parquet format, and place the output files into an S3 bucket.

D.Create an AWS Glue extract, transform, and load (ETL) job to convert the.csv files to Parquet format and place the output files into an S3 bucket.Create an AWS Lambda function for each S3PUT event to invoke the ETL job.

(单选题) Q162.A company is creating a prototype of an ecommerce website on AWS.The website consists of an Application Load Balancer, an Auto Scaling group of Amazon EC2 instances for web servers, and an Amazon RDS for MySQL DB instance that runs with the Single-AZ configuration.The website is slow to respond during searches of the product catalog.The product catalog is a group of tables in the MySQL database that the company does not update frequently.A solutions architect has determined that the CPU utilization on the DB instance is high when product catalog searches occur.What should the solutions architect recommend to

improve the performance of the website during searches of the product catalog ?

A.Migrate the product catalog to an Amazon Redshift database.Use the COPY command to load the product catalog tables.

B.Implement an Amazon ElastiCache for Redis cluster to cache the product catalog, Use lazy loading to populate the cache.

C.Add an additional scaling policy to the Auto Scaling group to launch additionalEC2 instances when database response is slow.

D.Tun on the Multi-AZ configuration for the DB instance.Configure the EC2 instances to throttle the product catalog queries that are sent to the database.

(单选题) Q163.A company hosts a data lake on AWS.The data lake consists of data in Amazon S3 and Amazon RDS for PostgreSQL.The company needs a reporting solution that provides data visualization and includes all the data sources within the data lake.Only the company's management team should have full access to all the visualizations.The rest of the company should have only limited access.Which solution will meet these requirements ?

A.Create an analysis in Amazon QuickSight.Connect all the data sources and create new datasets.Pubish dashboards to visualize the data.Share the dashboards with the appropriate IAM roles.

B.Create an analysis in Amazon QuickSight.Connect all the data sources and create new datasets.Pubish dashboards to visualize the data.Share the dashboards with the appropriate users and groups.

C.Create an AWS Glue table and crawler for the data in Amazon S3.Create an AWS Glue extract, transform, and load (ETL) job to produce reports.Publish the reports to Amazon S3.Use S3 bucket policies to limit access to the reports.

D.Create an AWS Glue table and crawler for the data in Amazon S3.Use Amazon Athena Federated Query to access data within Amazon RDS for PostgreSQL.Generate reports by using Amazon Athena.Publish the reports to Amazon S3.Use S3 bucket policies to limit access to the reports

(单选题) Q164.A company needs to move data from an Amazon EC2 instance to an Amazon S3 bucket.The company must ensure that no API calls and no data are

routed through public internet routes.Only the EC2 instance can have access to upload data to the S3 bucket.Which solution will meet these requirements ?

A.Create an interface VPC endpoint for Amazon S3 in the subnet where the EC2 instance is located.Attach a resource policy to the S3 bucket to only allow the EC2 instance's IAM role for access.

B.Create a gateway VPC endpoint for Amazon S3 in the Availability Zone where the EC2 instance is located.Attach appropriate security groups to the endpoint.Attach a resource policy to the S3 bucket to only allow the EC2 instance's IAM role for access.

C.Run the nslookup tool from inside the EC2 instance to obtain the private IP address of the S3 bucket's service API endpoint. Create a route in the VPC route table to provide the EC2 instance with access to the S3 bucket. Attach a resource policy to the S3 bucket to only allow the EC2 instance's IAM role for access.

D.Use the AWS provided, publicly available ip-ranges.json file to obtain the private IP address of the S3 bucket's service API endpoint.Create a route in the VPC route table to provide the EC2 instance with access to the S3 bucket.

Attach a resource policy to the S3 bucket to only allow the EC2 instance's IAM role for access.

(单选题) Q165.A company runs a global web application on Amazon EC2 instances behind an Application Load Balancer.The application stores data in Amazon Aurora The company needs to create a disaster recovery solution and can tolerate up to 30 minutes of downtime and potential data loss.The solution does not need to handle the load when the primary infrastructure is healthy.What should a solutions architect do to meet these requirements ?

A.Deploy the application with the required infrastructure elements in place.Use Amazon Route 53 to configure active-passive failover.Create an Aurora Replica in a second AWS Region

B.Host a scaled-down deployment of the application in a second AWS Region.Use Amazon Route 53 to configure active-active failover.Create an Aurora Replica in the second Region.

C.Replicate the primary infrastructure in a second AWS Region.Use Amazon Route

53to configure active-active failover.Create an Aurora database that is restored from the latest snapshot

D.Back up data with AWS Backup.Use the backup to create the required infrastructure in a second AWS Region.Use Amazon Route 53to configure active-passive failover.Create an Aurora second primary instance in the second Region.

(单选题) Q166.A company performs monthly maintenance on its AWS infrastructure.During these maintenance activities, the company needs to rotate the credentials for its Amazon RDS for MySQL databases across multiple AWS Regions.Which solution will meet these requirements with the LEAST operational overhead ?

A.Store the credentials as secrets in AWS Secrets Manager.Use multi-Region secret replication for the required Regions.Configure Secrets Manager to rotate the secrets on a schedule

B.Store the credentials as secrets in AWS Systems Manager by creating a secure string parameter.Use multi-Region secret replication for the required Regions.Configure Systems Manager to rotate the secrets on a schedule.

C.Store the credentials in an Amazon S3 bucket that has server-side encryption (SSE) enabled.Use Amazon EventBridge (Amazon CloudWatch Events) to invoke an AWS Lambda function to rotate the credentials.

D.Encrypt the credentials as secrets by using AWS Key Management Service (AWS KMS) multi-Region customer managed keys. Store the secrets in an Amazon DynamoDB global table. Use an AWS Lambda function to retrieve the secrets from DynamoDB.Use the RDS API to rotate the secrets.

(单选题) Q167.A payment processing company records all voice communication with its customers and stores the audio files in an Amazon S3 bucket.The company needs to capture the text from the audio files.The company must remove from the text any personally identifiable information (PII) that belongs to customers.What should a solutions architect do to meet these requirements ?

A.Process the audio files by using Amazon Kinesis Video Streams.Use an AWS Lambda function to scan for known PII patterns

- B. When an audio file is uploaded to the S3 bucket, invoke an AWS Lambda function to start an Amazon Textract task to analyze the call recordings
- C. Configure an Amazon Transcribe transcription job with PII redaction turned on. When an audio file is uploaded to the S3 bucket, invoke an AWS Lambda function to start the transcription job. Store the output in a separate S3 bucket
- D. Create an Amazon Connect contact flow that ingests the audio files with transcription turned on. Embed an AWS Lambda function to scan for known PII patterns. Use Amazon EventBridge (Amazon CloudWatch Events) to start the contact flow when an audio file is uploaded to the S3 bucket.

(单选题) Q168. A company has a business system that generates hundreds of reports each day. The business system saves the reports to a network share in CSV format. The company needs to store this data in the AWS Cloud in near-real time for analysis. Which solution will meet these requirements with the LEAST administrative overhead?

- A. Use AWS DataSync to transfer the files to Amazon S3. Create a scheduled task that runs at the end of each day
- B. Create an Amazon S3 File Gateway. Update the business system to use a new network share from the S3 File Gateway.
- C. Use AWS DataSync to transfer the files to Amazon S3. Create an application that uses the DataSync API in the automation workflow
- D. Deploy an AWS Transfer for SFTP endpoint. Create a script that checks for new files on the network share and uploads the new files by using SFTP

(单选题) Q169. A company stores confidential data in an Amazon Aurora PostgreSQL database in the ap-southeast-3 Region. The database is encrypted with an AWS Key Management Service (AWS KMS) customer managed key. The company was recently acquired and must securely share a backup of the database with the acquiring company's AWS account in ap-southeast-3. What should a solutions architect do to meet these requirements?

- A. Create a database snapshot. Copy the snapshot to a new unencrypted snapshot. Share the new snapshot with the acquiring company's AWS account

B.Create a database snapshot.Add the acquiring company's AWS account to the KMS key policy.Share the snapshot with the acquiring company's AWS account.

C.Create a database snapshot that uses a different AWS managed KMS key.Add the acquiring company's AWS account to the KMS key alias.Share the snapshot with the acquiring company's AWS account.

D.Create a database snapshot.Download the database snapshot.Upload the database snapshot to an Amazon S3 bucket.Update the S3 bucket policy to allow access from the acquiring company's AWS account.

(单选题) Q170.A company hosts a serverless application on AWS.The application uses Amazon API Gateway, AWS Lambda, and an Amazon RDS for PostgreSQL database.The company notices an increase in application errors that result from database connection timeouts during times of peak traffic or unpredictable traffic.The company needs a solution that reduces the application failures with the least amount of change to the code.What should a solutions architect do to meet these requirements ?

A.Reduce the Lambda concurrency rate.

B.Enable RDS Proxy on the RDS DB instance.

C.Resize the RDS DB instance class to accept more connections.

D.Migrate the database to Amazon DynamoDB with on-demand scaling.

(单选题) Q171.A company is building a new dynamic ordering website.The company wants to minimize server maintenance and patching.The website must be highly available and must scale read and write capacity as quickly as possible to meet changes in user demand.Which solution will meet these requirements ?

A.Host static content in Amazon S3.Host dynamic content by using Amazon API Gateway and AWS Lambda.Use Amazon DynamoDB with on-demand capacity for the database.Configure Amazon CloudFront to deliver the website content.

B.Host static content in Amazon S3.Host dynamic content by using Amazon API Gateway and AWS Lambda.Use Amazon Aurora with Aurora Auto Scaling for the database.Configure Amazon CloudFront to deliver the websit content.

C.Host all the website content on Amazon EC2 instances.Create an Auto Scaling

group to scale the EC2 instances. Use an Application Load Balancer to distribute traffic. Use Amazon DynamoDB with provisioned write capacity for the database.

D. Host all the website content on Amazon EC2 instances. Create an Auto Scaling group to scale the EC2 instances. Use an Application Load Balancer to distribute traffic. Use Amazon Aurora with Aurora Auto Scaling for the database.

(单选题) Q172. An online retail company has more than 50 million active customers and receives more than 25,000 orders each day. The company collects purchase data for customers and stores this data in Amazon S3. Additional customer data is stored in Amazon RDS. The company wants to make all the data available to various teams so that the teams can perform analytics. The solution must provide the ability to manage fine-grained permissions for the data and must minimize operational overhead. Which solution will meet these requirements?

A. Migrate the purchase data to write directly to Amazon RDS. Use RDS access controls to limit access.

B. Schedule an AWS Lambda function to periodically copy data from Amazon RDS to Amazon S3. Create an AWS Glue crawler. Use Amazon Athena to query the data. Use S3 policies to limit access.

C. Create a data lake by using AWS Lake Formation. Create an AWS Glue JDBC connection to Amazon RDS. Register the S3 bucket in Lake Formation. Use Lake Formation access controls to limit access.

D. Create an Amazon Redshift cluster. Schedule an AWS Lambda function to periodically copy data from Amazon S3 and Amazon RDS to Amazon Redshift. Use Amazon Redshift access controls to limit access.

(单选题) Q173. A solutions architect needs to implement a solution to reduce a company's storage costs. All the company's data is in the Amazon S3 Standard storage class. The company must keep all data for at least 25 years. Data from the most recent 2 years must be highly available and immediately retrievable. Which solution will meet these requirements?

A. Set up an S3 Lifecycle policy to transition objects to S3 Glacier Deep Archive immediately.

B.Set up an S3 Lifecycle policy to transition objects to S3 Glacier Deep Archive after 2 years.

C.Use S3 Intelligent-Tiering.Activate the archiving option to ensure that data is archived in S3 Glacier Deep Archive.

D.Set up an S3 Lifecycle policy to transition objects to S3 One Zone-Infrequent Access (S3 One Zone-IA) immediately and to S3 Glacier Deep Archive after 2 years.

(单选题) Q174.A company wants to migrate its existing on-premises monolithic application to AWS.The company wants to keep as much of the front-end code and the backend code as possible.However, the company wants to break the application into smaller applications.A different team will manage each application.The company needs a highly scalable solution that minimizes operational overhead.Which solution will meet these requirements ?

A.Host the application on AWS Lambda.Integrate the application with Amazon API Gateway

B.Host the application with AWS Amplify.Connect the application to an Amazon API Gateway API that is integrated with AWS Lambda

C.Host the application on Amazon EC2 instances.Set up an Application Load Balancer with EC2 instances in an Auto Scaling group as targets

D.Host the application on Amazon Elastic Container Service (Amazon ECS) .Set up an Application Load Balancer with Amazon ECS as the target

(多选题) Q175.A company is preparing a new data platform that will ingest real-time streaming data from multiple sources.The company needs to transform the data before writing the data to Amazon S3.The company needs the ability to use SQL to query the transformed data.Which solutions will meet these requirements ? (Select Two.)

A.Use Amazon Kinesis Data Streams to stream the data.Use Amazon Kinesis Data Analytics to transform the data.Use Amazon Kinesis Data Firehose to write the data to Amazon S3.Use Amazon Athena to query the transformed data from Amazon S3

B.Use Amazon Managed Streaming for Apache Kafka (Amazon MSK) to stream the data.Use AWS Glue to transform the data and to write the data to Amazon S3.Use

Amazon Athena to query the transformed data from Amazon S3

C.Use AWS Database Migration Service (AWS DMS) to ingest the data.Use Amazon EMR to transform the data and to write the data to Amazon S3.Use Amazon Athena to query the transformed data from Amazon S3

D.Use Amazon Managed Streaming for Apache Kafka (Amazon MSK) to stream the data.Use Amazon Kinesis Data Analytics to transform the data and to write the data to Amazon S3.Use the Amazon RDS query editor to query the transformed data from Amazon S3

E.Use Amazon Kinesis Data Streams to stream the data.Use AWS Glue to transform the data.Use Amazon Kinesis Data Firehose to write the data to Amazon S3.Use the Amazon RDS query editor to query the transformed data from Amazon S3

(单选题) Q176.A company has a three-tier web application that is deployed on AWS.The web servers are deployed in a public subnet in a VPC.The application servers and database servers are deployed in private subnets in the same VPC.The company has deployed a third-party virtual firewall appliance from AWS Marketplace in an inspection VPC.The appliance is configured with an IP interface that can accept IP packets.A solutions architect needs to integrate the web application with the appliance to inspect all traffic to the application before the traffic reaches the web server.Which solution will meet these requirements with the LEAST operational overhead ?

A.Create a Network Load Balancer in the public subnet of the application's VPC to route the traffic to the appliance for packet inspection.

B.Create an Application Load Balancer in the public subnet of the application's VPC to route the traffic to the appliance for packet inspection.

C.Deploy a transit gateway in the inspection VPC.Configure route tables to route the incoming packets through the transit gateway.

D.Deploy a Gateway Load Balancer in the inspection VPC.Create a Gateway Load Balancer endpoint to receive the incoming packets and forward the packets to the appliance.

(单选题) Q177.A company recently launched a variety of new workloads on

Amazon EC2 instances in its AWS account. The company needs to create a strategy to access and administer the instances remotely and securely. The company needs to implement a repeatable process that works with native AWS services and follows the AWS Well-Architected Framework. Which solution will meet these requirements with the LEAST operational overhead?

A. Use the EC2 serial console to directly access the terminal interface of each instance for administration.

B. Attach the appropriate IAM role to each existing instance and new instance. Use AWS Systems Manager Session Manager to establish a remote SSH session.

C. Create an administrative SSH key pair. Load the public key into each EC2 instance. Deploy a bastion host in a public subnet to provide a tunnel for administration of each instance.

D. Establish an AWS Site-to-Site VPN connection. Instruct administrators to use their local on-premises machines to connect directly to the instances by using SSH keys across the VPN tunnel.

(单选题) Q178. A company is building a solution that will report Amazon EC2 Auto Scaling events across all the applications in an AWS account. The company needs to use a serverless solution to store the EC2 Auto Scaling status data in Amazon S3. The company then will use the data in Amazon S3 to provide near-real-time updates in a dashboard. The solution must not affect the speed of EC2 instance launches. How should the company move the data to Amazon S3 to meet these requirements?

A. Use an Amazon CloudWatch metric stream to send the EC2 Auto Scaling status data to Amazon Kinesis Data Firehose. Store the data in Amazon S3.

B. Launch an Amazon EMR cluster to collect the EC2 Auto Scaling status data and send the data to Amazon Kinesis Data Firehose. Store the data in Amazon S3.

C. Create an Amazon EventBridge (Amazon CloudWatch Events) rule to invoke an AWS Lambda function on a schedule. Configure the Lambda function to send the EC2 Auto Scaling status data directly to Amazon S3.

D. Use a bootstrap script during the launch of an EC2 instance to install Amazon Kinesis Agent. Configure Kinesis Agent to collect the EC2 Auto Scaling status data and send the data to Amazon Kinesis Data Firehose. Store the data in Amazon S3.

(多选题) Q179.A company uses a popular content management system (CMS) for its corporate website.However, the required patching and maintenance are burdensome.The company is redesigning its website and wants a new solution.The website will be updated four times a year and does not need to have any dynamic content available.The solution must provide high scalability and enhanced security.Which combination of changes will meet these requirements with the LEAST operational overhead ? (Select TWO.)

A.Configure Amazon CloudFront in front of the website to use HTTPS functionality.
B.Deploy an AWS WAF web ACL in front of the website to provide HTTPS functionality. C.Create and deploy an AWS Lambda function to manage and serve the website content.

D.Create the new website and an Amazon S3 bucket.Deploy the website on the S3 bucket with static website hosting enabled.

D.Create the new website.Deploy the website by using an Auto Scaling group of Amazon EC2 instances behind an Application Load Balancer.

(单选题) Q180.A company is building a containerized application on premises and decides to move the application to AWS.The application will have thousands of users soonafter it is deployed.The company is unsure how to manage the deployment of containers at scale.The company needs to deploy the containerized application in a highly available architecture that minimizes operational overhead.Which solution will meet these requirements ?

A.Store container images in an Amazon Elastic Container Registry (Amazon ECR) repository.Use an Amazon Elastic Container Service (Amazon ECS) cluster with the AWS Fargate launch type to run the containers.Use target tracking to scale automatically based on demand.

B.Store container images in an Amazon Elastic Container Registry (Amazon ECR) repository.Use an Amazon Elastic Container Service (Amazon ECS) cluster with the Amazon EC2 launch type to run the containers.Use target tracking to scale automatically based on demand.

C.Store container images in a repository that runs on an Amazon EC2 instance.Run

the containers on EC2 instances that are spread across multiple Availability Zones. Monitor the average CPU utilization in Amazon CloudWatch. Launch new EC2 instances as needed.

D. Create an Amazon EC2 Amazon Machine Image (AMI) that contains the container image. Launch EC2 instances in an Auto Scaling group across multiple Availability Zones. Use an Amazon CloudWatch alarm to scale out EC2 instances when the average CPU utilization threshold is breached.

(单选题) Q181. A company is running a business-critical web application on Amazon EC2 instances behind an Application Load Balancer. The EC2 instances are in an Auto Scaling group. The application uses an Amazon Aurora PostgreSQL database that is deployed in a single Availability Zone. The company wants the application to be highly available with minimum downtime and minimum loss of data. Which solution will meet these requirements with the LEAST operational effort?

- A. Place the EC2 instances in different AWS Regions. Use Amazon Route 53 health checks to redirect traffic. Use Aurora PostgreSQL Cross-Region Replication.
- B. Configure the Auto Scaling group to use multiple Availability Zones. Configure the database as Multi-AZ. Configure an Amazon RDS Proxy instance for the database.
- C. Configure the Auto Scaling group to use one Availability Zone. Generate hourly snapshots of the database. Recover the database from the snapshots in the event of a failure.
- D. Configure the Auto Scaling group to use multiple AWS Regions. Write the data from the application to Amazon S3. Use S3 Event Notifications to launch an AWS Lambda function to write the data to the database.

(多选题) Q182. A company is migrating its on-premises PostgreSQL database to Amazon Aurora PostgreSQL. The on-premises database must remain online and accessible during the migration. The Aurora database must remain synchronized with the on-premises database. Which combination of actions must a solutions architect take to meet these requirements? (Select TWO.)

- A. Create an ongoing replication task.
- B. Create a database backup of the on-premises database.

- C.Create an AWS Database Migration Service (AWS DMS) replication server.
- D.Convert the database schema by using the AWS Schema Conversion Tool (Aws sCT)
- E.Create an Amazon EventBridge (Amazon CloudWatch Events) rule to monitor the database synchronization.

(多选题) Q183.A company is developing an ecommerce application that will consist of a load-balanced front end, a container-based application, and a relational database.A solutions architect needs to create a highly available solution that operates with as little manual intervention as possible.Which solutions meet these requirements ? (Select Two.)

- A.Create an Amazon RDS DB instance in Multi-AZ mode.
- B.Create an Amazon RDS DB instance and one or more replicas in another Availability Zone.
- C.Create an Amazon EC2 instance-based Docker cluster to handle the dynamic application load.
- D.Create an Amazon Elastic Container Service (Amazon ECS) cluster with a Fargate launch type to handle the dynamic application load.
- E.Create an Amazon Elastic Container Service (Amazon ECS) cluster with an Amazon EC2 launch type to handle the dynamic application load.

(单选题) Q184. A company's HTTP application is behind a Network Load Balancer (NLB). The NLB's target group is configured to use an Amazon EC2 Auto Scaling group with multiple EC2 instances that run the web service. The company notices that the NLB is not detecting HTTP errors for the application. These errors require a manual restart of the EC2 instances that run the web service. The company needs to improve the application's availability without writing custom scripts or code. What should a solutions architect do to meet these requirements?

- A. Enable HTTP health checks on the NLB supplying the URL of the company's application
- B. Add a cron job to the EC2 instances to check the local application's logs once each minute. If HTTP errors are detected, the application will restart

C. Replace the NLB with an Application Load Balancer. Enable HTTP health checks by supplying the URL of the company's application. Configure an Auto Scaling action to replace unhealthy instances

D. Create an Amazon CloudWatch alarm that monitors the UnhealthyHostCount metric for the NLB. Configure an Auto Scaling action to replace unhealthy instances when the alarm is in the ALARM state

(单选题) Q185. A company uses a three-tier web application to provide training to new employees. The application is accessed for only 12 hours every day. The company is using an Amazon RDS for MySQL DB instance to store information and wants to minimize costs. What should a solutions architect do to meet these requirements?

A. Configure an IAM policy for AWS Systems Manager Session Manager. Create an IAM role for the policy. Update the trust relationship of the role. Set up automatic start and stop for the DB instance.

B. Create an Amazon ElastiCache for Redis cache cluster that gives users the ability to access the data from the cache when the DB instance is stopped. Invalidate the cache after the DB instance is started.

C. Launch an Amazon EC2 instance. Create an IAM role that grants access to Amazon RDS. Attach the role to the EC2 instance. Configure a cron job to start and stop the EC2 instance on the desired schedule.

D. Create AWS Lambda functions to start and stop the DB instance. Create Amazon EventBridge (Amazon CloudWatch Events) scheduled rules to invoke the Lambda functions. Configure the Lambda functions as event targets for the rules.

(单选题) Q186. A company needs to save the results from a medical trial to an Amazon S3 repository. The repository must allow a few scientists to add new files and must restrict all other users to read-only access. No users can have the ability to modify or delete any files in the repository. The company must keep every file in the repository for a minimum of 1 year after its creation date. Which solution will meet these requirements?

A. Use S3 Object Lock in governance mode with a legal hold of 1 year.

B. Use S3 Object Lock in compliance mode with a retention period of 365 days.

C.Use an IAM role to restrict all users from deleting or changing objects in the S3 bucket.Use an S3 bucket policy to only allow the IAM role.

D.Configure the S3 bucket to invoke an AWS Lambda function every time an object is added.Configure the function to track the hash of the saved object so that modified objects can be marked accordingly.

(单选题) Q187.A company is migrating an application from on-premises servers to Amazon EC2 instances.As part of the migration design requirements, a solutions architect must implement infrastructure metric alarms.The company does not need to take action if CPU utilization increases to more than 50% for a short burst of time.However, if the CPU utilization increases to more than 50% and read IOPS on the disk are high at the same time, the company needs to act as soon as possible.The solutions architect also must reduce false alarms.What should the solutions architect do to meet these requirements ?

A.Create Amazon CloudWatch composite alarms where possible.

B.Create Amazon CloudWatch dashboards to visualize the metrics and react to issues quickly.

C.Create Amazon CloudWatch Synthetics canaries to monitor the application and raise an alarm.

D.Create single Amazon CloudWatch metric alarms with multiple metric thresholds where possible.

(单选题) Q188.A solutions architect needs to securely store a database user name and password that an application uses to access an Amazon RDS DB instance.The application that accesses the database runs on an Amazon EC2 instance.The solutions architect wants to create a secure parameter in AWS Systems Manager Parameter Store.What should the solutions architect do to meet this requirement ?

A.Create an IAM role that has read access to the Parameter Store parameter.Allow Decrypt access to an AWS Key Management Service (AWS KMS) key that is used to encrypt the parameter.Assign this IAM role to the EC2 instance

B.Create an IAM policy that allows read access to the Parameter Store parameter.Allow Decrypt access to an AWS Key Management Service (AWS KMS)

key that is used to encrypt the parameter. Assign this IAM policy to the EC2 instance.
C. Create an IAM trust relationship between the Parameter Store parameter and the EC2 instance. Specify Amazon RDS as a principal in the trust policy.
D. Create an IAM trust relationship between the DB instance and the EC2 instance. Specify Systems Manager as a principal in the trust policy

(单选题) Q189. A company hosts a website analytics application on a single Amazon EC2 On-Demand Instance. The analytics software is written in PHP and uses a MySQL database. The analytics software, the web server that provides PHP, and the database server are all hosted on the EC2 instance. The application is showing signs of performance degradation during busy times and is presenting 5xx errors. The company needs to make the application scale seamlessly. Which solution will meet these requirements MOST cost-effectively?

A. Migrate the database to an Amazon RDS for MySQL DB instance. Create an AMI of the web application. Use the AMI to launch a second EC2 On-Demand Instance. Use an Application Load Balancer to distribute the load to each EC2 instance.

B. Migrate the database to an Amazon RDS for MySQL DB instance. Create an AMI of the web application. Use the AMI to launch a second EC2 On-Demand Instance. Use Amazon Route 53 weighted routing to distribute the load across the two EC2 instances.

C. Migrate the database to an Amazon Aurora MySQL DB instance. Create an AWS Lambda function to stop the EC2 instance and change the instance type. Create an Amazon CloudWatch alarm to invoke the Lambda function when CPU utilization surpasses 75%.

D. Migrate the database to an Amazon Aurora MySQL DB instance. Create an AMI of the web application. Apply the AMI to a launch template. Create an Auto Scaling group with the launch template. Configure the launch template to use a Spot Fleet. Attach an Application Load Balancer to the Auto Scaling group.

(单选题) Q190. An ecommerce company hosts its analytics application in the AWS Cloud. The application generates about 300 MB of data each month. The data is stored in JSON format. The company is evaluating a disaster recovery solution to back up

the data.The data must be accessible in milliseconds if it is needed, and the data must be kept for 30 days.Which solution meets these requirements MOST cost-effectively ?

- A.Amazon OpenSearch Service (Amazon Elasticsearch Service)
- B.Amazon S3 Glacier
- C.Amazon S3 Standard
- D.Amazon RDS for PostgreSQL

(单选题) Q191.A company runs workloads on AWS.The company needs to connect to a service from an external provider.The service is hosted in the provider's VPC.According to the company's security team, the connectivity must be private and must be restricted to the target service.The connection must be initiated only from the company's VPC.Which solution will meet these requirements ?

- A.Create a VPC peering connection between the company's VPC and the provider's VPC.Update the route table to connect to the target service.
- B.Ask the provider to create a virtual private gateway in its VPC.Use AWS PrivateLink to connect to the target service.
- C.Create a NAT gateway in a public subnet of the company's VPC.Update the route table to connect to the target service.
- D.Ask the provider to create a VPC endpoint for the target service.Use AWS PrivateLink to connect to the target service.

(单选题) Q192.A company has created an image analysis application in which users can upload photos and add photo frames to their images.The users upload images and metadata to indicate which photo frames they want to add to their images.The application uses a single Amazon EC2instance and Amazon DynamoDB to store the metadata.The application is becoming more popular, and the number of users is increasing.The company expects the number of concurrent users to vary significantly depending on the time of day and day of week.The company must ensure that the application can scale to meet the needs of the growing user base.Which solution meets these requirements ?

- A.Use AWS Lambda to process the photos.Store the photos and metadata in

DynamoDB.

B.Use Amazon Kinesis Data Firehose to process the photos and to store the photos and metadata

C.Use AWS Lambda to process the photos.Store the photos in Amazon S3.Retain DynamoDB to store the metadata

D.increase the number of EC2 instances to three.Use Provisioned IOPS SSD (io2) Amazon Elastic Block Store (Amazon EBS) volumes to store the photos and metadata

(单选题) Q193.A company is developing a two-tier web application on AWS.The company's developers have deployed the application on an Amazon EC2 instance that connects directly to a backend Amazon RDS database.The company must not hardcode database credentials in the application.The company must also implement a solution to automatically rotate the database credentials on a regular basis.Which solution will meet these requirements with the LEAST operational overhead ?

A.Store the database credentials in the instance metadata.Use Amazon EventBridge (Amazon CloudWatch Events) rules to run a scheduled AWS Lambda function that updates the RDS credentials and instance metadata at the same time.

B.Store the database credentials in a configuration file in an encrypted Amazon S3 bucket.Use Amazon EventBridge (Amazon CloudWatch Events) rules to run a scheduled AWS Lambda function that updates the RDS credentials and the credentials in the configuration file at the same time.Use S3 Versioning to ensure the ability to fall back to previous values.

C.Store the database credentials as a secret in AWS Secrets Manager.Turn on automatic rotation for the secret.Attach the required permission to the EC2 role to grant access to the secret.

D.Store the database credentials as encrypted parameters in AWS Systems Manager Parameter Store.Turn on automatic rotation for the encrypted parameters.Attach the required permission to the EC2 role to grant access to the encrypted parameters.

(单选题) Q194.A company is planning to use an Amazon DynamoDB table for data storage.The company is concerned about cost optimization.The table will not be used

on most mornings. In the evenings, the read and write traffic will often be unpredictable. When traffic spikes occur, they will happen very quickly. What should a solutions architect recommend?

- A. Create a DynamoDB table in on-demand capacity mode.
- B. Create a DynamoDB table with a global secondary index.
- C. Create a DynamoDB table with provisioned capacity and auto scaling.
- D. Create a DynamoDB table in provisioned capacity mode, and configure it as a global table.

(单选题) Q195. A company runs an application on Amazon EC2 instances. The company needs to implement a disaster recovery (DR) solution for the application. The DR solution needs to have a recovery time objective (RTO) of less than 4 hours. The DR solution also needs to use the fewest possible AWS resources during normal operations. Which solution will meet these requirements in the MOST operationally efficient way?

- A. Create Amazon Machine Images (AMIs) to back up the EC2 instances. Copy the AMIs to a secondary AWS Region. Automate infrastructure deployment in the secondary Region by using AWS Lambda and custom scripts.
- B. Create Amazon Machine Images (AMIs) to back up the EC2 instances. Copy the AMIs to a secondary AWS Region. Automate infrastructure deployment in the secondary Region by using AWS CloudFormation.
- C. Launch EC2 instances in a secondary AWS Region. Keep the EC2 instances in the secondary Region active at all times.
- D. Launch EC2 instances in a secondary Availability Zone. Keep the EC2 instances in the secondary Availability Zone active at all times.

(单选题) Q196. A telemarketing company is designing its customer call center functionality on AWS. The company needs a solution that provides multiple speaker recognition and generates transcript files. The company wants to query the transcript files to analyze the business patterns. The transcript files must be stored for 7 years for auditing purposes. Which solution will meet these requirements?

- A. Use Amazon Rekognition for multiple speaker recognition. Store the transcript files

in Amazon S3. Use machine learning models for transcript file analysis.

B. Use Amazon Transcribe for multiple speaker recognition. Use Amazon Athena for transcript file analysis.

C. Use Amazon Translate for multiple speaker recognition. Store the transcript files in Amazon Redshift. Use SQL queries for transcript file analysis.

D. Use Amazon Rekognition for multiple speaker recognition. Store the transcript files in Amazon S3. Use Amazon Textract for transcript file analysis.

(单选题) Q197. A company wants to manage Amazon Machine Images (AMIs). The company currently copies AMIs to the same AWS Region where the AMIs were created. The company needs to design an application that captures AWS API calls and sends alerts whenever the Amazon EC2 CreateImage API operation is called within the company's account. Which solution will meet these requirements with the LEAST operational overhead?

A. Create an AWS Lambda function to query AWS CloudTrail logs and to send an alert when a CreateImage API call is detected.

B. Configure AWS CloudTrail with an Amazon Simple Notification Service (Amazon SNS) notification that occurs when updated logs are sent to Amazon S3. Use Amazon Athena to create a new table and to query on CreateImage when an API call is detected.

C. Create an Amazon EventBridge (Amazon CloudWatch Events) rule for the CreateImage API call. Configure the target as an Amazon Simple Notification Service (Amazon SNS) topic to send an alert when a CreateImage API call is detected. D. Configure an Amazon Simple Queue Service (Amazon SQS) FIFO queue as a target for AWS CloudTrail logs. Create an AWS Lambda function to send an alert to an Amazon Simple Notification Service (Amazon SNS) topic when a CreateImage API call is detected.

(多选题) Q198. A hospital wants to create digital copies for its large collection of historical written records. The hospital will continue to add hundreds of new documents each day. The hospital's data team will scan the documents and will upload the documents to the AWS Cloud. A solutions architect must implement a

solution to analyze the documents, extract the medical information, and store the documents so that an application can run SQL queries on the data. The solution must maximize scalability and operational efficiency. Which combination of steps should the solutions architect take to meet these requirements? (Select TWO.)

A. Write the document information to an Amazon EC2 instance that runs a MySQL database

B. Write the document information to an Amazon S3 bucket. Use Amazon Athena to query the data

C. Create an Auto Scaling group of Amazon EC2 instances to run a custom application that processes the scanned files and extracts the medical information

D. Create an AWS Lambda function that runs when new documents are uploaded. Use Amazon Rekognition to convert the documents to raw text. Use Amazon Transcribe Medical to detect and extract relevant medical information from the text

E. Create an AWS Lambda function that runs when new documents are uploaded. Use Amazon Textract to convert the documents to raw text. Use Amazon Comprehend Medical to detect and extract relevant medical information from the text

(单选题) Q199. A company serves a dynamic website from a fleet of Amazon EC2 instances behind an Application Load Balancer (ALB). The website needs to support multiple languages to serve customers around the world.

The website's architecture is running in the us-west-1 Region and is exhibiting high request latency for users that are located in other parts of the world. The website needs to serve requests quickly and efficiently regardless of a user's location. However, the company does not want to recreate the existing architecture across multiple Regions. What should a solutions architect do to meet these requirements?

A. Replace the existing architecture with a website that is served from an Amazon S3 bucket. Configure an Amazon CloudFront distribution with the S3 bucket as the origin. Set the cache behavior settings to cache based on the Accept-Language request header.

B. Configure an Amazon CloudFront distribution with the ALB as the origin. Set the cache behavior settings to cache based on the Accept-Language request header

C. Create an Amazon API Gateway API that is integrated with the ALB. Configure the

API to use the HTTP integration type. Setup an API Gateway stage to enable the API cache based on the Accept-Language request header

D. Launch an EC2 instance in each additional Region and configure NGINX to act as a cache server for that Region. Put all the EC2 instances and the ALB behind an Amazon Route 53 record set with a geolocation routing policy.

(单选题) Q200. A company runs a shopping application that uses Amazon DynamoDB to store customer information. In case of data corruption, a solutions architect needs to design a solution that meets a recovery point objective (RPO) of 15 minutes and a recovery time objective (RTO) of 1 hour. What should the solutions architect recommend to meet these requirements?

A. Configure DynamoDB global tables. For RPO recovery, point the application to a different AWS Region.

B. Configure DynamoDB point-in-time recovery. For RPO recovery, restore to the desired point in time.

C. Export the DynamoDB data to Amazon S3 Glacier on a daily basis. For RPO recovery, import the data from S3 Glacier to DynamoDB.

D. Schedule Amazon Elastic Block Store (Amazon EBS) snapshots for the DynamoDB table every 15 minutes. For RPO recovery, restore the DynamoDB table by using the EBS snapshot.

(单选题) Q201. A company has implemented a self-managed DNS solution on three Amazon EC2 instances behind a Network Load Balancer (NLB) in the us-west-2 Region. Most of the company's users are located in the United States and Europe. The company wants to improve the performance and availability of the solution. The company launches and configures three EC2 instances in the eu-west-1 Region and adds the EC2 instances as targets for a new NLB. Which solution can the company use to route traffic to all the EC2 instances?

A. Create an Amazon Route 53 geolocation routing policy to route requests to one of the two NLBs. Create an Amazon CloudFront distribution. Use the Route 53 record as the distribution's origin.

B. Create a standard accelerator in AWS Global Accelerator. Create endpoint groups in

us-west-2 and eu-west-1. Add the two NLBs as endpoints for the endpoint groups.
C. Attach Elastic IP addresses to the six EC2 instances. Create an Amazon Route 53 geolocation routing policy to route requests to one of the six EC2 instances. Create an Amazon CloudFront distribution. Use the Route 53 record as the distribution's origin.

D. Replace the two NLBs with two Application Load Balancers (ALBs). Create an Amazon Route 53 latency routing policy to route requests to one of the two ALBs. Create an Amazon CloudFront distribution. Use the Route 53 record as the distribution's origin.

(单选题) Q202. A company is building an ecommerce web application on AWS. The application sends information about new orders to an Amazon API Gateway REST API to process. The company wants to ensure that orders are processed in the order that they are received. Which solution will meet these requirements?

A. Use an API Gateway integration to publish a message to an Amazon Simple Notification Service (Amazon SNS) topic when the application receives an order. Subscribe an AWS Lambda function to the topic to perform processing.

B. Use an API Gateway integration to send a message to an Amazon Simple Queue Service (Amazon SQS) FIFO queue when the application receives an order. Configure the SQS FIFO queue to invoke an AWS Lambda function for processing.

C. Use an API Gateway authorizer to block any requests while the application processes an order.

D. Use an API Gateway integration to send a message to an Amazon Simple Queue Service (Amazon SQS) standard queue when the application receives an order. Configure the SQS standard queue to invoke an AWS Lambda function for processing.

(单选题) Q203. A reporting team receives files each day in an Amazon S3 bucket. The reporting team manually reviews and copies the files from this initial S3 bucket to an analysis S3 bucket each day at the same time to use with Amazon QuickSight. Additional teams are starting to send more files in larger sizes to the initial

S3 bucket. The reporting team wants to move the files automatically to the analysis S3 bucket as the files enter the initial S3 bucket. The reporting team also wants to use AWS Lambda functions to run pattern-matching code on the copied data. In addition, the reporting team wants to send the data files to a pipeline in Amazon SageMaker Pipelines. What should a solutions architect do to meet these requirements with the LEAST operational overhead?

A. Create a Lambda function to copy the files to the analysis S3 bucket. Create an S3 event notification for the analysis S3 bucket. Configure Lambda and SageMaker Pipelines as destinations of the event notification. Configure s3:ObjectCreated:Put as the event type.

B. Create a Lambda function to copy the files to the analysis S3 bucket. Configure the analysis S3 bucket to send event notifications to Amazon EventBridge (Amazon CloudWatch Events). Configure an ObjectCreated rule in EventBridge (CloudWatch Events). Configure Lambda and SageMaker Pipelines as targets for the rule.

C. Configure S3 replication between the S3 buckets. Create an S3 event notification for the analysis S3 bucket. Configure Lambda and SageMaker Pipelines as destinations of the event notification. Configure s3:ObjectCreated:Put as the event type.

D. Configure S3 replication between the S3 buckets. Configure the analysis S3 bucket to send event notifications to Amazon EventBridge (Amazon CloudWatch Events). Configure an ObjectCreated rule in EventBridge (CloudWatch Events). Configure Lambda and SageMaker Pipelines as targets for the rule.

(单选题) Q204. A company needs to keep user transaction data in an Amazon DynamoDB table. The company must retain the data for 7 years. What is the MOST operationally efficient solution that meets these requirements?

A. Use DynamoDB point-in-time recovery to back up the table continuously.

B. Use AWS Backup to create backup schedules and retention policies for the table.

C. Create an on-demand backup of the table by using the DynamoDB console. Store the backup in an Amazon S3 bucket. Set an S3 Lifecycle configuration for the S3 bucket.

D. Create an Amazon EventBridge (Amazon CloudWatch Events) rule to invoke an AWS Lambda function. Configure the Lambda function to back up the table and to store the backup in an Amazon S3 bucket. Set an S3 Lifecycle configuration for the

S3 bucket

(单选题) Q205.A company wants to move its application to a serverless solution.The serverless solution needs to analyze existing and new data by using SQL.The company stores the data in an Amazon S3 bucket The data requires encryption and must be replicated to a different AWS Region.Which solution will meet these requirements with the LEAST operational overhead ?

A.Create a new S3 bucket.Load the data into the new S3 bucket.Use S3 Cross-Region Replication (CRR) to replicate encrypted objects to an S3 bucket in another Region.Use server-side encryption with AwS KMS multi-Region keys (SSE-KMS) . Use Amazon Athena to query the data.

B.Create a new S3 bucket Load the data in to the new S3 bucket Use S3 Cross-Region Replication (CRR) to replicate encrypted objects to an S3 bucket in another Region.Use server-side encryption with AwS KMS multi-Region keys (SSE-KMS) Use Amazon RDS to query the data.

C.Load the data into the existing S3 bucket.Use S3 Cross-Region Replication (CRR) to replicate encrypted objects to an S3 bucket in another Region.Use server-side encryption with Amazon S3 managed encryption keys (SSE-s3) .Use Amazon Athena to query the data.

D.Load the data into the existing S3 bucket Use S3 Cross-Region Replication (CRR) to replicate encrypted objects to an S3 bucket in another Region.Use server-side encryption with Amazon S3 managed encryption keys (SSE-S3) .Use Amazon RDS to query the data.

(单选题) Q206.A company is planning to move its data to an Amazon S3 bucket.The data must be encrypted when it is stored in the S3 bucket.Additionally, the encryption key must be automatically rotated every year. Which solution will meet these requirements with the LEAST operational overhead ?

A.Move the data to the S3 bucket.Use server-side encryption with Amazon S3 managed encryption keys (SSE-S3) .Use the buit-in key rotation behavior of SSE-S3 encryption keys.

B.Create an AWS Key Management Service (AWS KMS) customer managed key

Enable automatic key rotation. Set the S3 bucket's default encryption behavior to use the customer managed KMS key. Move the data to the S3 bucket.

C. Create an AWS Key Management Service (AWS KMS) customer managed key. Set the S3 bucket's default encryption behavior to use the customer managed KMS key. Move the data to the S3 bucket. Manually rotate the KMS key every year.

D. Encrypt the data with customer key material before moving the data to the S3 bucket. Create an AWS Key Management Service (AWS KMS) key without key material. Import the customer key material into the KMS key. Enable automatic key rotation.

(单选题) Q207. A company's application integrates with multiple software-as-a-service (SaaS) sources for data collection. The company runs Amazon EC2 instances to receive the data and to upload the data to an Amazon S3 bucket for analysis. The same EC2 instance that receives and uploads the data also sends a notification to the user when an upload is complete. The company has noticed slow application performance and wants to improve the performance as much as possible. Which solution will meet these requirements with the LEAST operational overhead?

A. Create an Auto Scaling group so that EC2 instances can scale out. Configure an S3 event notification to send events to an Amazon Simple Notification Service (Amazon SNS) topic when the upload to the S3 bucket is complete.

B. Create an Amazon AppFlow flow to transfer data between each SaaS source and the S3 bucket. Configure an S3 event notification to send events to an Amazon Simple Notification Service (Amazon SNS) topic when the upload to the S3 bucket is complete.

C. Create an Amazon EventBridge (Amazon CloudWatch Events) rule for each SaaS source to send output data. Configure the S3 bucket as the rule's target. Create a second EventBridge (CloudWatch Events) rule to send events when the upload to the S3 bucket is complete. Configure an Amazon Simple Notification Service (Amazon SNS) topic as the second rule's target.

D. Create a Docker container to use instead of an EC2 instance. Host the containerized application on Amazon Elastic Container Service (Amazon ECS). Configure Amazon CloudWatch Container Insights to send events to an Amazon Simple Notification

Service (Amazon SNS) topic when the upload to the S3 bucket is complete

(单选题) Q208. A company is running a popular social media website. The website gives users the ability to upload images to share with other users. The company wants to make sure that the images do not contain inappropriate content. The company needs a solution that minimizes development effort. What should a solutions architect do to meet these requirements ?

A. Use Amazon Comprehend to detect inappropriate content. Use human review for low-confidence predictions.

B. Use Amazon Rekognition to detect inappropriate content. Use human review for low-confidence predictions

C. Use Amazon SageMaker to detect inappropriate content. Use ground truth to label low-confidence predictions.

D. Use AWS Fargate to deploy a custom machine learning model to detect inappropriate content. Use ground truth to label low-confidence predictions.

(单选题) Q209. A company needs a storage solution for an application that runs on a high performance computing (HPC) cluster. The cluster is hosted on AWS Fargate for Amazon Elastic Container Service (Amazon ECS). The company needs a mountable file system that provides concurrent access to files while delivering hundreds of GBps of throughput at sub-millisecond latencies.

Which solution meets these requirements?

A. Create an Amazon FSx for Lustre file share for the application data. Create an IAM role that allows Fargate to access the FSx for Lustre file share

B. Create an Amazon Elastic File System (Amazon EFS) file share for the application data. Create an IAM role that allows Fargate to access the EFS file share

C. Create an Amazon S3 bucket for the application data. Create an S3 bucket policy that allows Fargate to access the S3 bucket

D. Create an Amazon Elastic Block Store (Amazon EBS) Provisioned IOPS SSD (io2) volume for the application data. Create an IAM role that allows Fargate to access the volume.

(单选题) Q210.A company's infrastructure consists of Amazon EC2 instances and an Amazon RDS DB instance in a single AWS Region.The company wants to back up its data in a separate Region.Which solution wil meet these requirements with the LEAST operational overhead ?

- A.Use AWS Backup to copy EC2 backups and RDS backups to the separate Region.
- B.Use Amazon Data Lifecycle Manager (Amazon DLM) to copy EC2 backups and RDS backups to the separate Region.
- C.Create Amazon Machine Images (AMIs) of the EC2 instances.Copy the AMIs to the separate Region.Create a read replica for the RDS DB instance in the separate Region.
- D.Create Amazon Elastic Block Store (Amazon EBS) snapshots.Copy the EBS snapshots to the separate Region.Create RDS snapshots.Export the RDS snapshots to Amazon S3.Configure S3 Cross-Region Replication (CRR) to the separate Region.

(单选题) Q211.An ecommerce company wants to launch a one-deal-a-day website on AWS.Each day will feature exactly one product on sale for a period of 24 hours.The company wants to be able to handle millions of requests each hour with millisecond latency during peak hours.Which solution will meet these requirements with the LEAST operational overhead ?

- A.Use Amazon S3 to host the full website in different S3 buckets.Add Amazon CloudFront distributions.Set the S3 buckets as origins for the distributions.Store the order data in Amazon S3.
- B.Deploy the full website on Amazon EC2instances that run in Auto Scaling groups across multiple Availabiity Zones.Add an Application Load Balancer (ALB) to distribute the website traffic.Add another ALB for the backend APIs.Store the data in Amazon RDS for MySQL.
- C.Migrate the full application to run in containers.Host the containers on Amazon Elastic Kubernetes Service (Amazon EKS) .Use the Kubernetes Cluster Autoscaler to increase and decrease the number of pods to process bursts in traffic.Store the data in Amazon RDS for MySQL.
- D.Use an Amazon S3 bucket to host the website's static content.Deploy an Amazon CloudFront distribution.Set the S3 bucket as the origin.Use Amazon API Gateway and

AWS Lambda functions for the backend APIs.Store the data in Amazon DynamoDB.

(多选题) Q212.A company wants to migrate its on-premises data center to AWS.According to the company's compliance requirements, the company can use only the ap-northeast-3 Region.Company administrators are not permitted to connect VPCs to the internetWhich solutions will meet these requirements ? (Select Two)

A.Use AWS Control Tower to implement data residency guardrails to deny internet access and deny access to all AWS Regions except ap-northeast-3

B.Use rules in AWS WAF to prevent internet access.Deny access to all AWS Regions exceptap-northeast-3 in the AWS account settings.

C.Use AWS Organizations to configure service control policies (SCPs) that prevent VPCs from gaining internet access.Deny access to all AWS Regions except ap-northeast-3.

D.Create an outbound rule for the network ACL in each VPC to deny all traffic from 0.0.0.0/0.Create an IAM policy for each user to prevent the use of any AWS Region other than ap-northeast-3

E.Use AWS Config to activate managed rules to detect and alert for internet gateways and to detect and alert for new resources deployed outside of ap-northeast-3.

(单选题) Q213.A company runs an Oracle database on premises.As part of the company's migration to AWS, the company wants to upgrade the database to the most recent available version.The company also wants to set up disaster recovery (DR) for the database.The company needs to minimize the operational overhead for normal operations and DR setup.The company also needs to maintain access to the database's underlying operating system.Which solution will meet these requirements ?

A.Migrate the Oracle database to an Amazon EC2 instance.Set up database replication to a different AWS Region.

B.Migrate the Oracle database to Amazon RDS for Oracle.Activate cross-Region automated backups to replicate the snapshots to another AWS Region

C.Migrate the Oracle database to Amazon RDS Custom for Oracle.Create a read

replica for the database in another AWS Region.

D.Migrate the Oracle database to Amazon RDS for Oracle.Create a standby database in another Availability Zone.

(单选题) Q214.A company hosts its application on AWS.The company uses Amazon Cognito to manage users.When users log in to the application, the application fetches required data from Amazon DynamoDB by using a REST API that is hosted in Amazon API Gateway.The company wants an AWS managed solution that will control access to the REST API to reduce development efforts.Which solution will meet these requirements with the LEAST operational overhead ?

A.Configure an AWS Lambda function to be an authorizer in API Gateway to validate which user made the request.

B.For each user, create and assign an API key that must be sent with each request.Validate the key by using an AWS Lambda function.

C.Send the user's email address in the header with every request.Invoke an AWS Lambda function to validate that the user with that email address has proper access.

D.Configure an Amazon Cognito user pool authorizer in API Gateway to allow Amazon Cognito to validate each request.

(单选题) Q215.A hospital recently deployed a RESTful API with Amazon API Gateway and AWS Lambda.The hospital uses API Gateway and Lambda to upload reports that are in PDF format and JPEG format.The hospital needs to modify the Lambda code to identify protected health information (PHI) in the reports.Which solution will meet these requirements with the LEAST operational overhead ?

A.Use existing Python libraries to extract the text from the reports and to identify the PHI from the extracted text.

B.Use Amazon Textract to extract the text from the reports.Use Amazon SageMaker to identify the PHI from the extracted text

C.Use Amazon Textract to extract the text from the reports.Use Amazon Comprehend Medical to identify the PHI from the extracted text.

D.Use Amazon Rekognition to extract the text from the reports.Use Amazon Comprehend Medical to identify the PHI from the extracted text

(单选题) Q216.A company needs to configure a real-time data ingestion architecture for its application.The company needs an API, a process that transforms data as the data is streamed, and a storage solution for the data.

Which solution will meet these requirements with the LEAST operational overhead ?

A.Deploy an Amazon EC2 instance to host an API that sends data to an Amazon Kinesis data stream.Create an Amazon Kinesis Data Firehose delivery stream that uses the Kinesis data stream as a data source.Use AWS Lambda functions to transform the data.Use the Kinesis Data Firehose delivery stream to send the data to Amazon S3

B.Deploy an Amazon EC2 instance to host an API that sends data to AWS Glue.Stop source/destination checking on the EC2 instance.Use AWS Glue to transform the data and to send the data to Amazon S3

C.Configure an Amazon API Gateway API to send data to an Amazon Kinesis data stream.Create an Amazon Kinesis Data Firehose delivery stream that uses the Kinesis data stream as a data source.Use AWS Lambda functions to transform the data.Use the Kinesis Data Firehose delivery stream to send the data to Amazon S3.

D.Configure an Amazon API Gateway API to send data to AWS Glue.Use AWS Lambda functions to transform the data.Use AWS Glue to send the data to Amazon S3.

(单选题) Q217.A company is launching a new application and will display application metrics on an Amazon CloudWatch dashboard.The company's product manager needs to access this dashboard periodically.The product manager does not have an AWS account.A solutions architect must provide access to the product manager by following the principle of least privilege.Which solution will meet these requirements ?

A.Share the dashboard from the CloudWatch console.Enter the product manager's email address, and complete the sharing steps.Provide a shareable link for the dashboard to the product manager.

B.Create an IAM user specifically for the product manager.Attach the CloudWatchReadOnlyAccess AWS managed policy to the user.Share the new login credentials with the product manager.Share the browser URL of the correct

dashboard with the product manager.

C.Create an IAM user for the company's employees Attach the ViewOnlyAccess AWS managed policy to the IAM user.Share the new login credentials with the product manager.Ask the product manager to navigate to the CloudWatch console and locate the dashboard by name in the Dashboards section

D.Deploy a bastion server in a public subnet.When the product manager requires access to the dashboard, start the server and share the RDP credentials.On the bastion server, ensure that the browser is configured to open the dashboard URL with cached AWS credentials that have appropriate permissions to view the dashboard.

(多选题) Q218.A company has a data ingestion workflow that consists of the following:"An Amazon Simple Notification Service (Amazon SNS) topic for notifications about new data deliveries"An AWS Lambda function to process the data and record metadata The company observes that the ingestion workflow fails occasionally because of network connectivity issues.When such a failure occurs, the Lambda function does not ingest the corresponding data unless the company manually reruns the job.Which combination of actions should a solutions architect take to ensure that the Lambda function ingests all data in the future ? (Select TWO.) A.Deploy the Lambda function in multiple Availability Zones.

B.Create an Amazon Simple Queue Service (Amazon SQS) queue, and subscribe it to the SNS topic.

C.Increase the CPU and memory that are allocated to the Lambda function.

D.Increase provisioned throughput for the Lambda function.

E.Modify the Lambda function to read from an Amazon Simple Queue Service (Amazon SQS) queue.

(单选题) Q219.A company wants to improve its ability to clone large amounts of production data into a test environment in the same AWS Region.The data is stored in Amazon EC2 instances on Amazon Elastic Block Store

(Amazon EBS) volumes.Modifications to the cloned data must not affect the production environment.The software that accesses this data requires consistently

high IO performance.A solutions architect needs to minimize the time that is required to clone the production data into the test environment.Which solution will meet these requirements ?

A.Take EBS snapshots of the production EBS volumes.Restore the snapshots onto EC2instance store volumes in the test environment

B.Configure the production EBS volumes to use the EBS Multi-Attach feature.Take EBS snapshots of the production EBS volumes.Attach the production EBS volumes to the EC2 instances in the test environment.

C.Take EBS snapshots of the production EBS volumes.Create and initialize new EBS volumes.Attach the new EBS volumes to EC2 instances in the test environment before restoring the volumes from the production EBS snapshots.

D.Take EBS snapshots of the production EBS volumes.Turn on the EBS fast snapshot restore feature on the EBS snapshots.Restore the snapshots into new EBS volumes.Attach the new EBS volumes to EC2instances in the test environment

(单选题) Q220.A company runs a highly available image-processing application on Amazon EC2 instances in a single VPC.The EC2instances run inside several subnets across multiple Availability Zones.The EC2 instances do not communicate with each other.However, the EC2 instances download images from Amazon S3 and upload images to Amazon S3 through a single NAT gateway.The company is concerned about data transfer charges.What is the MOST cost-effective way for the company to avoid Regional data transfer charges ?

A.Launch the NAT gateway in each Availability Zone.

B.Replace the NAT gateway with a NAT instance.

C.Deploy a gateway VPC endpoint for Amazon S3.

D.Provision an EC2 Dedicated Host to run the EC2 instances.

(单选题) Q221.A company runs an ecommerce application on Amazon EC2instances behind an Application Load Balancer.The instances run in an Amazon EC2Auto Scaling group across multiple Availabiity Zones.The Auto Scaling group scales based on CPU utiization metrics.The ecommerce application stores the transaction data in a MySQL 8.0 database that is hosted on a large EC2 instance.The

database's performance degrades quickly as application load increases. The application handles more read requests than write transactions. The company wants a solution that will automatically scale the database to meet the demand of unpredictable read workloads while maintaining high availability. Which solution will meet these requirements ?

- A. Use Amazon Redshift with a single node for leader and compute functionality.
- B. Use Amazon RDS with a Single-AZ deployment. Configure Amazon RDS to add reader instances in a different Availability Zone.
- C. Use Amazon Aurora with a Multi-AZ deployment. Configure Aurora Auto Scaling with Aurora Replicas.
- D. Use Amazon ElastiCache for Memcached with EC2 Spot Instances.

(单选题) Q222. A company has more than 5 TB of file data on Windows file servers that run on premises. Users and applications interact with the data each day. The company is moving its Windows workloads to AWS. As the company continues this process, the company requires access to AWS and on-premises file storage with minimum latency. The company needs a solution that minimizes operational overhead and requires no significant changes to the existing file access patterns. The company uses an AWS Site-to-Site VPN connection for connectivity to AWS. What should a solutions architect do to meet these requirements ?

- A. Deploy and configure Amazon FSx for Windows File Server on AWS. Move the on-premises file data to FSx for Windows File Server. Reconfigure the workloads to use FSx for Windows File Server on AWS.
- B. Deploy and configure an Amazon S3 File Gateway on premises. Move the on-premises file data to the S3 File Gateway. Reconfigure the on-premises workloads and the cloud workloads to use the S3 File Gateway.
- C. Deploy and configure an Amazon S3 File Gateway on premises. Move the on-premises file data to Amazon S3. Reconfigure the workloads to use either Amazon S3 directly or the S3 File Gateway, depending on each workload's location.
- D. Deploy and configure Amazon FSx for Windows File Server on AWS. Deploy and configure an Amazon FSx File Gateway on premises. Move the on-premises file data to the FSx File Gateway. Configure the cloud workloads to use FSx for Windows File

Server on AWS. Configure the on-premises workloads to use the FSx File Gateway.

(单选题) Q223. A company's containerized application runs on an Amazon EC2 instance. The application needs to download security certificates before it can communicate with other business applications. The company wants a highly secure solution to encrypt and decrypt the certificates in near real time. The solution also needs to store data in highly available storage after the data is encrypted. Which solution will meet these requirements with the LEAST operational overhead?

A. Create AWS Secrets Manager secrets for encrypted certificates. Manually update the certificates as needed. Control access to the data by using fine-grained IAM access

B. Create an AWS Lambda function that uses the Python cryptography library to receive and perform encryption operations. Store the function in an Amazon S3 bucket

C. Create an AWS Key Management Service (AWS KMS) customer managed key. Allow the EC2 role to use the KMS key for encryption operations. Store the encrypted data on Amazon S3

D. Create an AWS Key Management Service (AWS KMS) customer managed key. Allow the EC2 role to use the KMS key for encryption operations. Store the encrypted data on Amazon Elastic Block Store (Amazon EBS) volumes

(单选题) Q224. A company has registered its domain name with Amazon Route 53. The company uses Amazon API Gateway in the ca-central-1 Region as a public interface for its backend microservice APIs. Third-party services consume the APIs securely. The company wants to design its API Gateway URL with the company's domain name and corresponding certificate so that the third-party services can use HTTPS. Which solution will meet these requirements?

A. Create stage variables in API Gateway with Name="Endpoint-URL" and Value="Company Domain Name" to overwrite the default URL. Import the public certificate associated with the company's domain name into AWS Certificate Manager (ACM).

B. Create Route 53 DNS records with the company's domain name. Point the alias

record to the Regional API Gateway stage endpoint.Import the public certificate associated with the company's domain name into AWS Certificate Manager (ACM) in the us-east-1 Region.

C.Create a Regional API Gateway endpoint Associate the API Gateway endpoint with the company's domain name.Import the public certificate associated with the company's domain name into AWS Certificate Manager (ACM) in the same Region.Attach the certificate to the API Gateway endpoint.Configure Route 53to route traffic to the API Gateway endpoint.

D.Create a Regional API Gateway endpoint Associate the API Gateway endpoint with the company's domain name.Import the public certificate associated with the company's domain name into AWS Certificate Manager (ACM) in the us-east-1 Region.Attach the certificate to the API Gateway APIs.Create Route 53 DNS records with the company's domain name.Point an A record to the company's domain name.

(多选题) Q225.A company is developing an application that provides order shipping statistics for retrieval by a RESTAPI.The company wants to extract the shipping statistics, organize the data into an easy-to-read HTML format, and send the report to several email addresses at the same time every moning.Which combination of steps should a solutions architect take to meet these requirements? (Select TWO)

A.Configure the application to send the data to Amazon Kinesis Data Firehose.

B.Use Amazon Simple Email Service (Amazon SES) to format the data and to send the report by email.

C.Create an Amazon EventBridge (Amazon CloudWatch Events) scheduled event that invokes an AWS Glue job to query the application's API for the data.

D.Create an Amazon EventBridge (Amazon CloudWatch Events) scheduled event that invokes an AWS Lambda function to query the application's API for the data

E.Store the application data in Amazon S3 Create an Amazon Simple Notification Service (Amazon SNS) topic as an S3 event destination to send the report by email

(单选题) Q226.A company has a web application that is based on Java and

PHP.The company plans to move the application from on premises to AWS.The company needs the ability to test new site features frequently.

The company also needs a highly available and managed solution that requires minimum operational overhead.Which solution will meet these requirements ?

A.Create an Amazon S3 bucket.Enable static web hosting on the S3 bucket Upload the static content to the S3 bucket Use AWS Lambda to process all dynamic content.

B.Deploy the web application to an AWS Elastic Beanstalk environment.Use URL swapping to switch between multiple Elastic Beanstalk environments for feature testing

C.Deploy the web application to Amazon EC2 instances that are configured with Java and PHP.Use Auto Scaling groups and an Application Load Balancer to manage the website's availability.

D.Containerize the web application.Deploy the web application to Amazon EC2 instances.Use the AWS Load Balancer Controller to dynamically route traffic between containers that contain the new site features for testing

(单选题) Q227.An image-processing company has a web application that users use to upload images.The application uploads the images into an Amazon S3 bucket.The company has set up S3 event notifications to publish the object creation events to an Amazon Simple Queue Service (Amazon SQS) standard queue.The SQS queue serves as the event source for an AWS Lambda function that processes the images and sends the results to users through email.Users report that they are receiving multiple email messages for every uploaded image.A solutions architect determines that SQS messages are invoking the Lambda function more than once, resulting in multiple email messages.What should the solutions architect do to resolve this issue with the LEAST operational overhead ?

A.Set up long polling in the SQS queue by increasing the ReceiveMessage wait time to 30 seconds

B.Change the SQS standard queue to an SQS FIFO queue.Use the message deduplication ID to discard duplicate messages

C.Increase the visibility timeout in the SQS queue to a value that is greater than the total of the function timeout and the batch window timeout.

D.Modify the Lambda function to delete each message from the SQS queue immediately after the message is read before processing.

(单选题) Q228.A company hosts a containerized web application on a fleet of on-premises servers that process incoming requests.The number of requests is growing quickly.The on-premises servers cannot handle the increased number of requests.The company wants to move the application to AWS with minimum code changes and minimum development effort.Which solution will meet these requirements with the LEAST operational overhead ?

A.Use AWS Fargate on Amazon Elastic Container Service (Amazon ECS) to run the containerized web application with Service Auto Scaling.Use an Application Load Balancer to distribute the incoming requests.

B.Use two Amazon EC2 instances to host the containerized web application.Use an Application Load Balancer to distribute the incoming requests.

C.Use AWS Lambda with a new code that uses one of the supported languages.Create multiple Lambda functions to support the load.Use Amazon API Gateway as an entry point to the Lambda functions.

D.Use a high performance computing (HPC) solution such as AWS ParalelCluster to establish an HPC cluster that can process the incoming requests at the appropriate scale.

(单选题) Q229.A company wants to migrate its on-premises application to AWS.The application produces output files that vary in size from tens of gigabytes to hundreds of terabytes.The application data must be stored in a standard file system structure.The company wants a solution that scales automatically, is highly available, and requires minimum operational overhead.Which solution will meet these requirements ?

A.Migrate the application to run as containers on Amazon Elastic Container Service (Amazon ECS) .Use Amazon S3for storage.

B.Migrate the application to run as containers on Amazon Elastic Kubernetes Service (Amazon EKS) .Use Amazon Elastic Block Store (Amazon EBS) for storage.

C.Migrate the application to Amazon EC2 instances in a Multi-AZ Auto Scaling

group.Use Amazon Elastic File System (Amazon EFS) for storage.

D.Migrate the application to Amazon EC2 instances in a Multi-AZ Auto Scaling group.Use Amazon Elastic Block Store (Amazon EBS) for storage.

(单选题) Q230.A company is running an SMB file server in its data center.The file server stores large files that are accessed frequently for the first few days after the files are created.After 7 days, the files are rarely accessed.The total data size is increasing and is close to the company's total storage capacity.A solutions architect must increase the company's available storage space without losing low-latency access to the most recently accessed files.The solutions architect must also provide file lifecycle management to avoid future storage issues.Which solution will meet these requirements ?

A.Use AWS DataSync to copy data that is older than 7 days from the SMB file server to AWS

B.Create an Amazon S3 File Gateway to extend the company's storage space.Create an S3 Lifecycle policy to transition the data to S3 Glacier Deep Archive after 7 days

C.Create an Amazon FSx for Windows File Server file system to extend the company's storage space.

D.Install a utility on each user's computer to access Amazon S3 Create an S3 Lifecycle policy to transition the data to S3 Glacier Flexible Retrieval after 7 days.

(单选题) Q231.A company is building an application in the AWS Cloud.The application will store data in Amazon S3buckets in two AWS Regions.The company must use an AWS Key Management Service (AWS KMS) customer managed key to encrypt all data that is stored in the S3 buckets The data in both S3 buckets must be encrypted and decrypted with the same KMS key.The data and the key must be stored in each of the two Regions.Which solution will meet these requirements with the LEAST operational overhead ?

A.Create an S3 bucket in each Region Configure the S3 buckets to use server-side encryption with Amazon S3managed encryption keys (SSE-S3) .Configure replication between the S3 buckets

B.Create a customer managed multi-Region KMS key.Create an S3 bucket in each

Region Configure replication between the S3 buckets. Configure the application to use the KMS key with client-side encryption.

C. Create a customer managed KMS key and an S3 bucket in each Region. Configure the S3 buckets to use server-side encryption with Amazon S3 managed encryption keys (SSE-S3). Configure replication between the S3 buckets.

D. Create a customer managed KMS key and an S3 bucket in each Region. Configure the S3 buckets to use server-side encryption with AWS KMS keys (SSE-KMS). Configure replication between the S3 buckets.

(单选题) Q232. A company needs to run a critical application on AWS. The company needs to use Amazon EC2 for the application's database. The database must be highly available and must fail over automatically if a disruptive event occurs. Which solution will meet these requirements?

A. Launch two EC2 instances, each in a different Availability Zone in the same AWS Region. Install the database on both EC2 instances. Configure the EC2 instances as a cluster. Set up database replication.

B. Launch an EC2 instance in an Availability Zone. Install the database on the EC2 instance. Use an Amazon Machine Image (AMI) to back up the data. Use AWS CloudFormation to automate provisioning of the EC2 instance if a disruptive event occurs.

C. Launch two EC2 instances, each in a different AWS Region. Install the database on both EC2 instances. Set up database replication. Fail over the database to a second Region.

D. Launch an EC2 instance in an Availability Zone. Install the database on the EC2 instance. Use an Amazon Machine Image (AMI) to back up the data. Use EC2 automatic recovery to recover the instance if a disruptive event occurs.

(单选题) Q233. An application runs on an Amazon EC2 instance in a VPC. The application processes logs that are stored in an Amazon S3 bucket. The EC2 instance needs to access the S3 bucket without connectivity to the internet. Which solution will provide private network connectivity to Amazon S3?

A. Create a gateway VPC endpoint to the S3 bucket.

- B.Stream the logs to Amazon CloudWatch Logs.Export the logs to the S3 bucket.
- C.Create an instance profile on Amazon EC2 to allow S3 access.
- D.Create an Amazon API Gateway API with a private link to access the S3 endpoint.

(单选题) Q234.A company has an application that runs on Amazon EC2 instances and uses an Amazon Aurora database.The EC2 instances connect to the database by using user names and passwords that are stored locally in a file.The company wants to minimize the operational overhead of credential management.What should a solutions architect do to accomplish this goal ?

- A.Use AWS Secrets Manager.Turn on automatic rotation.
- B.Use AWS Systems Manager Parameter Store.Turn on automatic rotation.
- C.Create an Amazon S3 bucket to store objects that are encrypted with an AWS Key Management Service (AWS KMS) encryption key.Migrate the credential file to the S3 bucket.Point the application to the S3 bucket.
- D.Create an encrypted Amazon Elastic Block Store (Amazon EBS) volume for each EC2 instance.Attach the new EBS volume to each EC2 instance.Migrate the credential file to the new EBS volume.Point the application to the new EBS volume.

(单选题) Q235.A company is implementing a new business application.The application runs on two Amazon EC2 instances and uses an Amazon S3 bucket for document storage.A solutions architect needs to ensure that the EC2 instances can access the S3 bucket.What should the solutions architect do to meet this requirement ?

- A.Create an IAM role that grants access to the S3 bucket.Attach the role to the EC2 instances.
- B.Create an IAM policy that grants access to the S3 bucket.Attach the policy to the EC2 instances.
- C.Create an IAM group that grants access to the S3 bucket.Attach the group to the EC2 instances.
- D.Create an IAM user that grants access to the S3 bucket.Attach the user account to the EC2 instances.

(单选题) Q236.A company is designing an application.The application uses an AWS Lambda function to receive information through Amazon API Gateway and to store the information in an Amazon Aurora PostgreSQL database.During the proof-of-concept stage, the company has to increase the Lambda quotas significantly to handle the high volumes of data that the company needs to load into the database.A solutions architect must recommend a new design to improve scalability and minimize the configuration effort.Which solution will meet these requirements ?

A.Refactor the Lambda function code to Apache Tomcat code that runs on Amazon EC2 instances.Connect the database by using native Java Database Connectivity (JDBC) drivers.

B.Change the platform from Aurora to Amazon DynamoDB.Provision a DynamoDB Accelerator (DAX) cluster.Use the DAX client SDK to point the existing DynamoDB API calls at the DAX cluster.

C.Set up two Lambda functions.Configure one function to receive the information.Configure the other function to load the information into the database.Integrate the Lambda functions by using Amazon Simple Notification Service (Amazon SNS) .

D.Set up two Lambda functions.Configure one function to receive the information.Configure the other function to load the information into the database.Integrate the Lambda functions by using an Amazon Simple Queue Service (Amazon SQS) queue.

(单选题) Q237.A company needs to review its AWS Cloud deployment to ensure that its Amazon S3 buckets do not have unauthorized configuration changes.What should a solutions architect do to accomplish this goal ?

A.Turn on AWS Config with the appropriate rules.

B.Turn on AWS Trusted Advisor with the appropriate checks.

C.Turn on Amazon Inspector with the appropriate assessment template.

D.Turn on Amazon S3 server access logging.Configure Amazon EventBridge (Amazon Cloud Watch Events) .

(单选题) Q238.A company is migrating applications to AWS.The applications are

deployed in different accounts. The company manages the accounts centrally by using AWS Organizations. The company's security team needs a single sign-on (SSO) solution across all the company's accounts. The company must continue managing the users and groups in its on-premises self-managed Microsoft Active Directory. Which solution will meet these requirements ?

A. Enable AWS Single Sign-On (AWS SSO) from the AWS SSO console. Create a one-way forest trust or a one-way domain trust to connect the company's self-managed Microsoft Active Directory with AWS SSO by using AWS Directory Service for Microsoft Active Directory.

B. Enable AWS Single Sign-On (AWS SSO) from the AWS SSO console. Create a two-way forest trust to connect the company's self-managed Microsoft Active Directory with AWS SSO by using AWS Directory Service for Microsoft Active Directory.

C. Use AWS Directory Service. Create a two-way trust relationship with the company's self-managed Microsoft Active Directory.

D. Deploy an identity provider (IdP) on premises. Enable AWS Single Sign-On (AWS SSO) from the AWS SSO console.

(单选题) Q239. A company provides a Voice over Internet Protocol (VoIP) service that uses UDP connections. The service consists of Amazon EC2 instances that run in an Auto Scaling group. The company has deployments across multiple AWS Regions. The company needs to route users to the Region with the lowest latency. The company also needs automated failover between Regions. Which solution will meet these requirements ?

A. Deploy a Network Load Balancer (NLB) and an associated target group. Associate the target group with the Auto Scaling group. Use the NLB as an AWS Global Accelerator endpoint in each Region.

B. Deploy an Application Load Balancer (ALB) and an associated target group. Associate the target group with the Auto Scaling group. Use the ALB as an AWS Global Accelerator endpoint in each Region.

C. Deploy a Network Load Balancer (NLB) and an associated target group. Associate the target group with the Auto Scaling group. Create an Amazon Route 53 latency

record that points to aliases for each NLB.Create an Amazon CloudFront distribution that uses the latency record as an origin.

D.Deploy an Application Load Balancer (ALB) and an associated target group.Associate the target group with the Auto Scaling group.Create an Amazon Route 53 weighted record that points to aliases for each ALB.Deploy an Amazon CloudFront distribution that uses the weighted record as an origin.

(单选题) Q240.A company has an on-premises application that generates a large amount of time-sensitive data that is backed up to Amazon S3.The application has grown and there are user complaints about internet bandwidth limitations.A solutions architect needs to design a long-term solution that allows for both timely backups to Amazon S3 and with minimal impact on internet connectivity for internal users.Which solution meets these requirements ?

A.Establish AWS VPN connections and proxy all traffic through a VPC gateway endpoint.

B.Establish a new AWS Direct Connect connection and direct backup traffic through this new connection.

C.Order daily AWS Snowball devices.Load the data onto the Snowball devices and return the devices to AWS each day.

D.Submit a support ticket through the AWS Management Console.Request the removal of S3 service limits from the account.

(多选题) Q241.A company has an Amazon S3 bucket that contains critical data.The company must protect the data from accidental deletion.Which combination of steps should a solutions architect take to meet these requirements ? (Choose two.)

A.Enable versioning on the S3 bucket.

B.Enable MFA Delete on the S3 bucket.

C.Create a bucket policy on the S3 bucket.

D.Enable default encryption on the S3 bucket.

E.Create a lifecycle policy for the objects in the S3 bucket.

(单选题) Q242.A company's website uses an Amazon EC2 instance store for its

catalog of items.The company wants to make sure that the catalog is highly available and that the catalog is stored in a durable location.What should a solutions architect do to meet these requirements ?

- A.Move the catalog to Amazon ElastiCache for Redis.
- B.Deploy a larger EC2 instance with a larger instance store.
- C.Move the catalog from the instance store to Amazon S3 Glacier Deep Archive.
- D.Move the catalog to an Amazon Elastic File System (Amazon EFS) file system.

(单选题) Q243.A company stores call transcript files on a monthly basis.Users access the files randomly within 1 year of the call, but users access the files infrequently after 1 year.The company wants to optimize its solution by giving users the ability to query and retrieve files that are less than 1-year-old as quickly as possible.A delay in retrieving older files is acceptable.Which solution will meet these requirements MOST cost-effectively ?

- A.Store individual files with tags in Amazon S3 Glacier Instant Retrieval.Query the tags to retrieve the files from S3 Glacier Instant Retrieval.在Amazon S3 Glacier Instant Retrieval
- B.Store individual files in Amazon S3 Intelligent-Tiering.Use S3 Lifecycle policies to move the files to S3 Glacier Flexible Retrieval after 1 year.Query and retrieve the files that are in Amazon S3 by using Amazon Athena.Query and retrieve the files that are in S3 Glacier by using S3 Glacier Select.
- C.Store individual files with tags in Amazon S3 Standard storage.Store search metadata for each archive in Amazon S3 Standard storage.Use S3 Lifecycle policies to move the files to S3 Glacier Instant Retrieval after 1 year.Query and retrieve the files by searching for metadata from Amazon S3.
- D.Store individual files in Amazon S3 Standard storage.Use S3 Lifecycle policies to move the files to S3 Glacier Deep Archive after 1 year.Store search metadata in Amazon RDS.Query the files from Amazon RDS.Retrieve the files from S3 Glacier Deep Archive.

(单选题) Q244.A company has a production workload that runs on 1,000 Amazon EC2 Linux instances.The workload is powered by third-party software.The

company needs to patch the third-party software on all EC2 instances as quickly as possible to remediate a critical security vulnerability.What should a solutions architect do to meet these requirements ?

- A.Create an AWS Lambda function to apply the patch to all EC2 instances.
- B.Configure AWS Systems Manager Patch Manager to apply the patch to all EC2 instances.
- C.Schedule an AWS Systems Manager maintenance window to apply the patch to all EC2 instances.
- D.Use AWS Systems Manager Run Command to run a custom command that applies the patch to all EC2 instances.

(单选题) Q245.A company runs multiple Windows workloads on AWS.The company's employees use Windows file shares that are hosted on two Amazon EC2 instances.The file shares synchronize data between themselves and maintain duplicate copies.The company wants a highly available and durable storage solution that preserves how users currently access the files.What should a solutions architect do to meet these requirements ?

- A.Migrate all the data to Amazon S3.Set up IAM authentication for users to access files.
- B.Set up an Amazon S3 File Gateway.Mount the S3 File Gateway on the existing EC2 instances.
- C.Extend the file share environment to Amazon FSx for Windows File Server with a Multi-AZ configuration.Migrate all the data to FSx for Windows File Server.
- D.Extend the file share environment to Amazon Elastic File System (Amazon EFS) with a Multi-AZ configuration.Migrate all the data to Amazon EFS.

(单选题) Q246.A company has an AWS Glue extract, transform, and load (ETL) job that runs every day at the same time.The job processes XML data that is in an Amazon S3 bucket.New data is added to the S3 bucket every day.A solutions architect notices that AWS Glue is processing all the data during each run.What should the solutions architect do to prevent AWS Glue from reprocessing old data ?

- A.Edit the job to use job bookmarks.
- B.Edit the job to delete data after the data is processed.
- C.Edit the job by setting the NumberOfWorkers field to 1,
- D.Use a FindMatches machine learning (ML) transform.

(单选题) Q247.A company uses 50 TB of data for reporting.The company wants to move this data from on premises to AWS.A custom application in the company's data center runs a weekly data transformation job.The company plans to pause the application until the data transfer is complete and needs to begin the transfer process as soon as possible.The data center does not have any available network bandwidth for additional workloads.A solutions architect must transfer the data and must configure the transformation job to continue to run in the AWS Cloud.Which solution will meet these requirements with the LEAST operational overhead ?

- A.Use AWS DataSync to move the data.Create a custom transformation job by using AWS Glue.使用AWS DataSync
- B.Order an AWS Snowcone device to move the data.Deploy the transformation application to the device.
- C.Order an AWS Snowball Edge Storage Optimized device.Copy the data to the device.Create a custom transformation job by using AWS Glue.
- D.Order an AWS Snowball Edge Storage Optimized device that includes Amazon EC2 compute.Copy the data to the device.Create a new EC2 instance on AWS to run the transformation application.

(单选题) Q248.A company is building a web-based application running on Amazon EC2 instances in multiple Availability Zones.The web application will provide access to a repository of text documents totaling about 900 TB in size.The company anticipates that the web application will experience periods of high demand.A solutions architect must ensure that the storage component for the text documents can scale to meet the demand of the application at all times.The company is concerned about the overall cost of the solution.Which storage solution meets these requirements MOST cost-effectively ?

- A.Amazon Elastic Block Store (Amazon EBS)

- B.Amazon Elastic File System (Amazon EFS)
- C.Amazon OpenSearch Service (Amazon Elasticsearch Service)
- D.Amazon S3

(单选题) Q249.A company uses AWS Organizations to create dedicated AWS accounts for each business unit to manage each business unit's account independently upon request.The root email recipient missed a notification that was sent to the root user email address of one account.The company wants to ensure that all future notifications are not missed.Future notifications must be limited to account administrators.Which solution will meet these requirements ?

- A.Configure the company's email server to forward notification email messages that are sent to the AWS account root user email address to all users in the organization.
- B.Configure all AWS account root user email addresses as distribution lists that go to a few administrators who can respond to alerts.Configure AWS account alternate contacts in the AWS Organizations console or programmatically.
- C.Configure all AWS account root user email messages to be sent to one administrator who is responsible for monitoring alerts and forwarding those alerts to the appropriate groups.
- D.Configure all existing AWS accounts and all newly created accounts to use the same root user email address.Configure AWS account alternate contacts in the AWS Organizations console or programmatically.

Answer:B

(单选题) Q250.A large media company hosts a web application on AWS.The company wants to start caching confidential media files so that users around the world will have reliable access to the files.The content is stored in Amazon S3 buckets.The company must deliver the content quickly, regardless of where the requests originate geographically.Which solution will meet these requirements ?

- A.Use AWS DataSync to connect the S3 buckets to the web application.
- B.Deploy AWS Global Accelerator to connect the S3 buckets to the web application.
- C.Deploy Amazon CloudFront to connect the S3 buckets to CloudFront edge servers.
- D.Use Amazon Simple Queue Service (Amazon SQS) to connect the S3 buckets to

the web application.

(单选题) Q251.An ecommerce company has an order-processing application that uses Amazon API Gateway and an AWS Lambda function.The application stores data in an Amazon Aurora PostgreSQL database.During a recent sales event, a sudden surge in customer orders occurred.Some customers experienced timeouts, and the application did not process the orders of those customers.A solutions architect determined that the CPU utilization and memory utilization were high on the database because of a large number of open connections.The solutions architect needs to prevent the timeout errors while making the least possible changes to the application.Which solution will meet these requirements ?

A.Configure provisioned concurrency for the Lambda function.Modify the database to be a global database in multiple AWS Regions.

B.Use Amazon RDS Proxy to create a proxy for the database.Modify the Lambda function to use the RDS Proxy endpoint instead of the database endpoint.

C.Create a read replica for the database in a different AWS Region.Use query string parameters in API Gateway to route traffic to the read replica,

D.Migrate the data from Aurora PostgreSQL to Amazon DynamoDB by using AWS Database Migration Service (AWS DMS) .Modify the Lambda function to use the DynamoDB table.

(单选题) Q252.A company's order system sends requests from clients to Amazon EC2 instances.The EC2 instances process the orders and then store the orders in a database on Amazon RDS.Users report that they must reprocess orders when the system fails.The company wants a resilient solution that can process orders automatically if a system outage occurs.What should a solutions architect do to meet these requirements ?

A.Move the EC2 instances into an Auto Scaling group.Create an Amazon EventBridge (Amazon CloudWatch Events) rule to target an Amazon Elastic Container Service (Amazon ECS) task.

B.Move the EC2 instances into an Auto Scaling group behind an Application Load Balancer (ALB) .Update the order system to send messages to the ALB endpoint.

C.Move the EC2 instances into an Auto Scaling group.Configure the order system to send messages to an Amazon Simple Queue Service (Amazon SQS) queue.Configure the EC2 instances to consume messages from the queue.

D.Create an Amazon Simple Notification Service (Amazon SNS) topic.Create an AWS Lambda function, and subscribe the function to the SNS topic.Configure the order system to send messages to the SNS topic.Send a command to the EC2 instances to process the messages by using AWS Systems Manager Run Command.

(单选题) Q253.A company is developing a marketing communications service that targets mobile app users.The company needs to send confirmation messages with Short Message Service (SMS) to its users.The users must be able to reply to the SMS messages.The company must store the responses for a year for analysis.What should a solutions architect do to meet these requirements ?

A.Create an Amazon Connect contact flow to send the SMS messages.Use AWS Lambda to process the responses.

B.Build an Amazon Pinpoint journey.Configure Amazon Pinpoint to send events to an Amazon Kinesis data stream for analysis and archiving.

C.Use Amazon Simple Queue Service (Amazon SQS) to distribute the SMS messages.Use AWS Lambda to process the responses.

D.Create an Amazon Simple Notification Service (Amazon SNS) FIFO topic.Subscribe an Amazon Kinesis data stream to the SNS topic for analysis and archiving.

(单选题) Q254.The customers of a finance company request appointments with financial advisors by sending text messages.A web application that runs on Amazon EC2 instances accepts the appointment requests.The text messages are published to an Amazon Simple Queue Service (Amazon SQS) queue through the web application.Another application that runs on EC2 instances then sends meeting invitations and meeting confirmation email messages to the customers.After successful scheduling, this application stores the meeting information in an Amazon DynamoDB database.As the company expands, customers report that their meeting invitations are taking longer to arrive.What should a solutions architect recommend

to resolve this issue ?

- A.Add a DynamoDB Accelerator (DAX) cluster in front of the DynamoDB database.
- B.Add an Amazon API Gateway API in front of the web application that accepts the appointment requests.
- C.Add an Amazon CloudFront distribution.Set the origin as the web application that accepts the appointment requests.添加Amazon CloudFront
- D.Add an Auto Scaling group for the application that sends meeting invitations.Configure the Auto Scaling group to scale based on the depth of the SQS queue.

Q255. A Solutions Architect is designing a microservices-based application using Amazon ECS. The application includes a WebSocket component, and the traffic needs to be distributed between microservices based on the URL.

Which service should the Architect choose to distribute the workload?

- A. ELB Classic Load Balancer
- B. Amazon Route 53 DNS
- C. ELB Application Load Balancer
- D. Amazon CloudFront

Q256. A Solutions Architect is designing the storage layer for a production relational database. The database will run on Amazon EC2. The database is accessed by an application that performs intensive reads and writes, so the database requires the LOWEST random I/O latency. Which data storage method fulfills the above requirements?

- A. Store data in a filesystem backed by Amazon Elastic File System (EFS).
- B. Store data in Amazon S3 and use a third-party solution to expose Amazon S3 as a filesystem to the database server.
- C. Store data in Amazon Dynamo DB and emulate relational database semantics.
- D.Stripe data across multiple Amazon EBS volumes using RAID .

仅个人学习参考 禁止外传
店长微信：hj feng128